

# Assembly Programming for the

Chris Patuzzo



# About me

London Computation Club

Software Developer

Hard search problems

# This talk

## How to get started

## Basics of CPU architecture

## Solving Hamiltonian path

By the end

Access to a powerful tool

Know more about your computer

Where to go next

Access to a powerful tool

Access to a powerful tool

216,000,000 iterations/second

```
Benchmark.realtime { 1.upto(216_000_000) {} }  
#=> 7.0587
```

# What is this?



# “System on a chip”

CPU



“Secure Enclave”

GPU

“Neural Engine”

Memory

# “System on a chip”

CPU



“Secure Enclave”

GPU

“Neural Engine”

Memory

# CPU

Runs ARM64 (v8.5-A)

RISC architecture

Low power, less heat

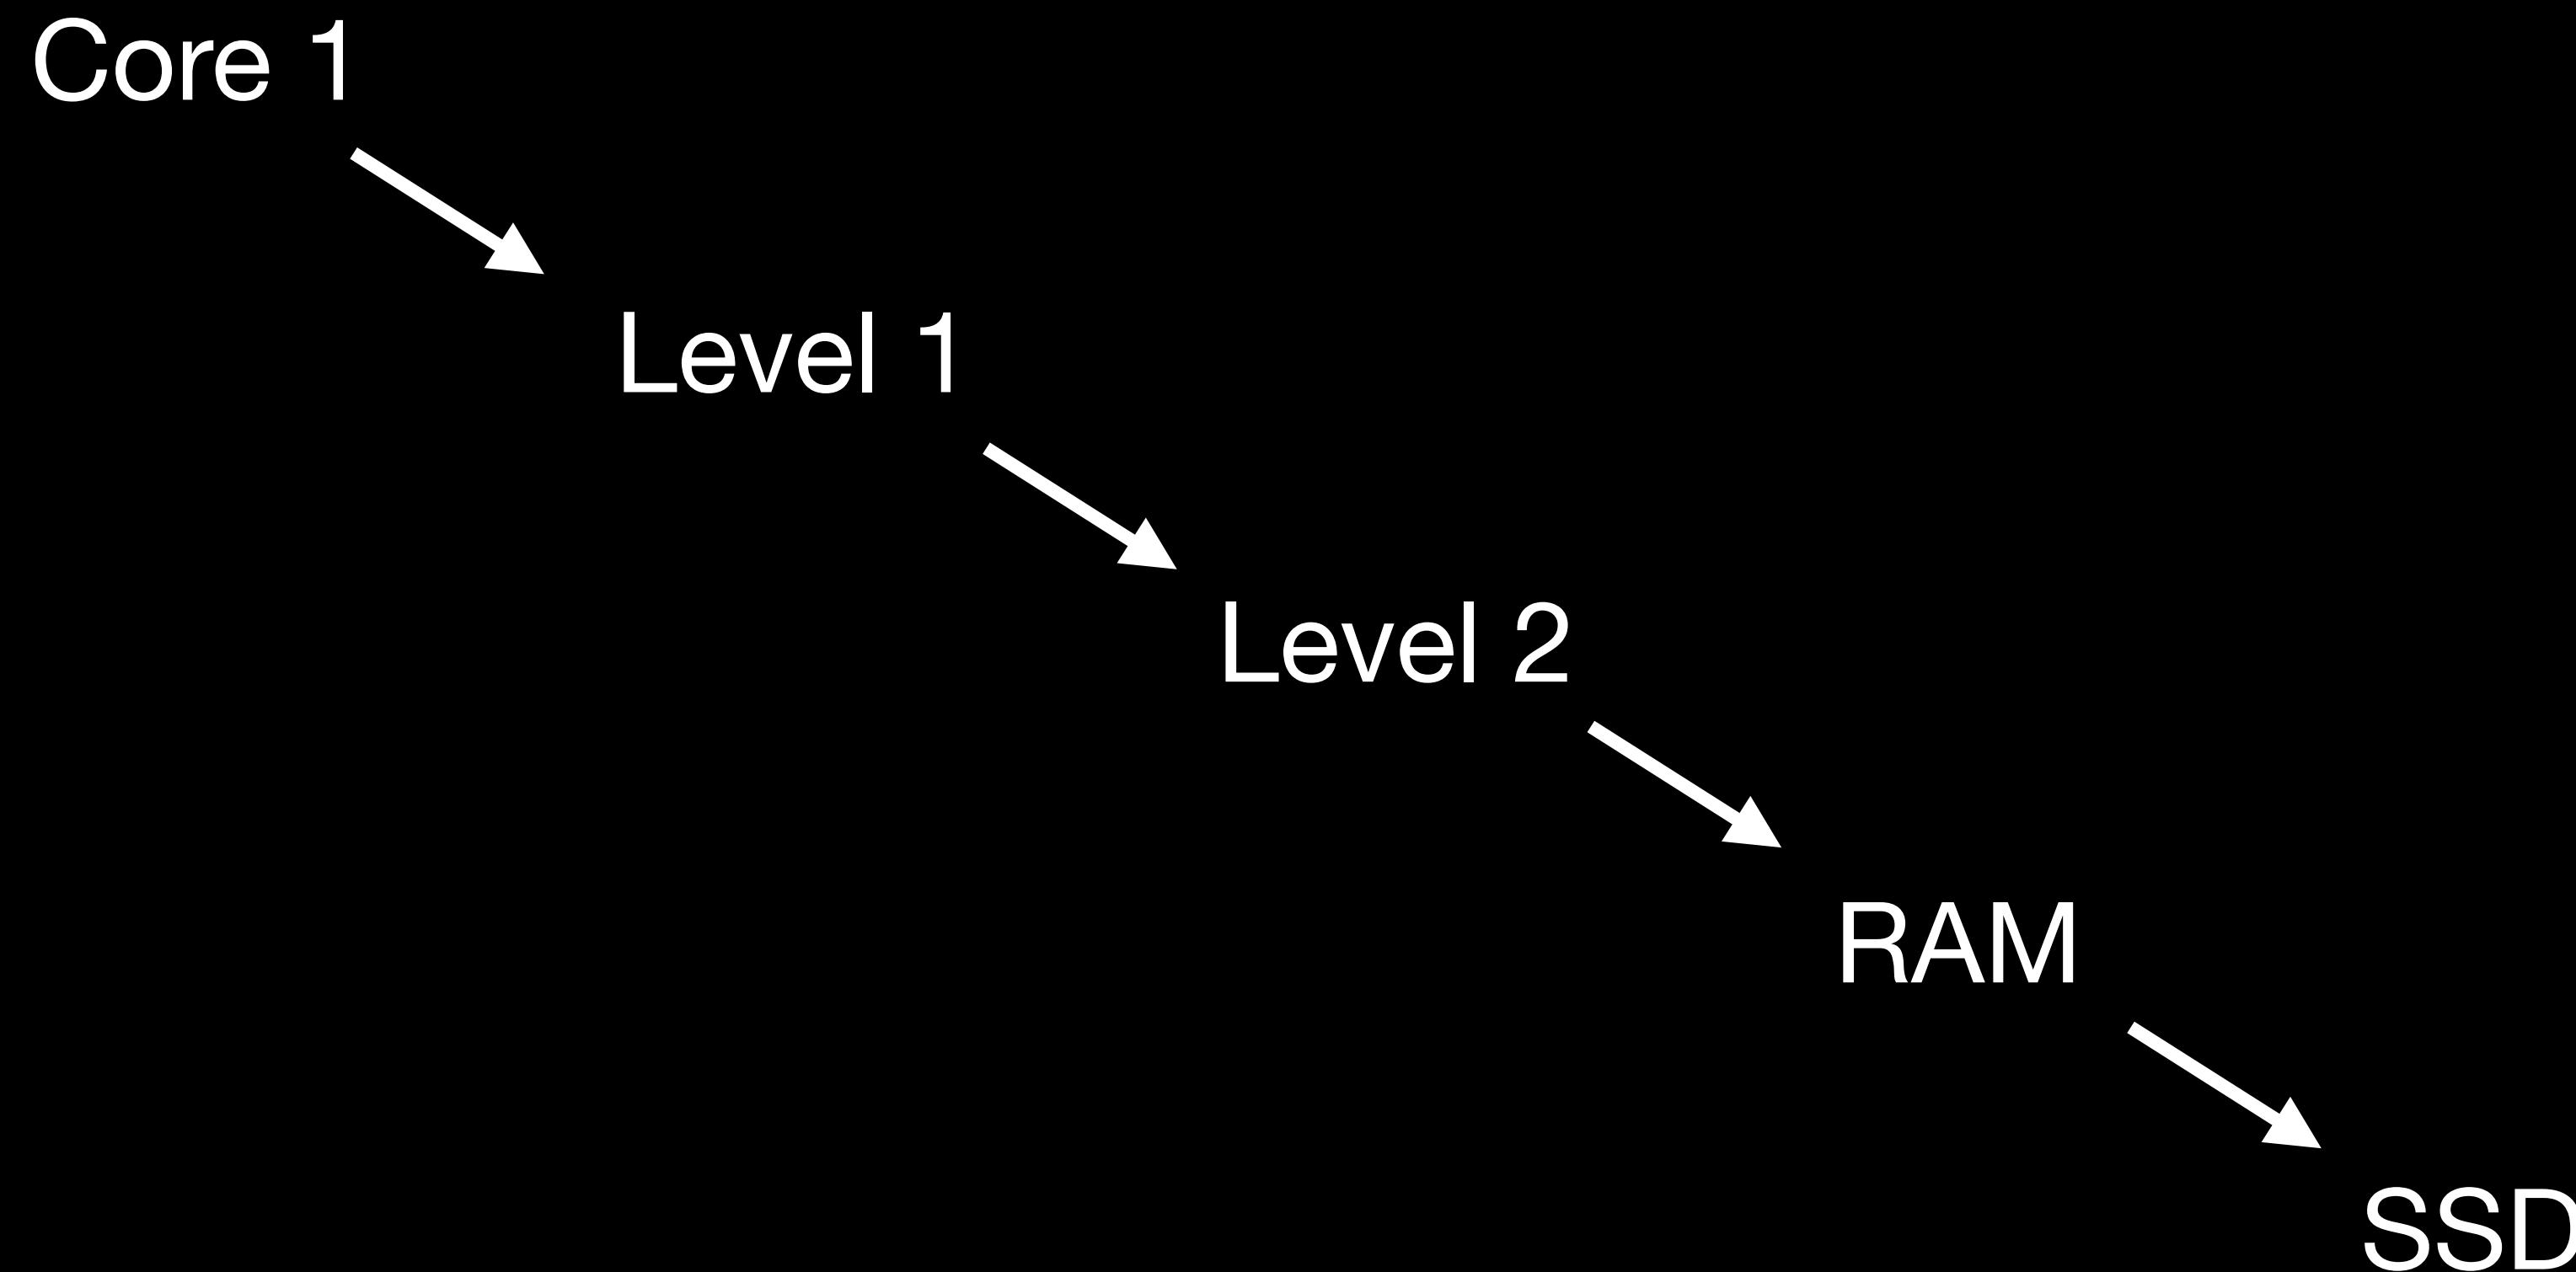
# CPU

4 “Firestorm” cores

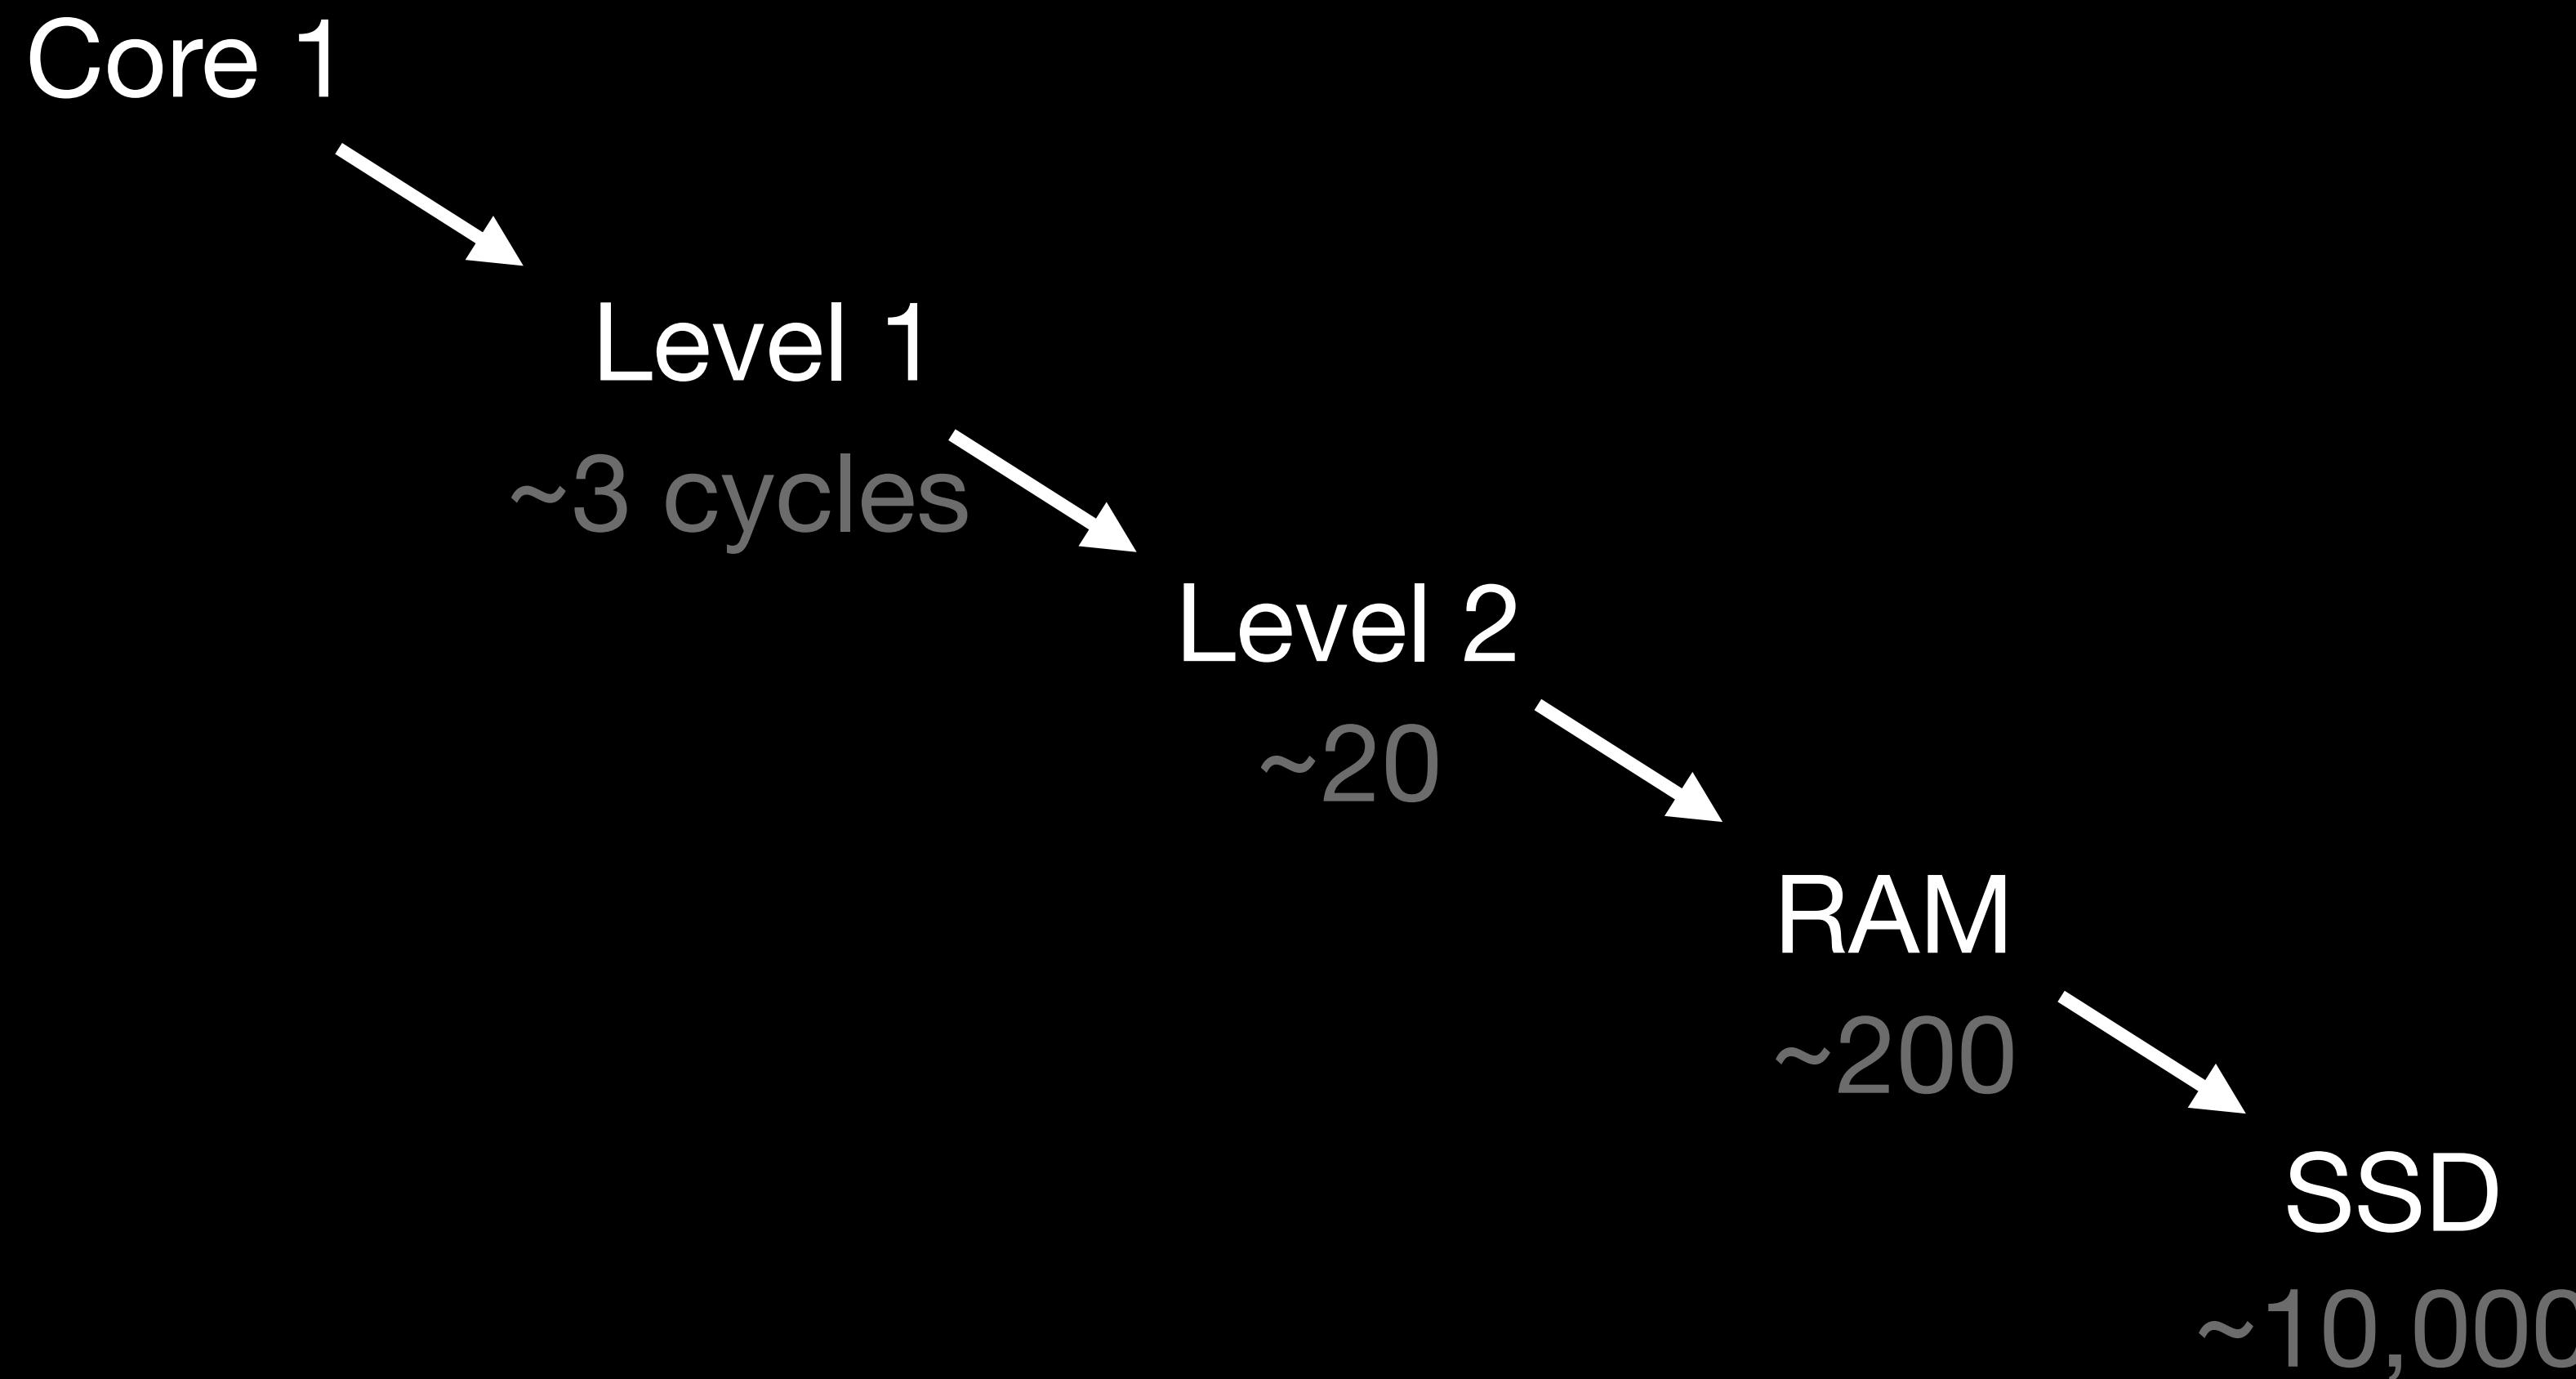
4 “Icestorm” cores

(different speeds / cache)

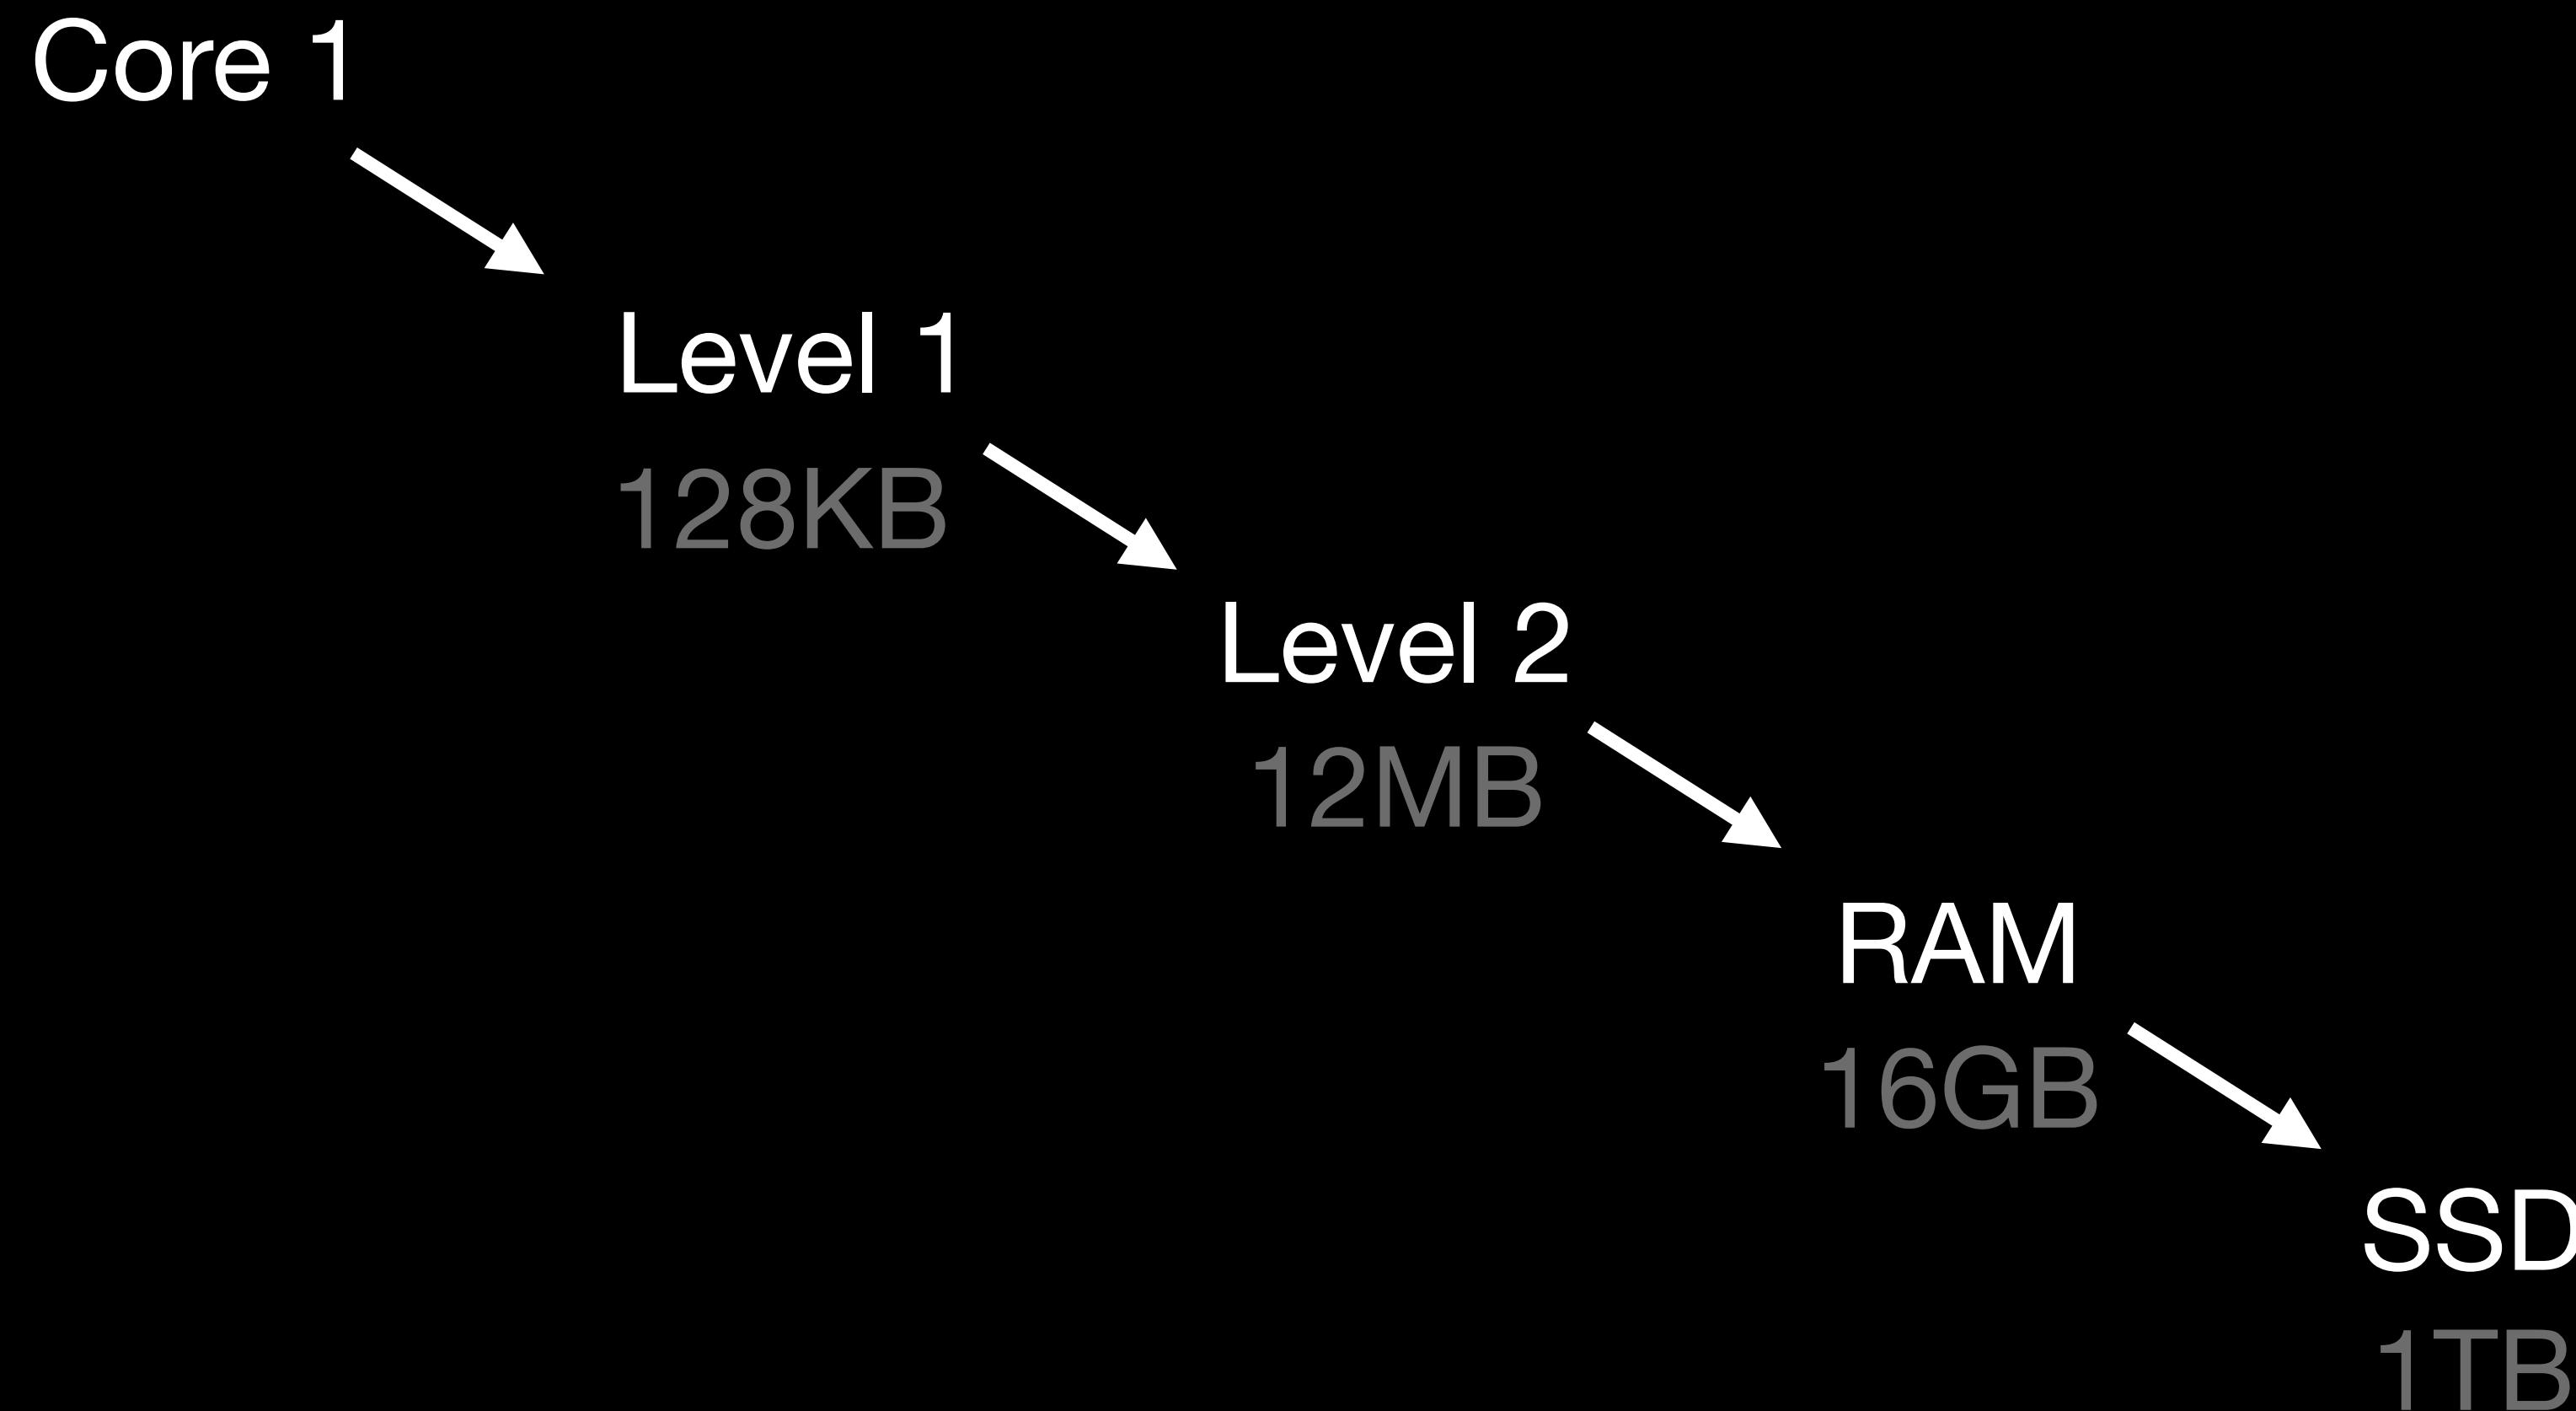
# What is cache?



# What is cache?

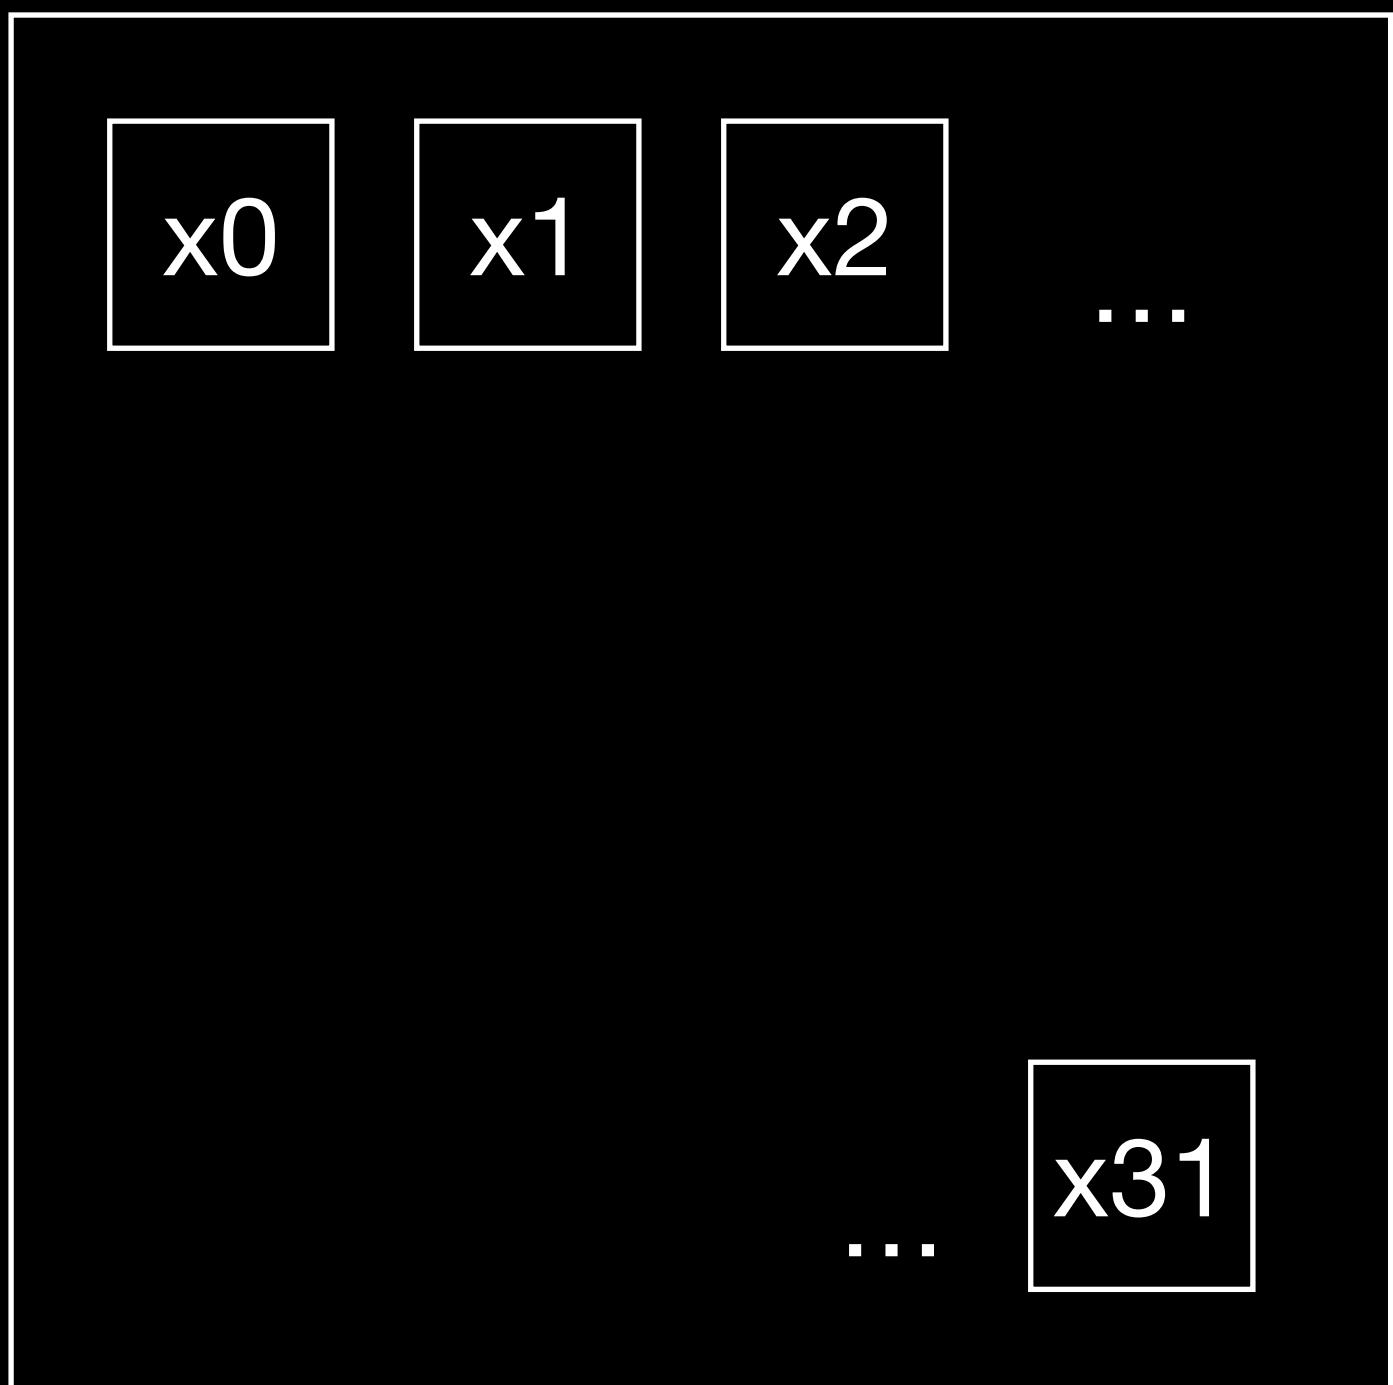


# What is cache?



# Registers

Core 1



32 registers (1 cycle)

64 bits each

Some are special

# Load Store Architecture

Load from memory

Calculate something

Store into memory

```
ldr x0, [x5]
mov x1, #12345
add x0, x0, x1
str x0, [x5]
```

```
ldr x0, [x5] ← x0 = *x5
mov x1, #12345
add x0, x0, x1
str x0, [x5]
```

```
ldr x0, [x5]
mov x1, #12345 ← x1 = 12345
add x0, x0, x1
str x0, [x5]
```

```
ldr x0, [x5]
mov x1, #12345
add x0, x0, x1
str x0, [x5]
```

$$x0 = x0 + x1$$


```
ldr x0, [x5]
mov x1, #12345
add x0, x0, x1
str x0, [x5] ← *x5 = x0
```

# Service Calls

Supervising process

Ask it to do something

e.g. print, exit

```
.global _main  
.align 2
```

```
_main:
```

```
    mov x16, #4  
    mov x0, #1  
    adr x1, hello  
    mov x2, #14  
    svc 0
```

```
    mov x16, #1  
    mov x0, #0  
    svc 0
```

```
hello: .ascii "Hello, world!\n"
```

```
.global _main
.align 2
}
```

\_main:

```
    mov x16, #4
    mov x0, #1
    adr x1, hello
    mov x2, #14
    svc 0
```

```
    mov x16, #1
    mov x0, #0
    svc 0
```

hello: .ascii "Hello, world!\n"

} boilerplate

```
.global _main  
.align 2
```

```
_main:
```

```
    mov x16, #4  
    mov x0, #1  
    adr x1, hello  
    mov x2, #14  
    svc 0
```

supervisor calls

```
    mov x16, #1  
    mov x0, #0  
    svc 0
```

```
hello: .ascii "Hello, world!\n"
```

```
.global _main  
.align 2
```

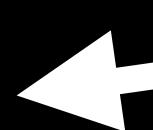
```
_main:
```

```
    mov x16, #4  
    mov x0, #1  
    adr x1, hello  
    mov x2, #14  
    svc 0
```

```
    mov x16, #1  
    mov x0, #0  
    svc 0
```

```
hello: .ascii "Hello, world!\n"
```

print



```
.global _main  
.align 2
```

```
_main:  
    mov x16, #4  
    mov x0, #1 ← stdout  
    adr x1, hello  
    mov x2, #14  
    svc 0
```

```
    mov x16, #1  
    mov x0, #0  
    svc 0
```

```
hello: .ascii "Hello, world!\n"
```

```
.global _main  
.align 2
```

```
_main:
```

```
    mov x16, #4  
    mov x0, #1  
    adr x1, hello  
    mov x2, #14  
    svc 0
```

```
    mov x16, #1  
    mov x0, #0  
    svc 0
```

```
hello: .ascii "Hello, world!\n"
```

address of the string



```
.global _main  
.align 2
```

```
_main:  
    mov x16, #4  
    mov x0, #1  
    adr x1, hello  
    mov x2, #14  
    svc 0
```

```
    mov x16, #1  
    mov x0, #0  
    svc 0
```

```
hello: .ascii "Hello, world!\n"
```

length of the string



```
.global _main  
.align 2
```

```
_main:
```

```
    mov x16, #4  
    mov x0, #1  
    adr x1, hello  
    mov x2, #14  
    svc 0
```

exit

```
    mov x16, #1 ←  
    mov x0, #0  
    svc 0
```

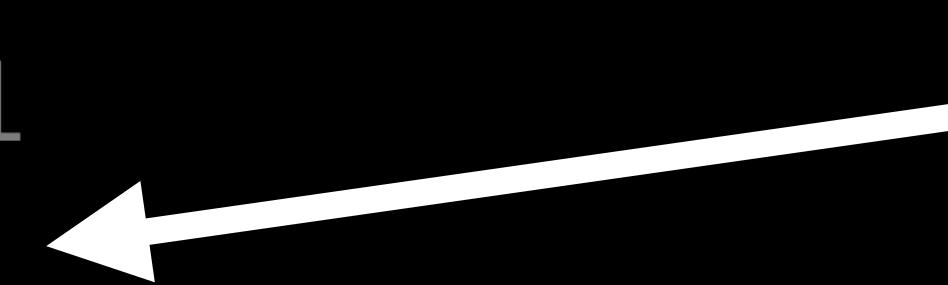
```
hello: .ascii "Hello, world!\n"
```

```
.global _main  
.align 2
```

```
_main:  
    mov x16, #4  
    mov x0, #1  
    adr x1, hello  
    mov x2, #14  
    svc 0
```

```
    mov x16, #1  
    mov x0, #0  
    svc 0
```

exit status



```
hello: .ascii "Hello, world!\n"
```

# Assemble and Link

```
#!/bin/bash

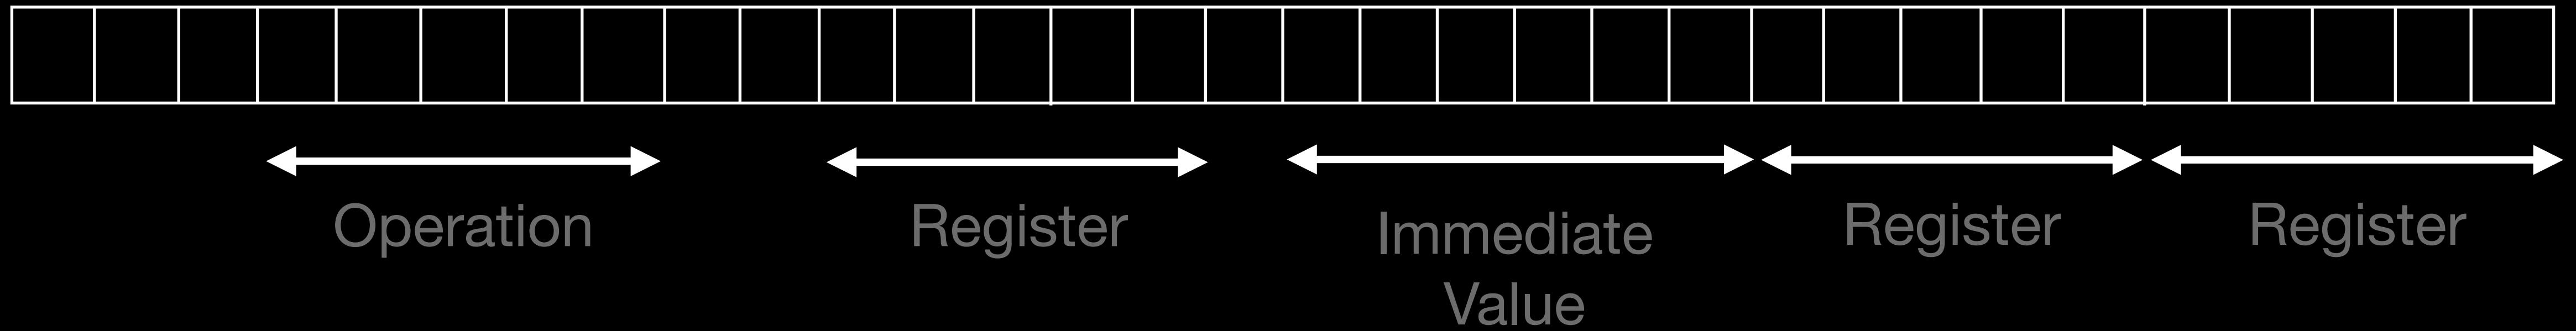
sdk=`xcrun -sdk macOS --show-sdk-path`

as -arch arm64 hello.s -o hello.o
ld -arch arm64 hello.o -o hello -lSystem -syslibroot $sdk
```

```
$ ./hello  
Hello, world!
```

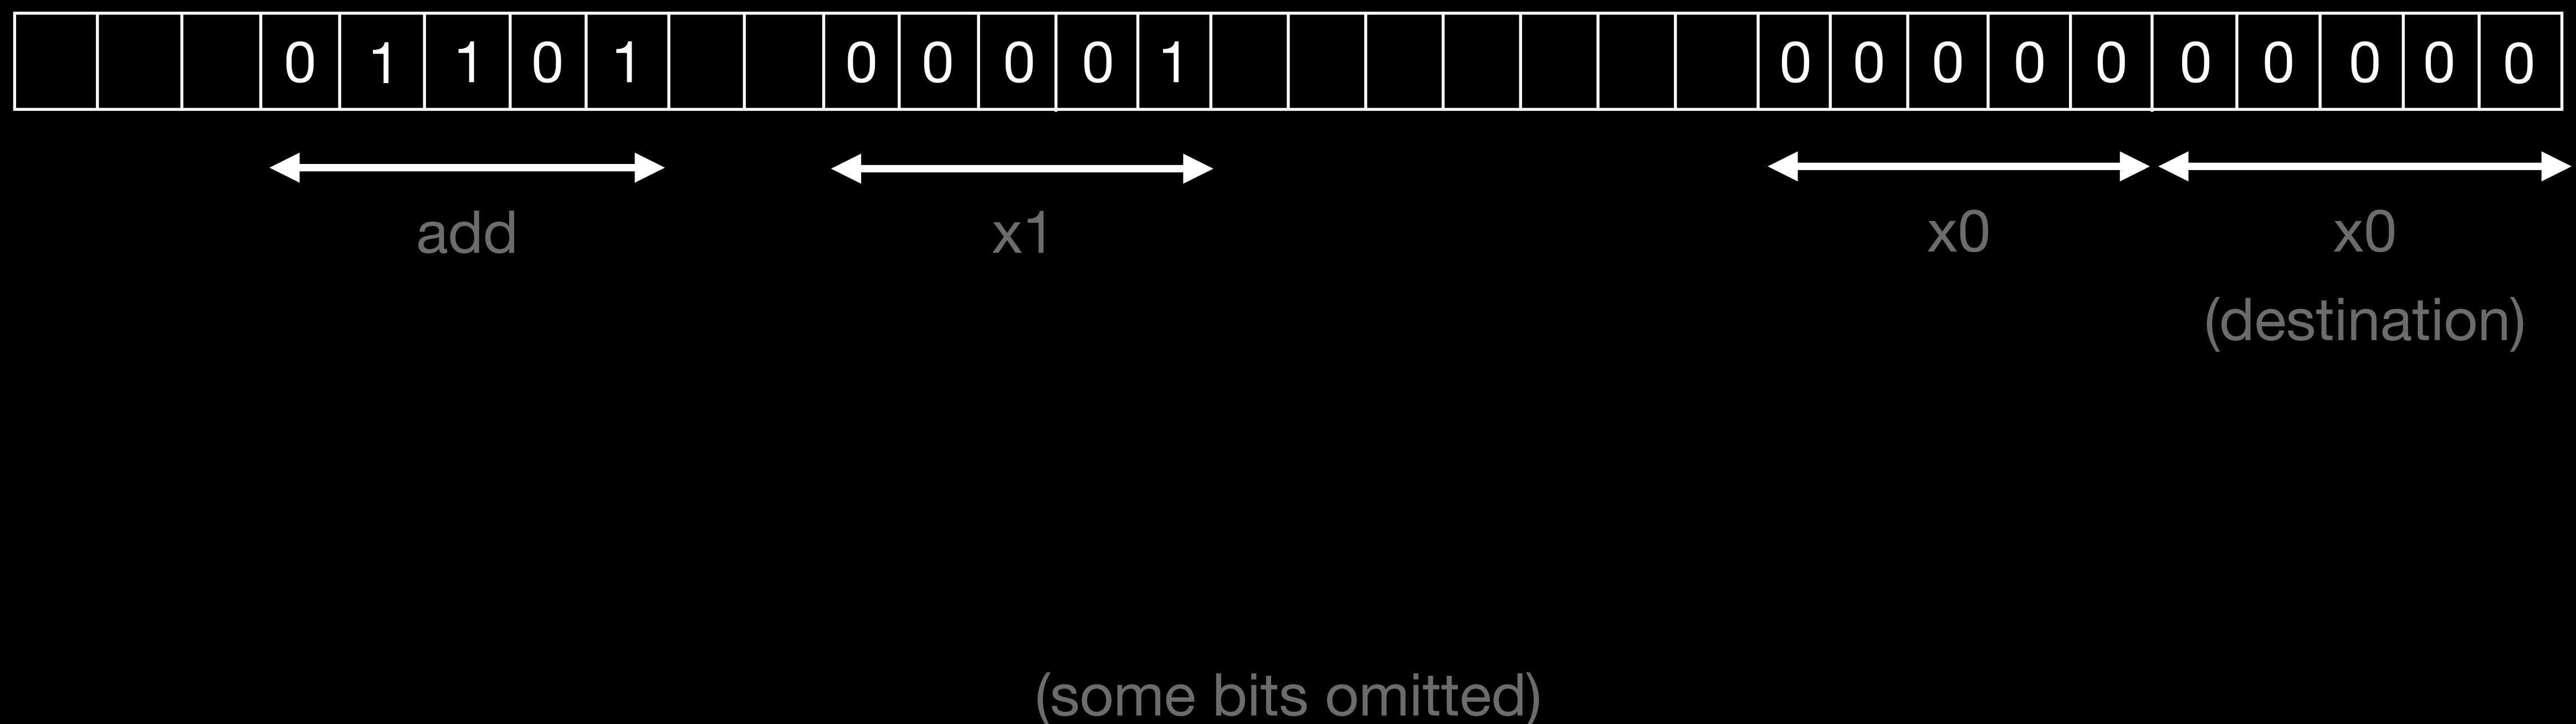
# Instructions

32 bits wide



(some bits omitted)

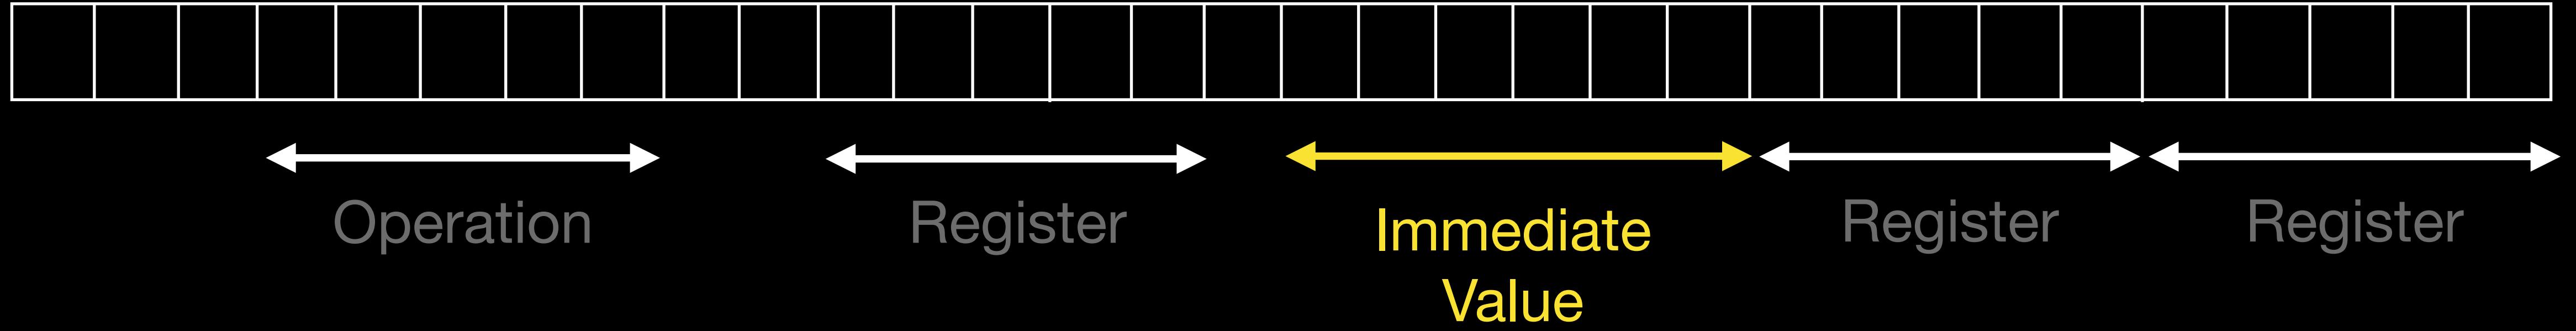
```
add x0, x0, x1
```



```
mov x1, #12345  
add x0, x0, x1
```

versus.

```
add x0, x0, #12345
```



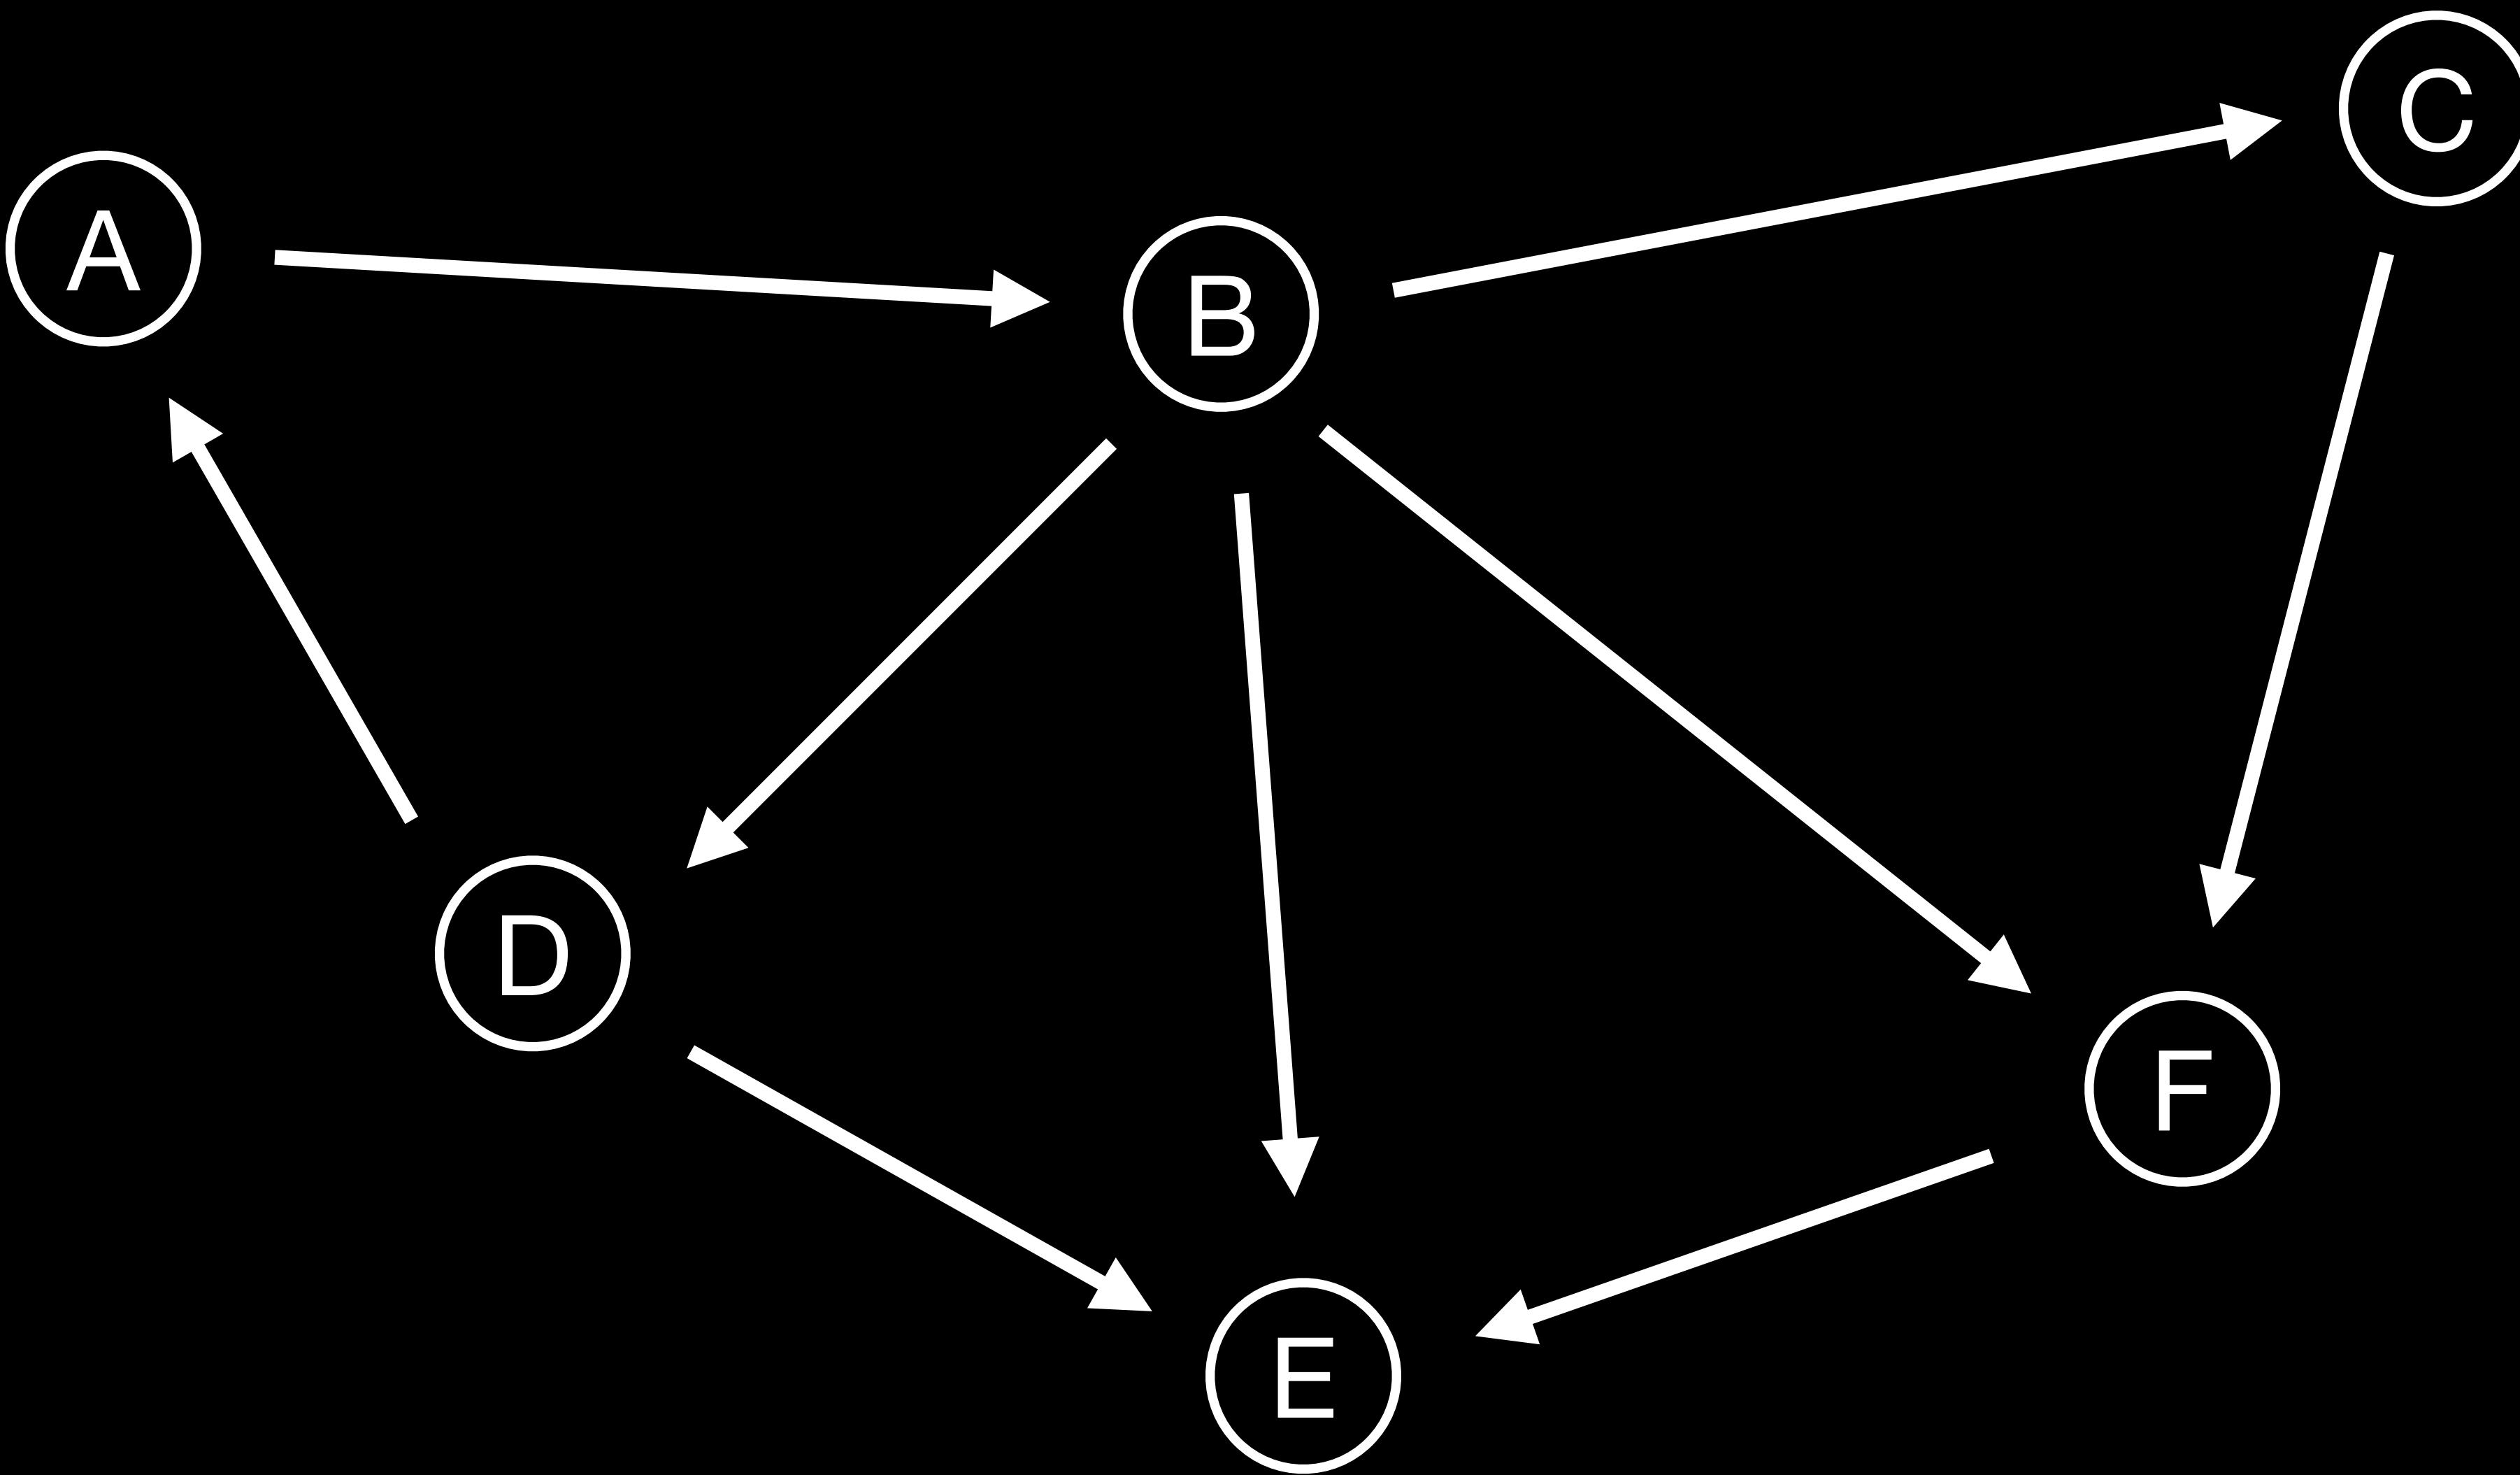
# What next?

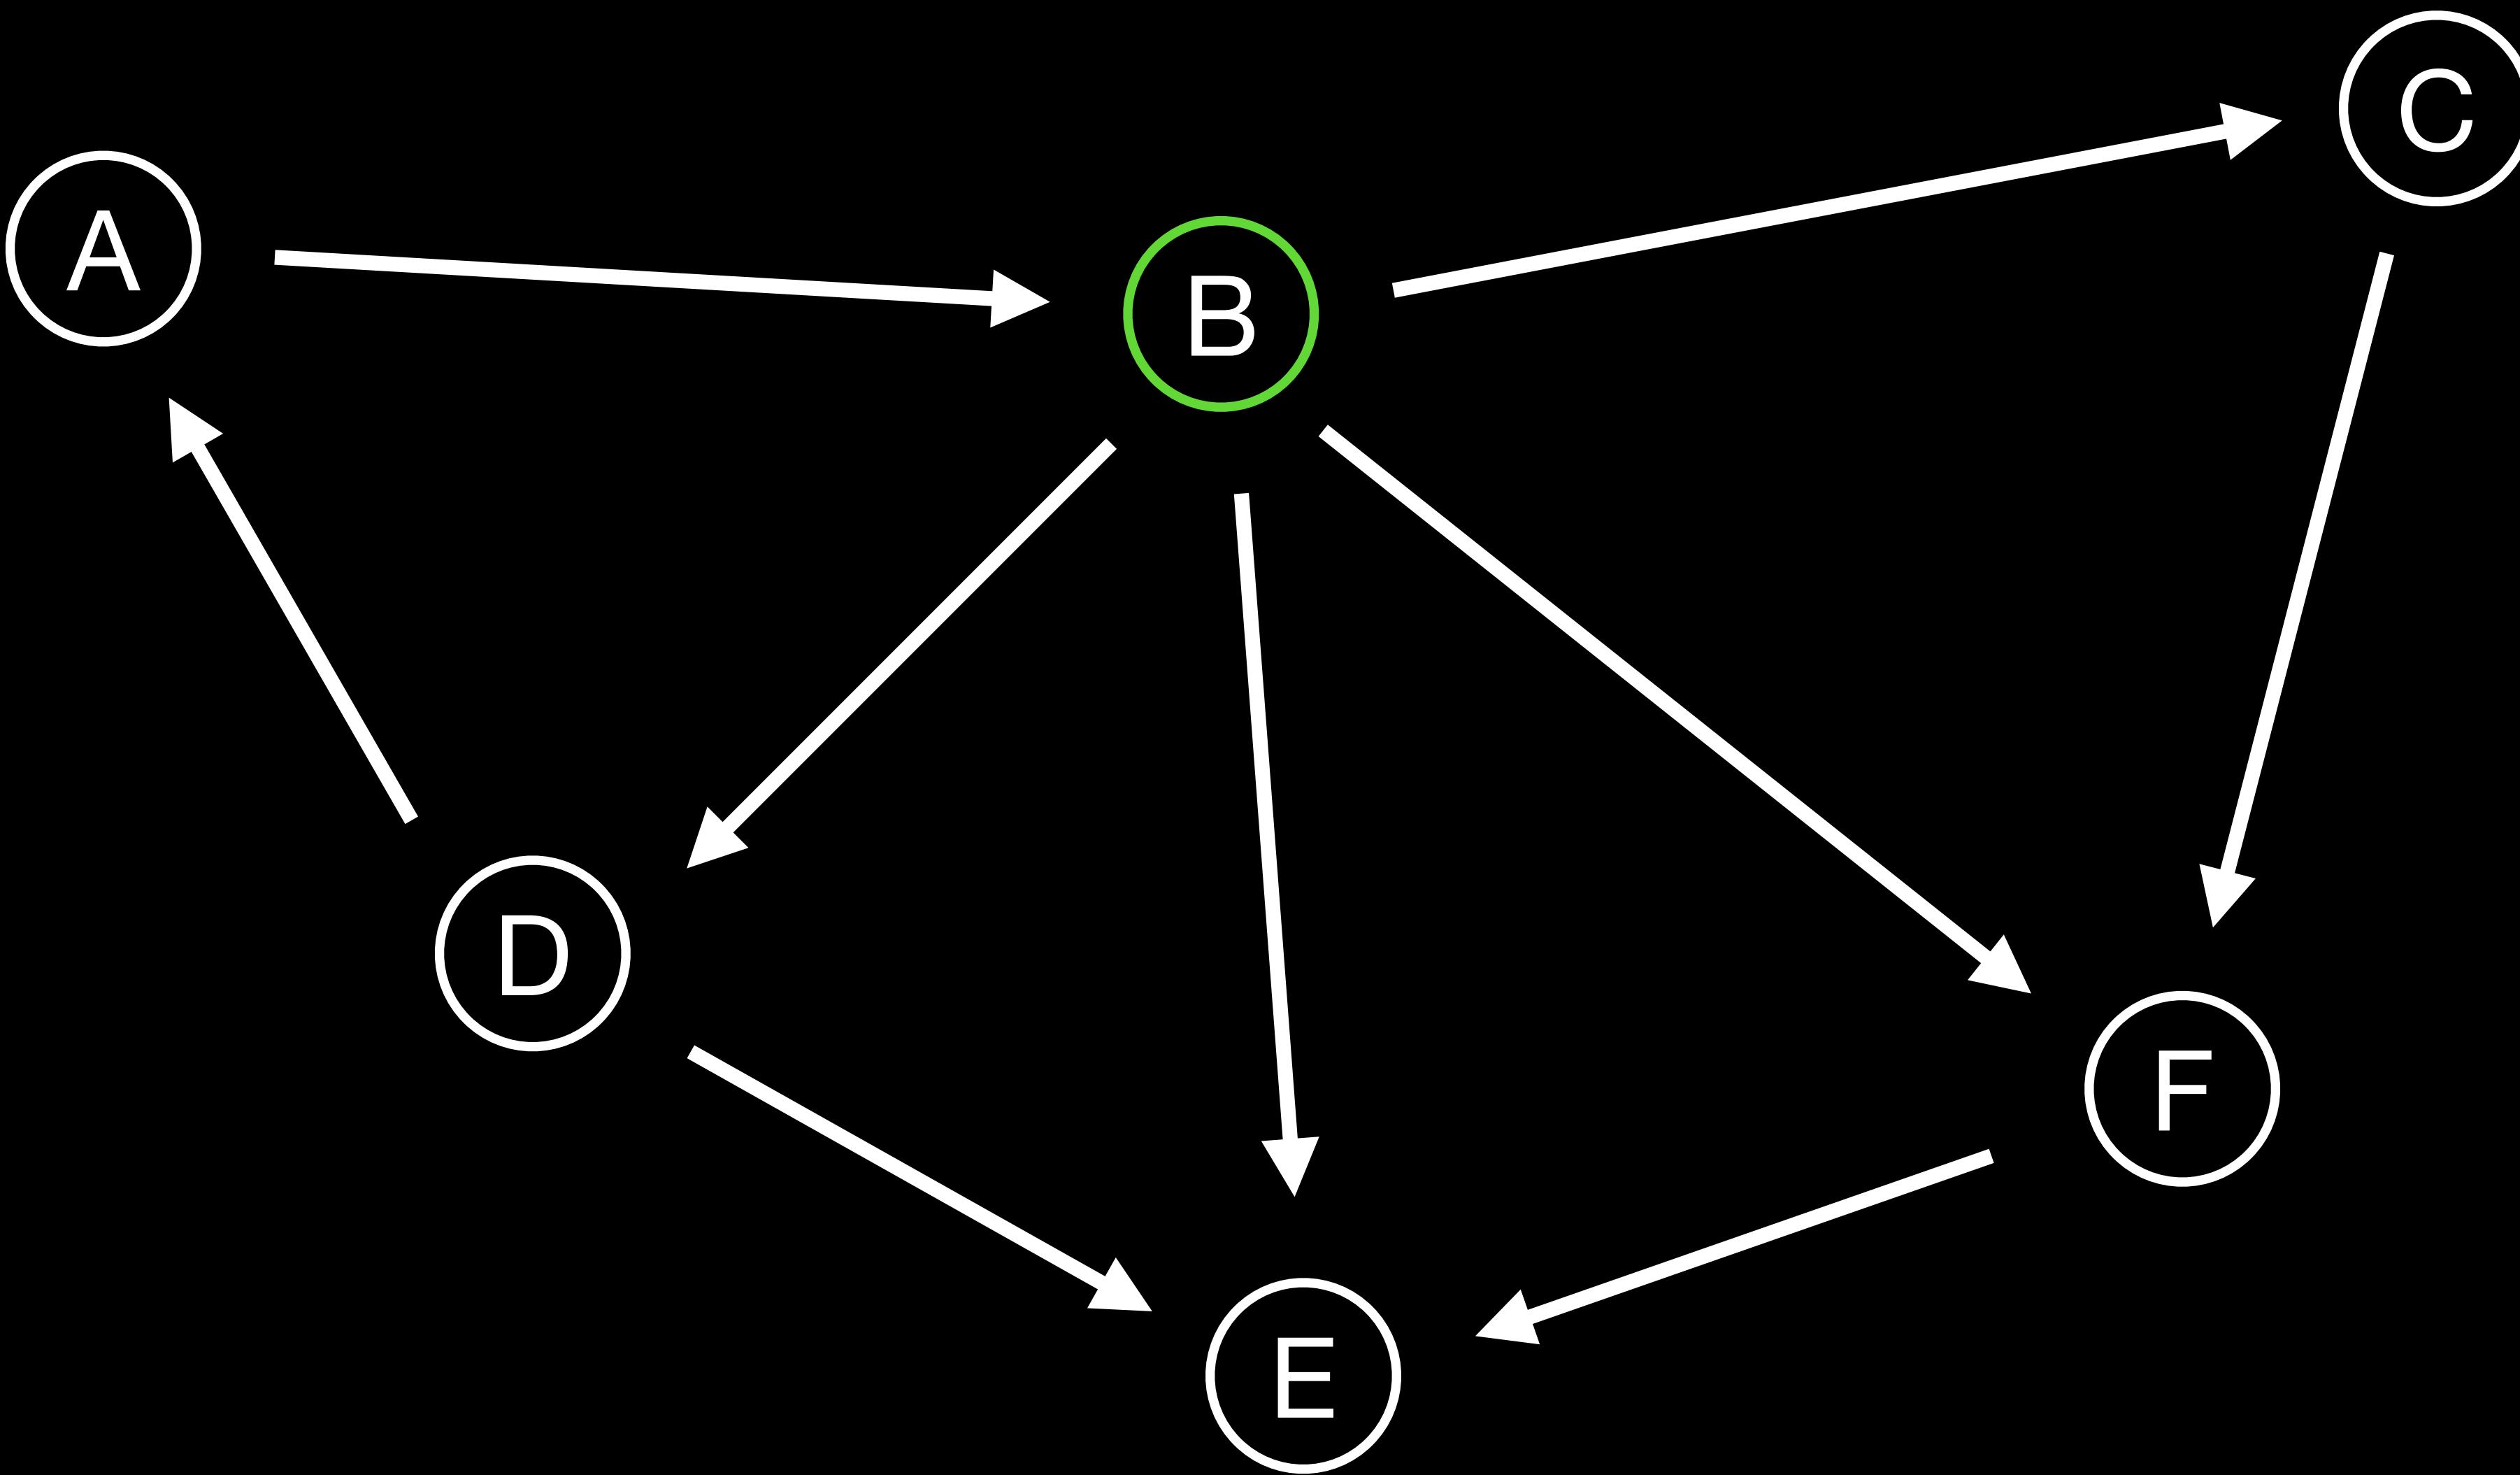
Branching?

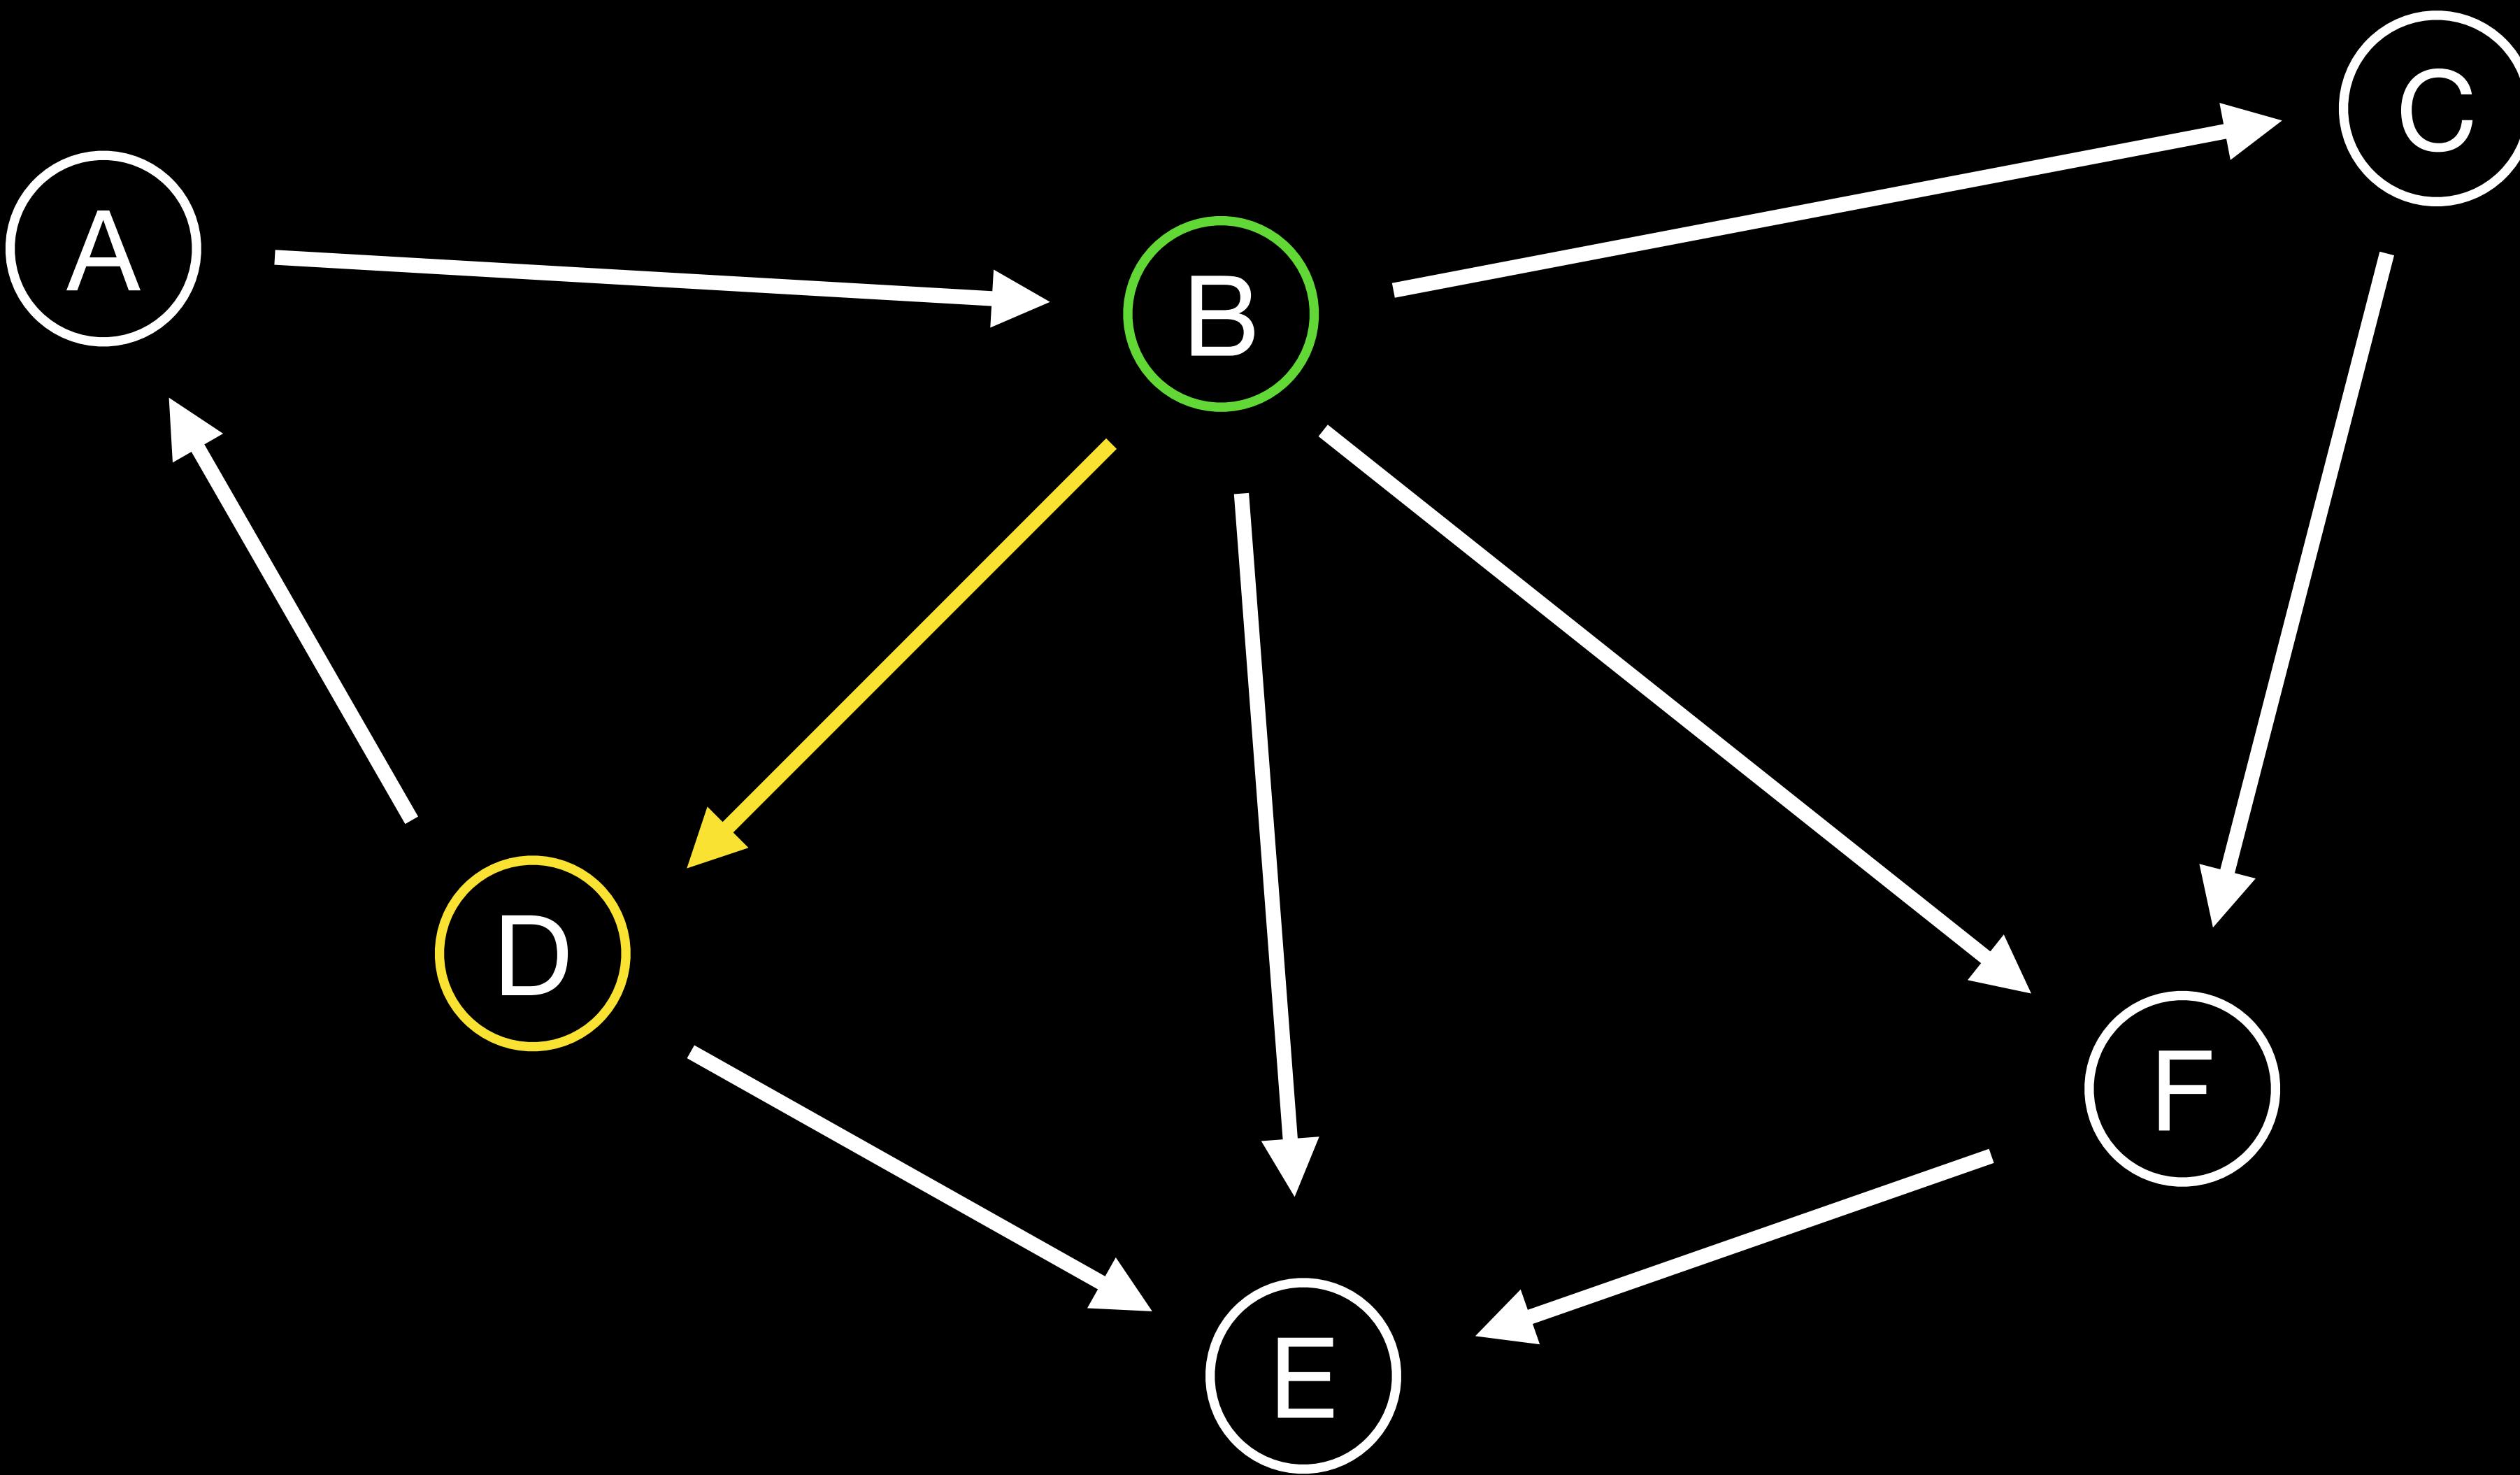
Logical operations?

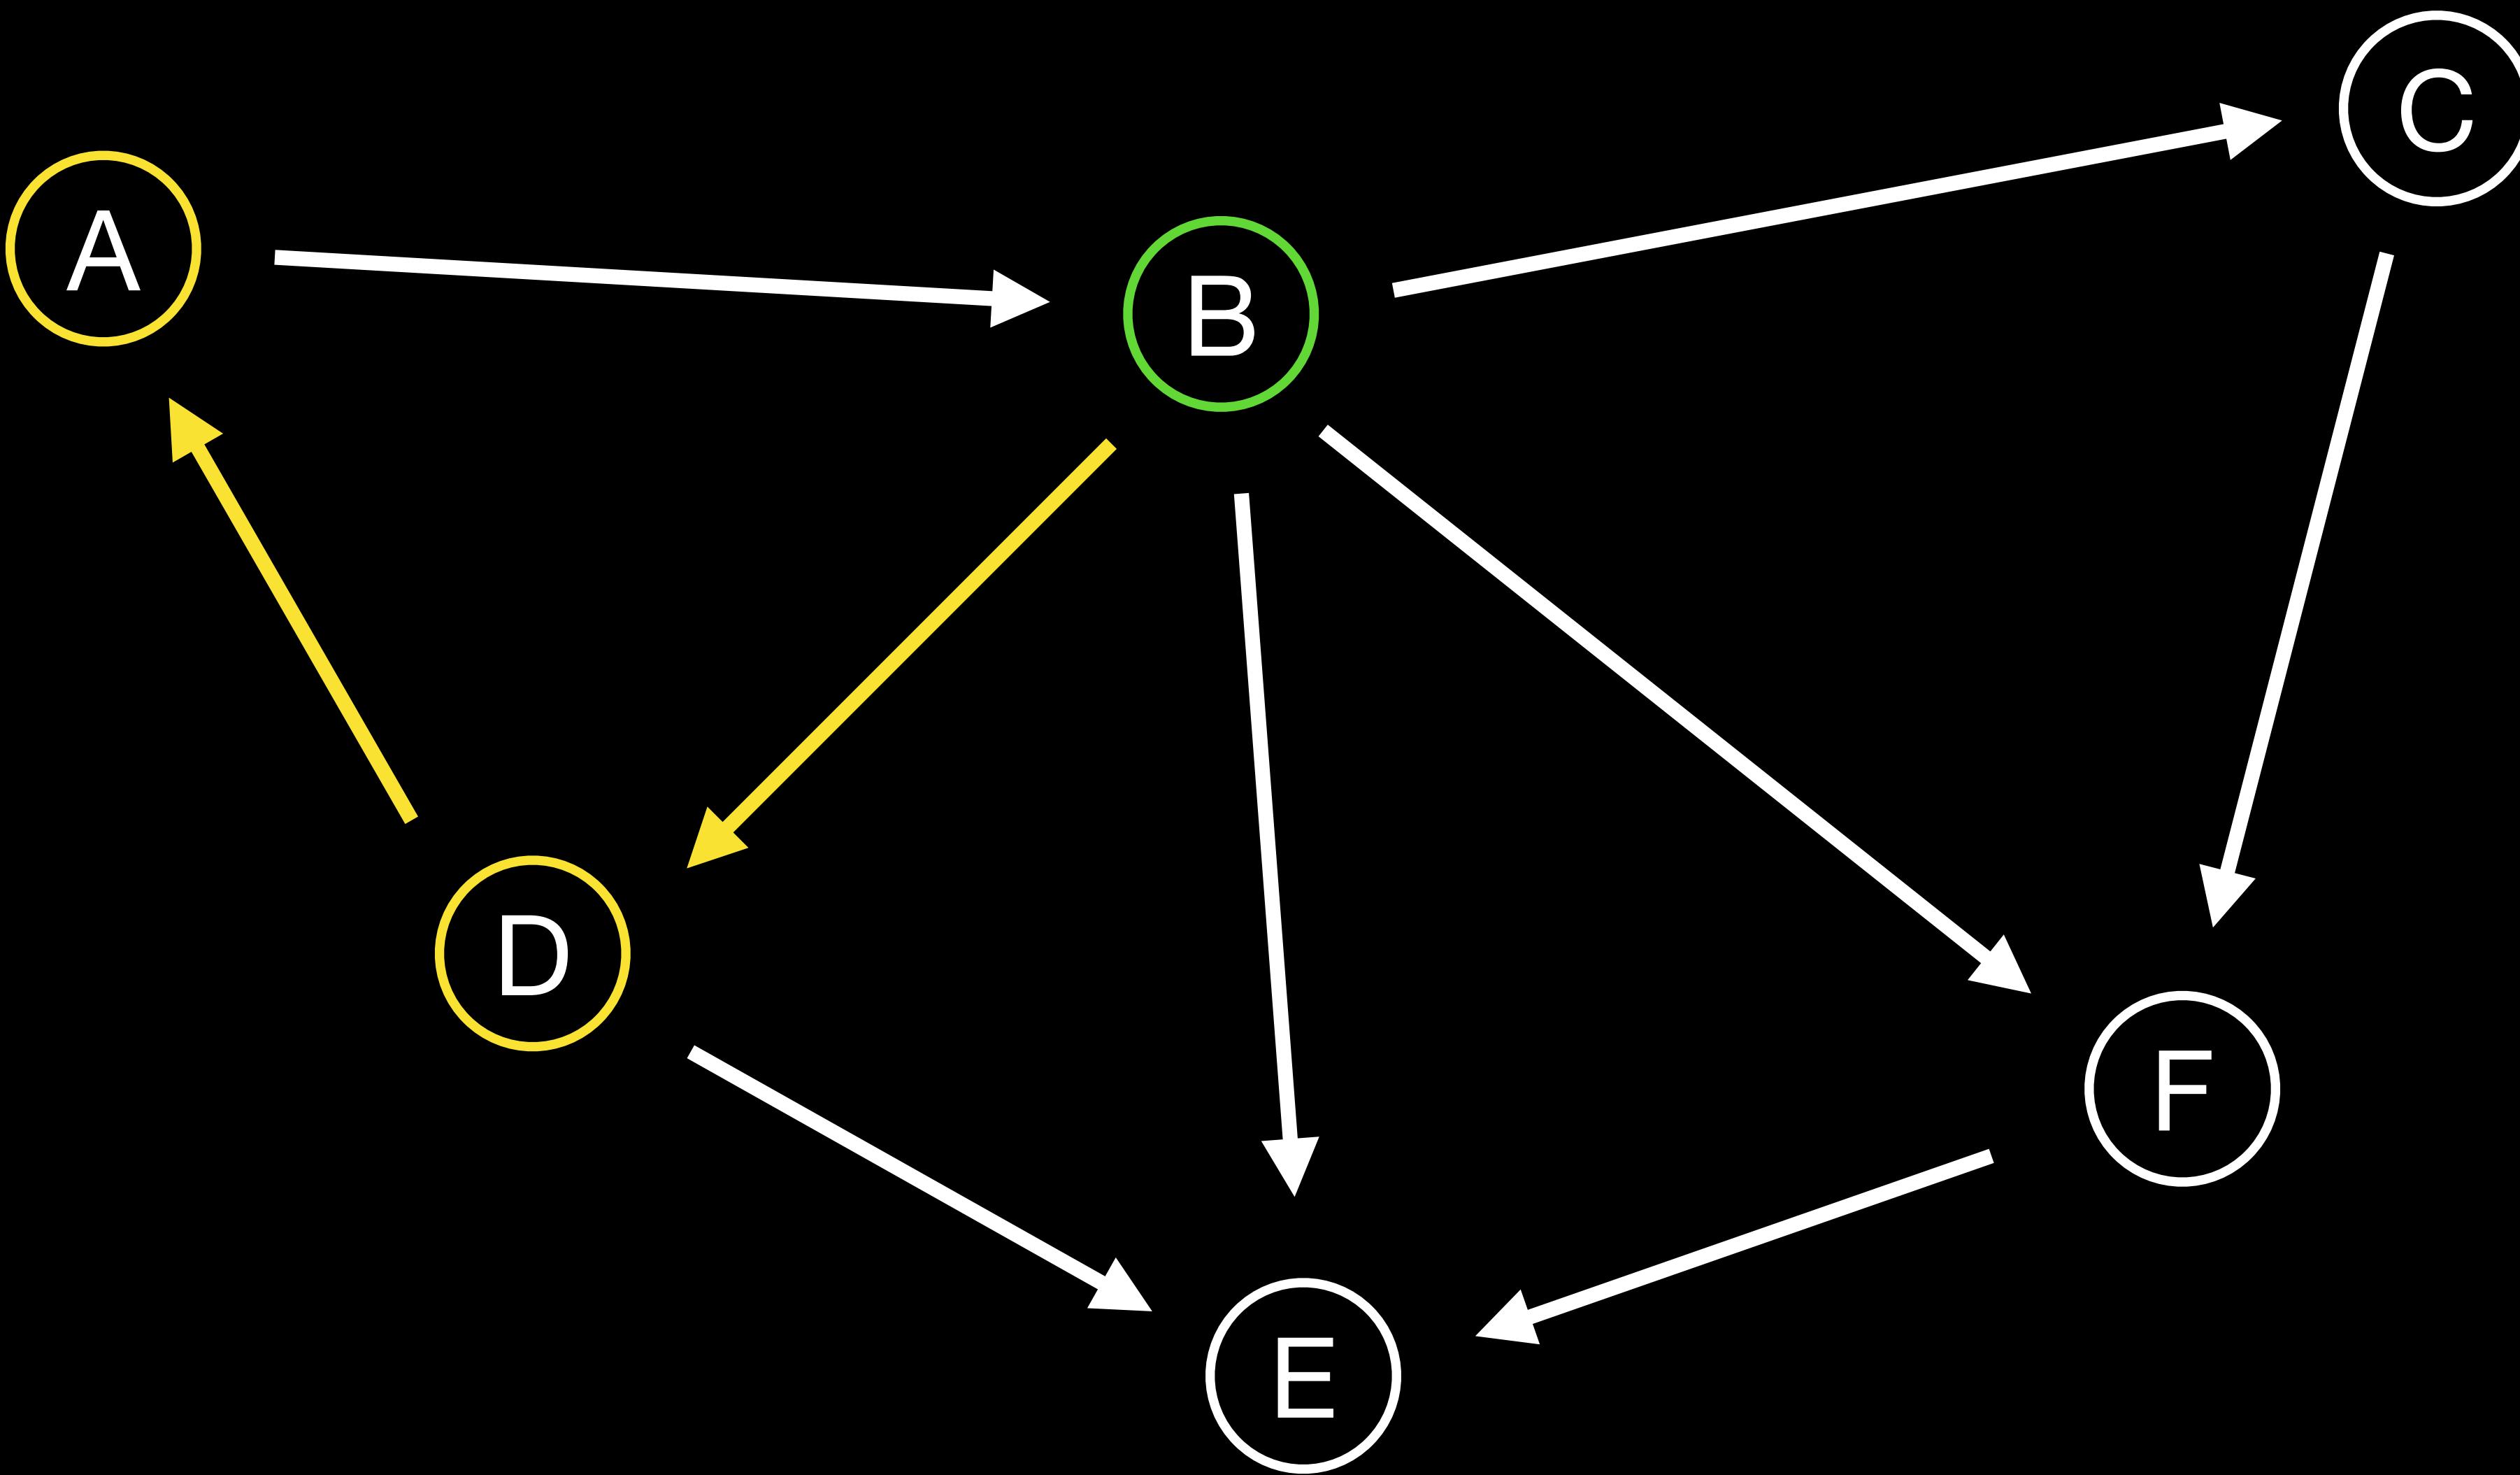
Memory operations?

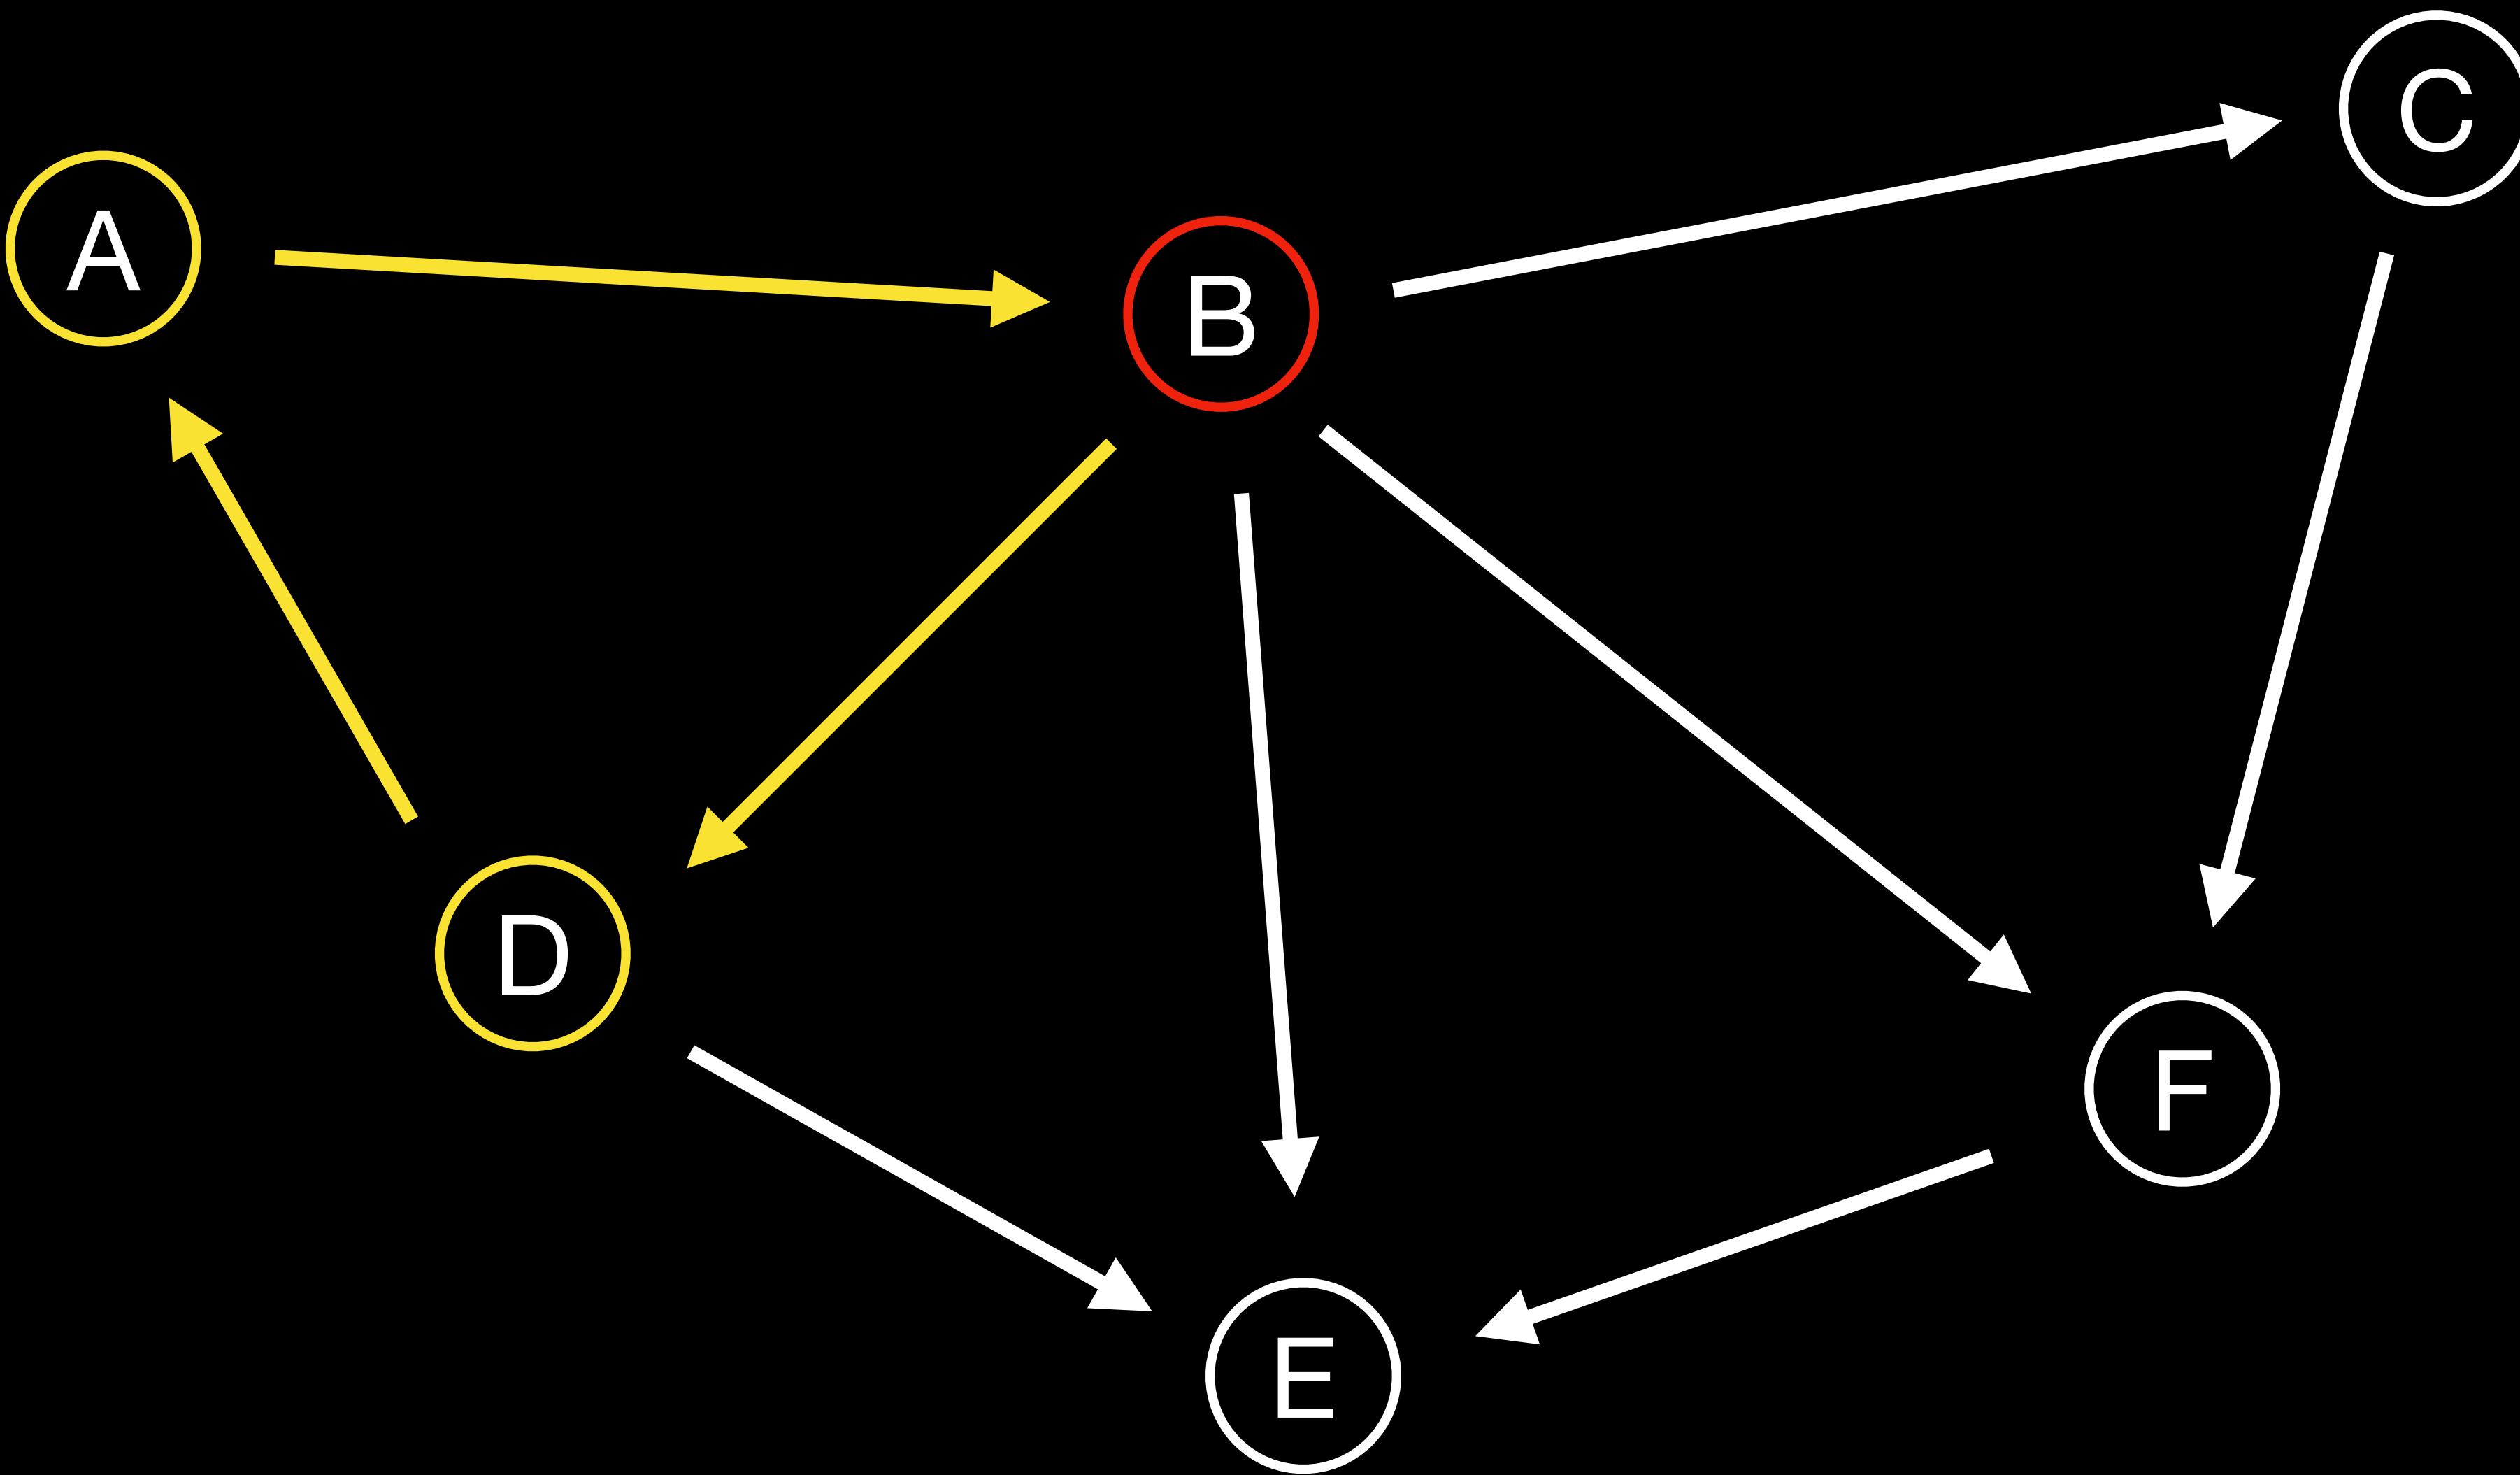
# Hamiltonian path

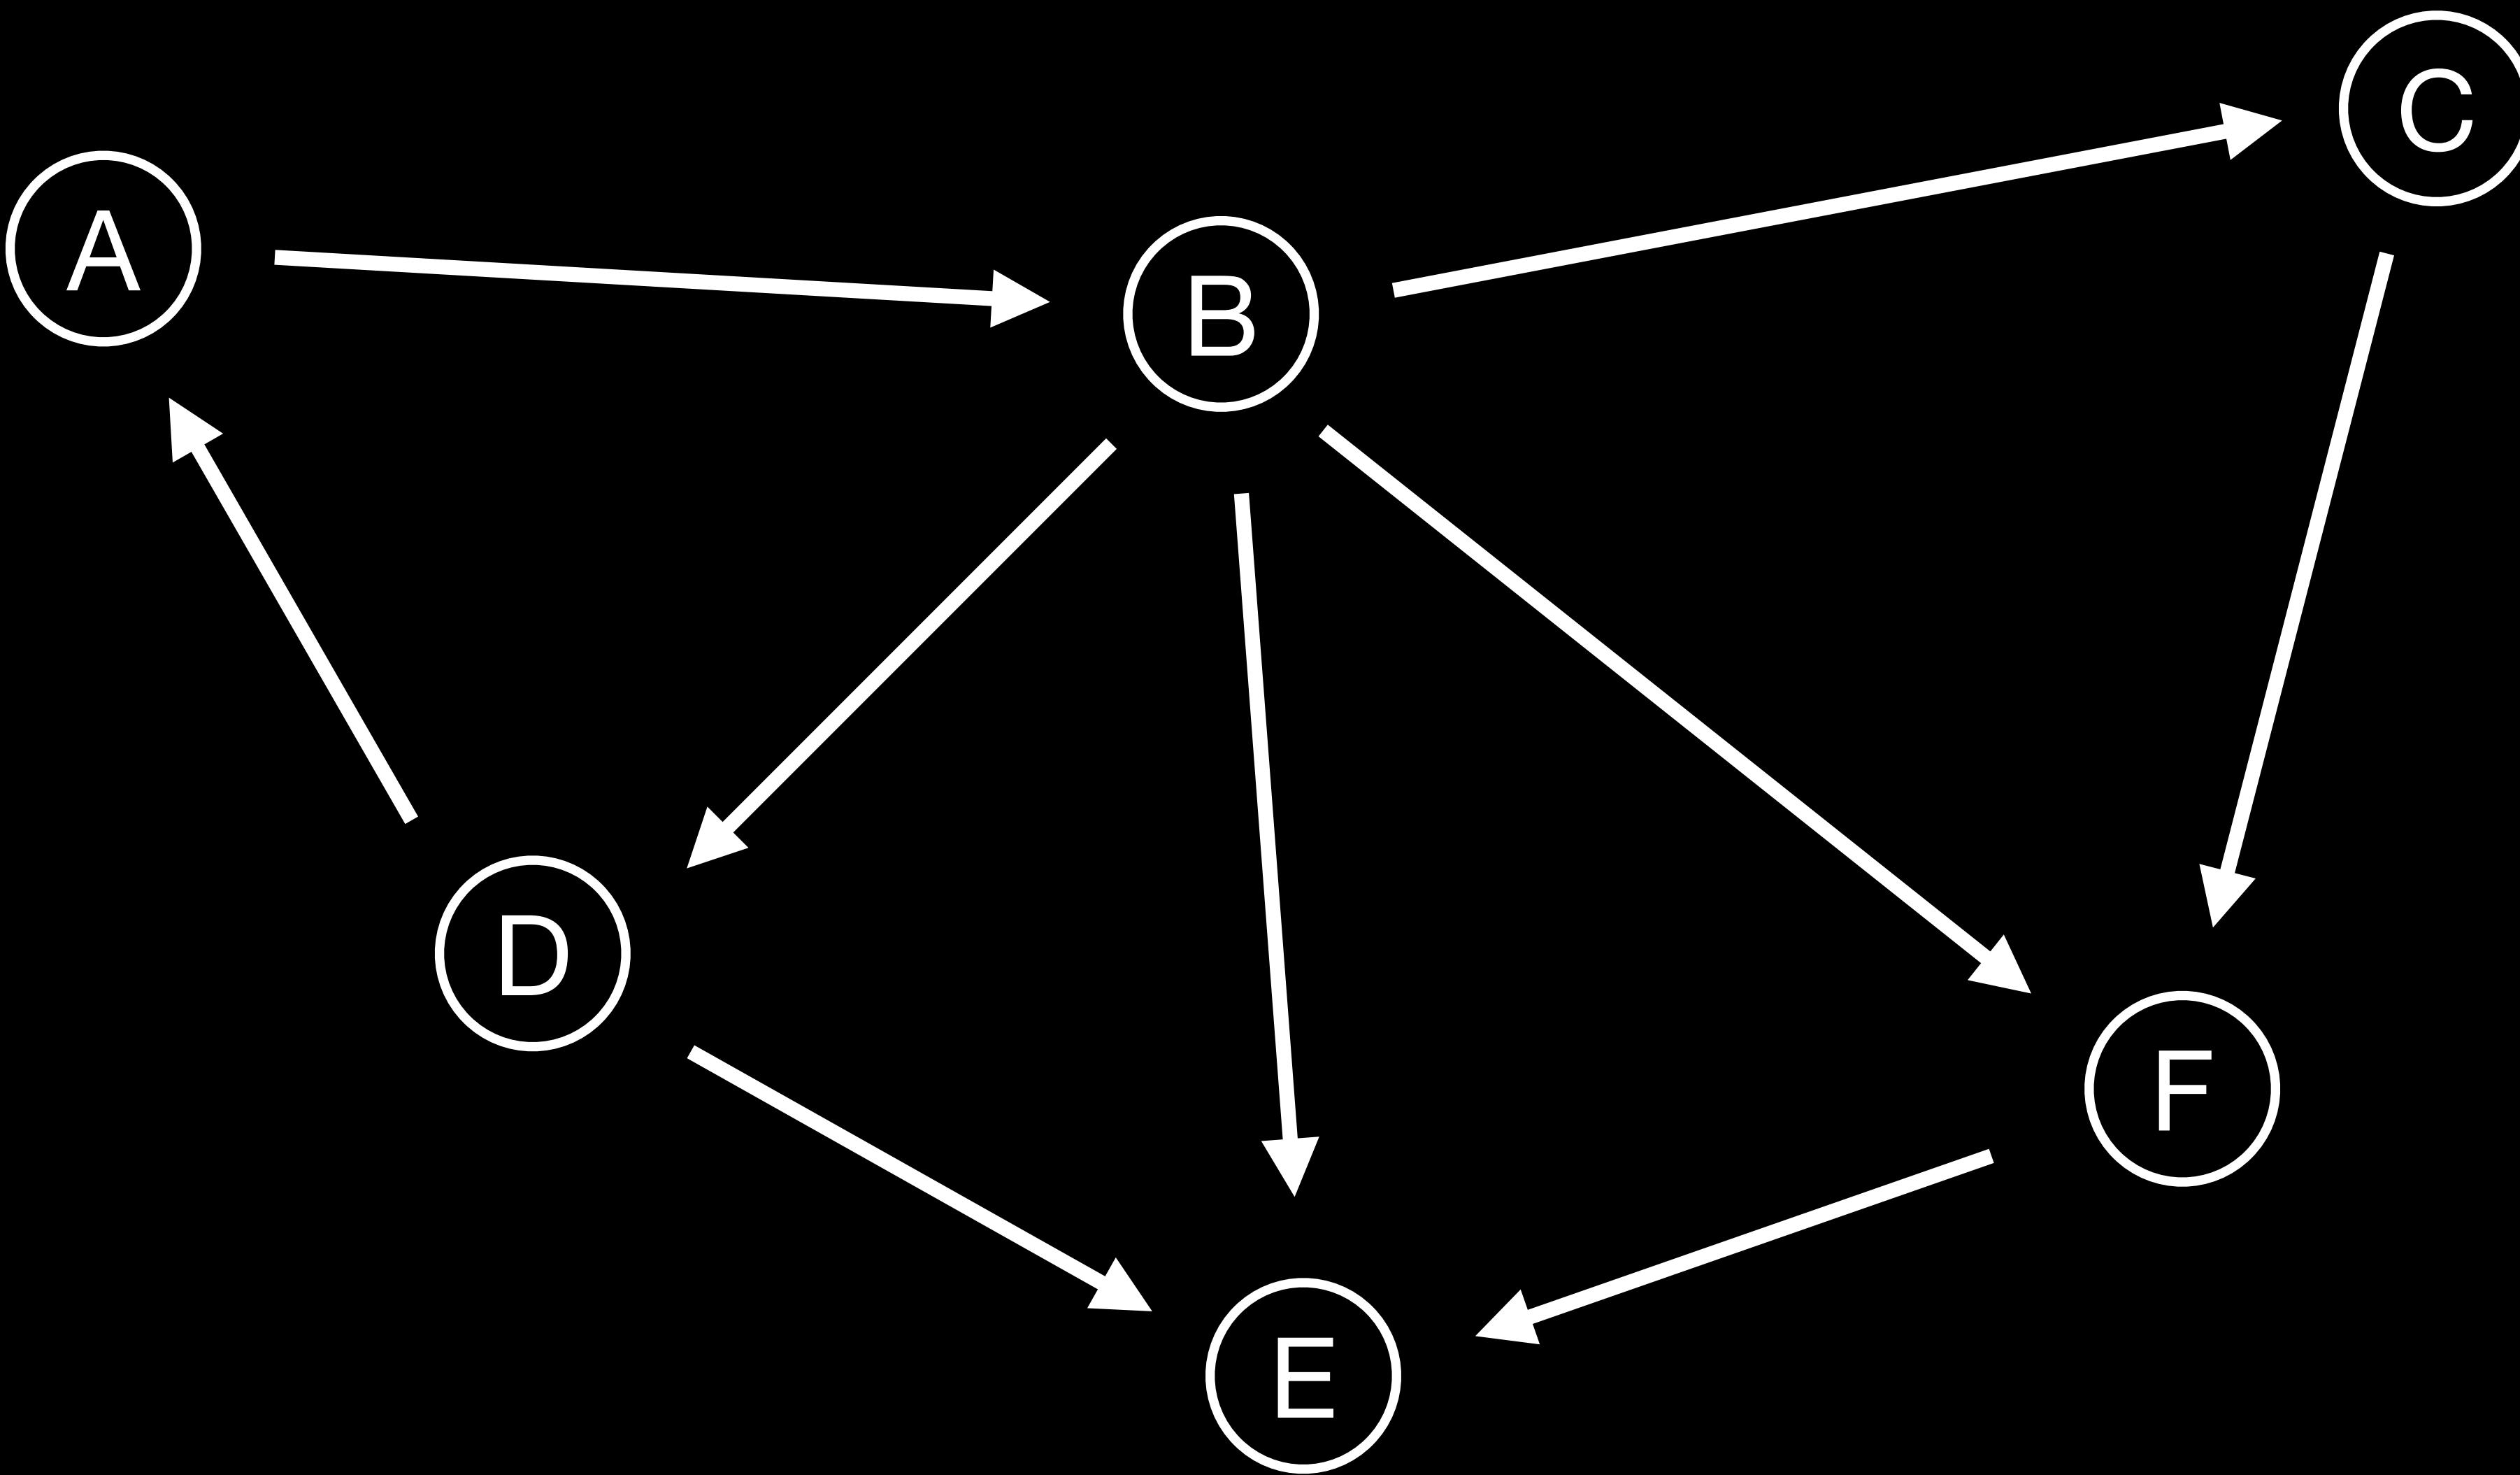


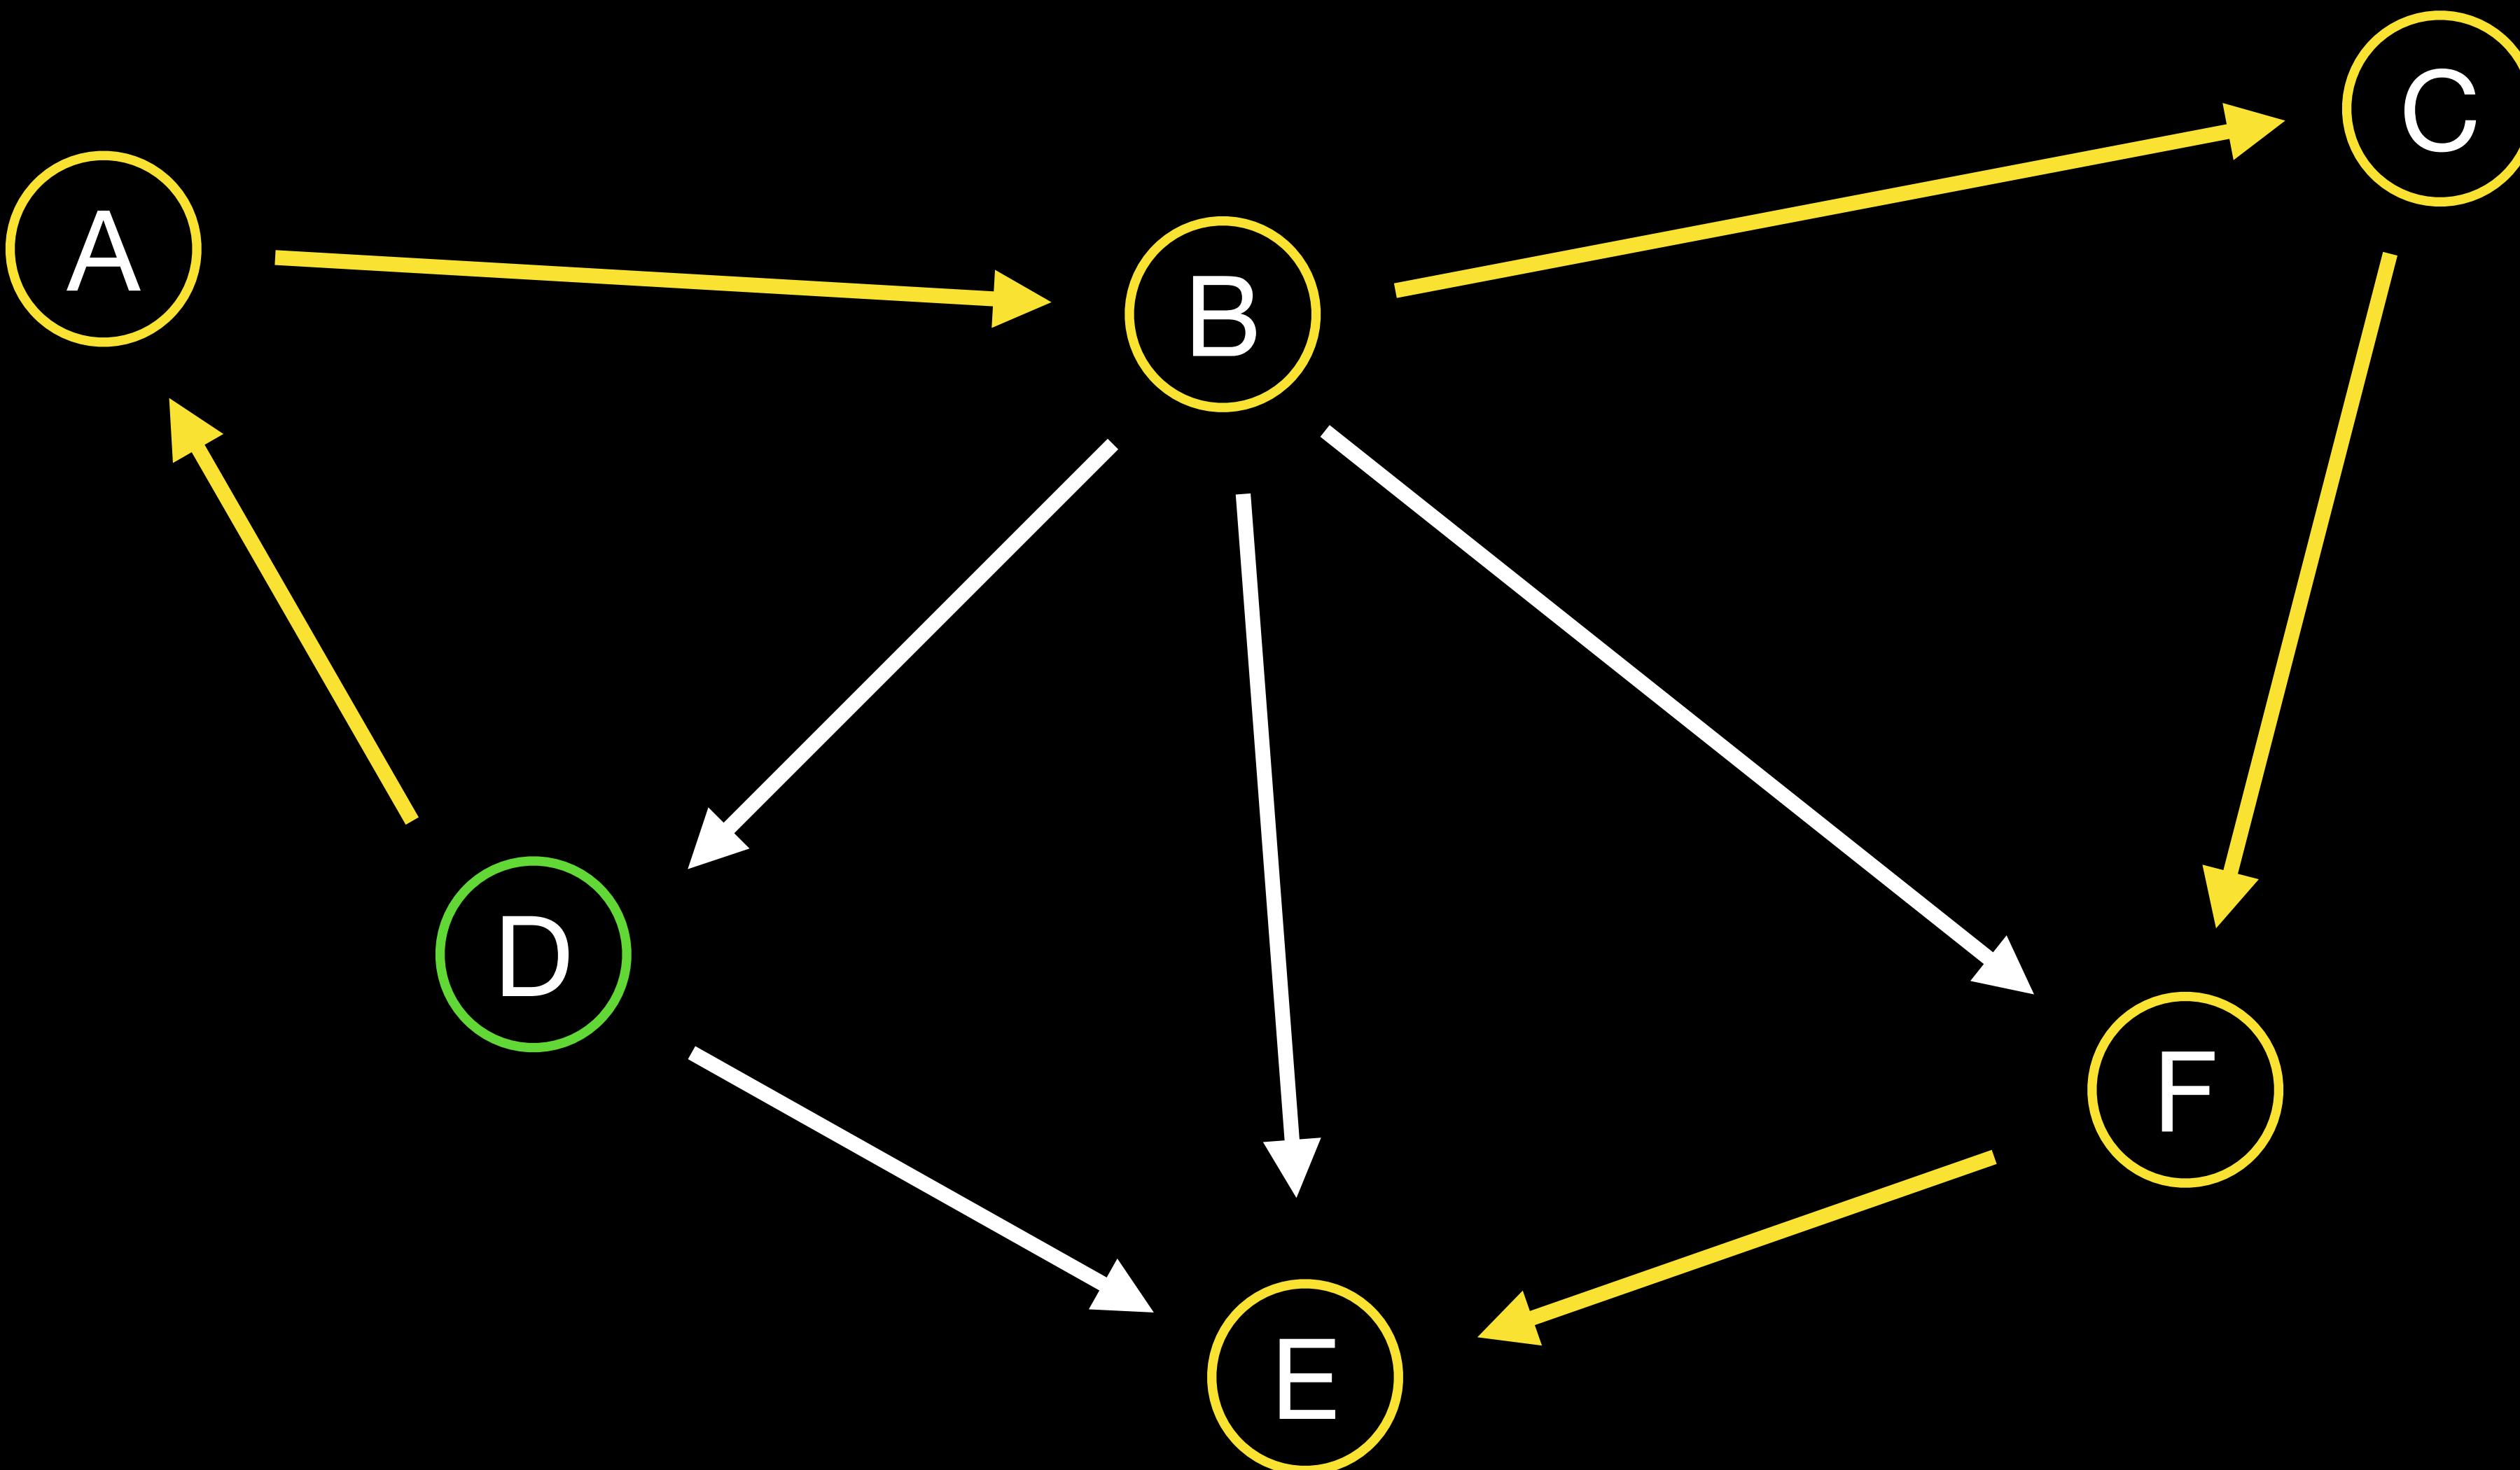


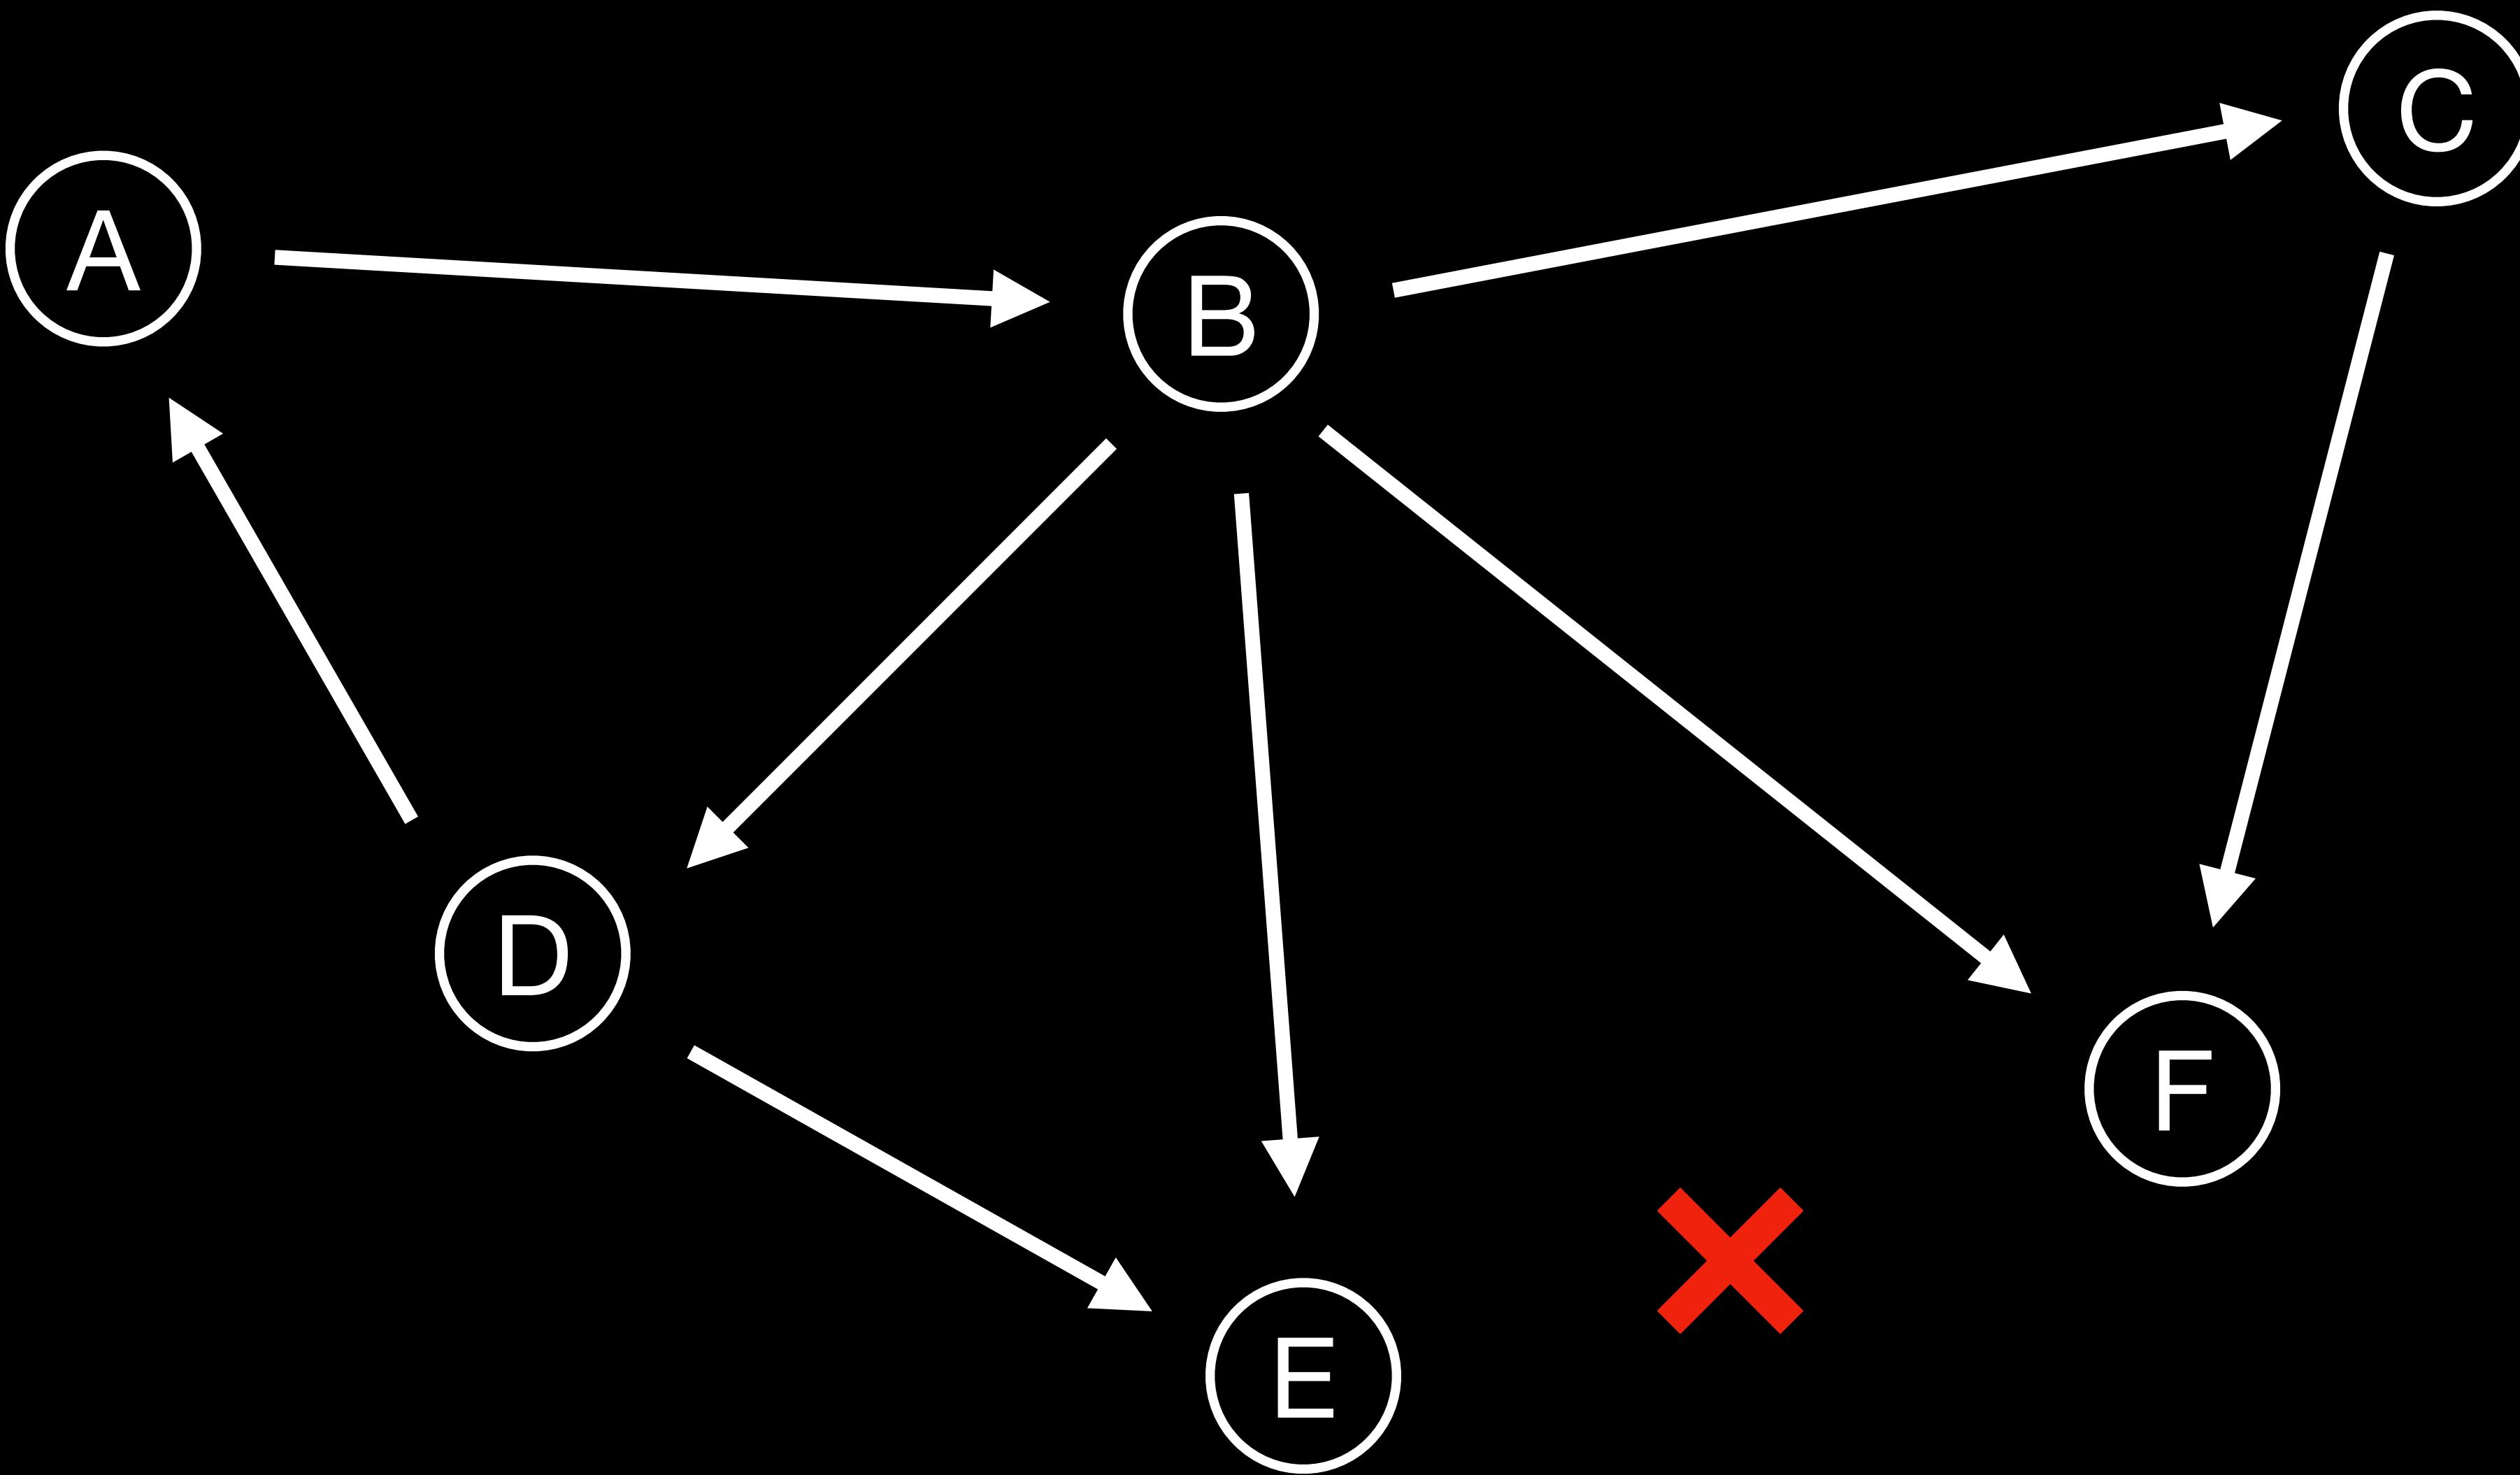












# Assembly

# How to signal if a path exists?

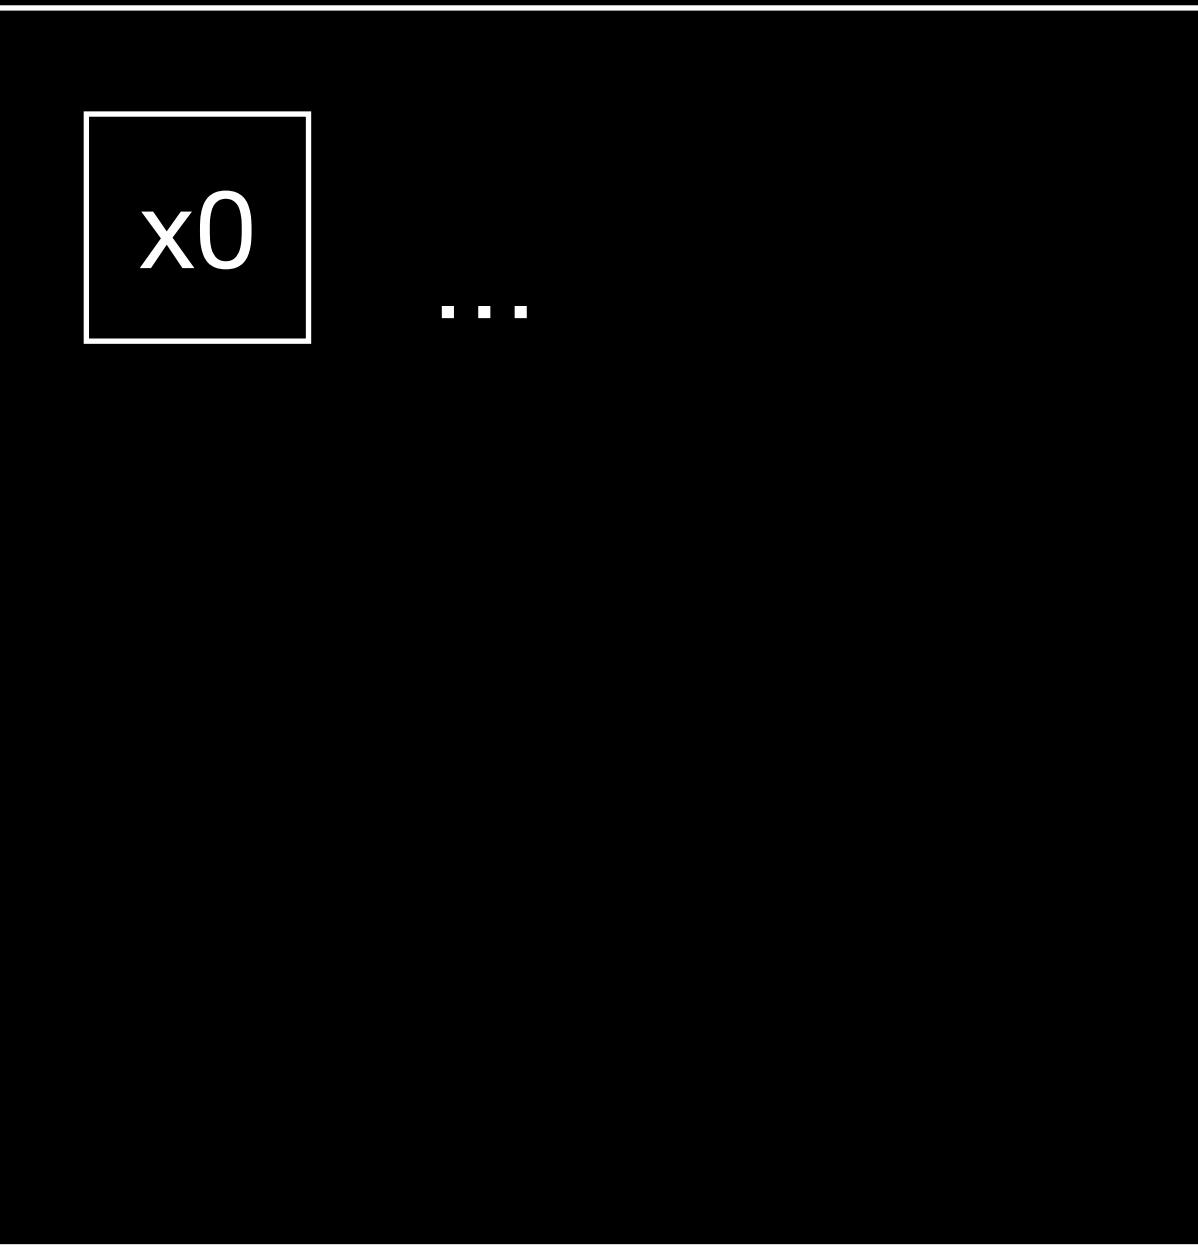
# How to signal if a path exists?

Exit status (0 or 1)

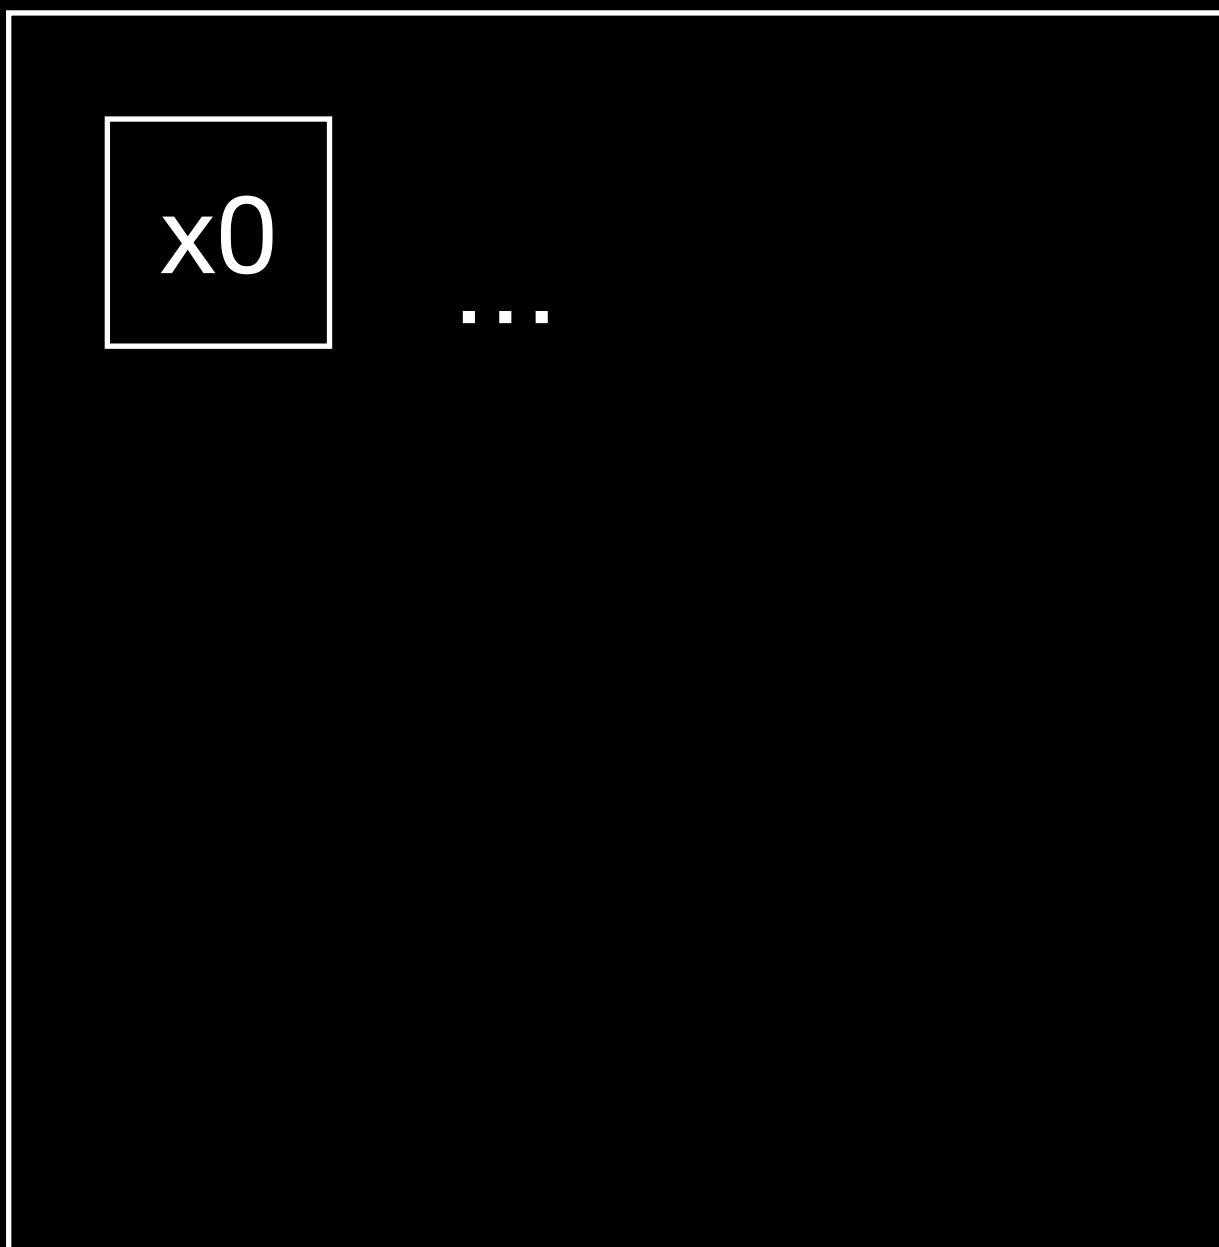
```
mov x16, #1          exit status  
mov x0, #0           ←  
svc 0
```

# How to store the “visited” set?

# How to store the “visited” set?

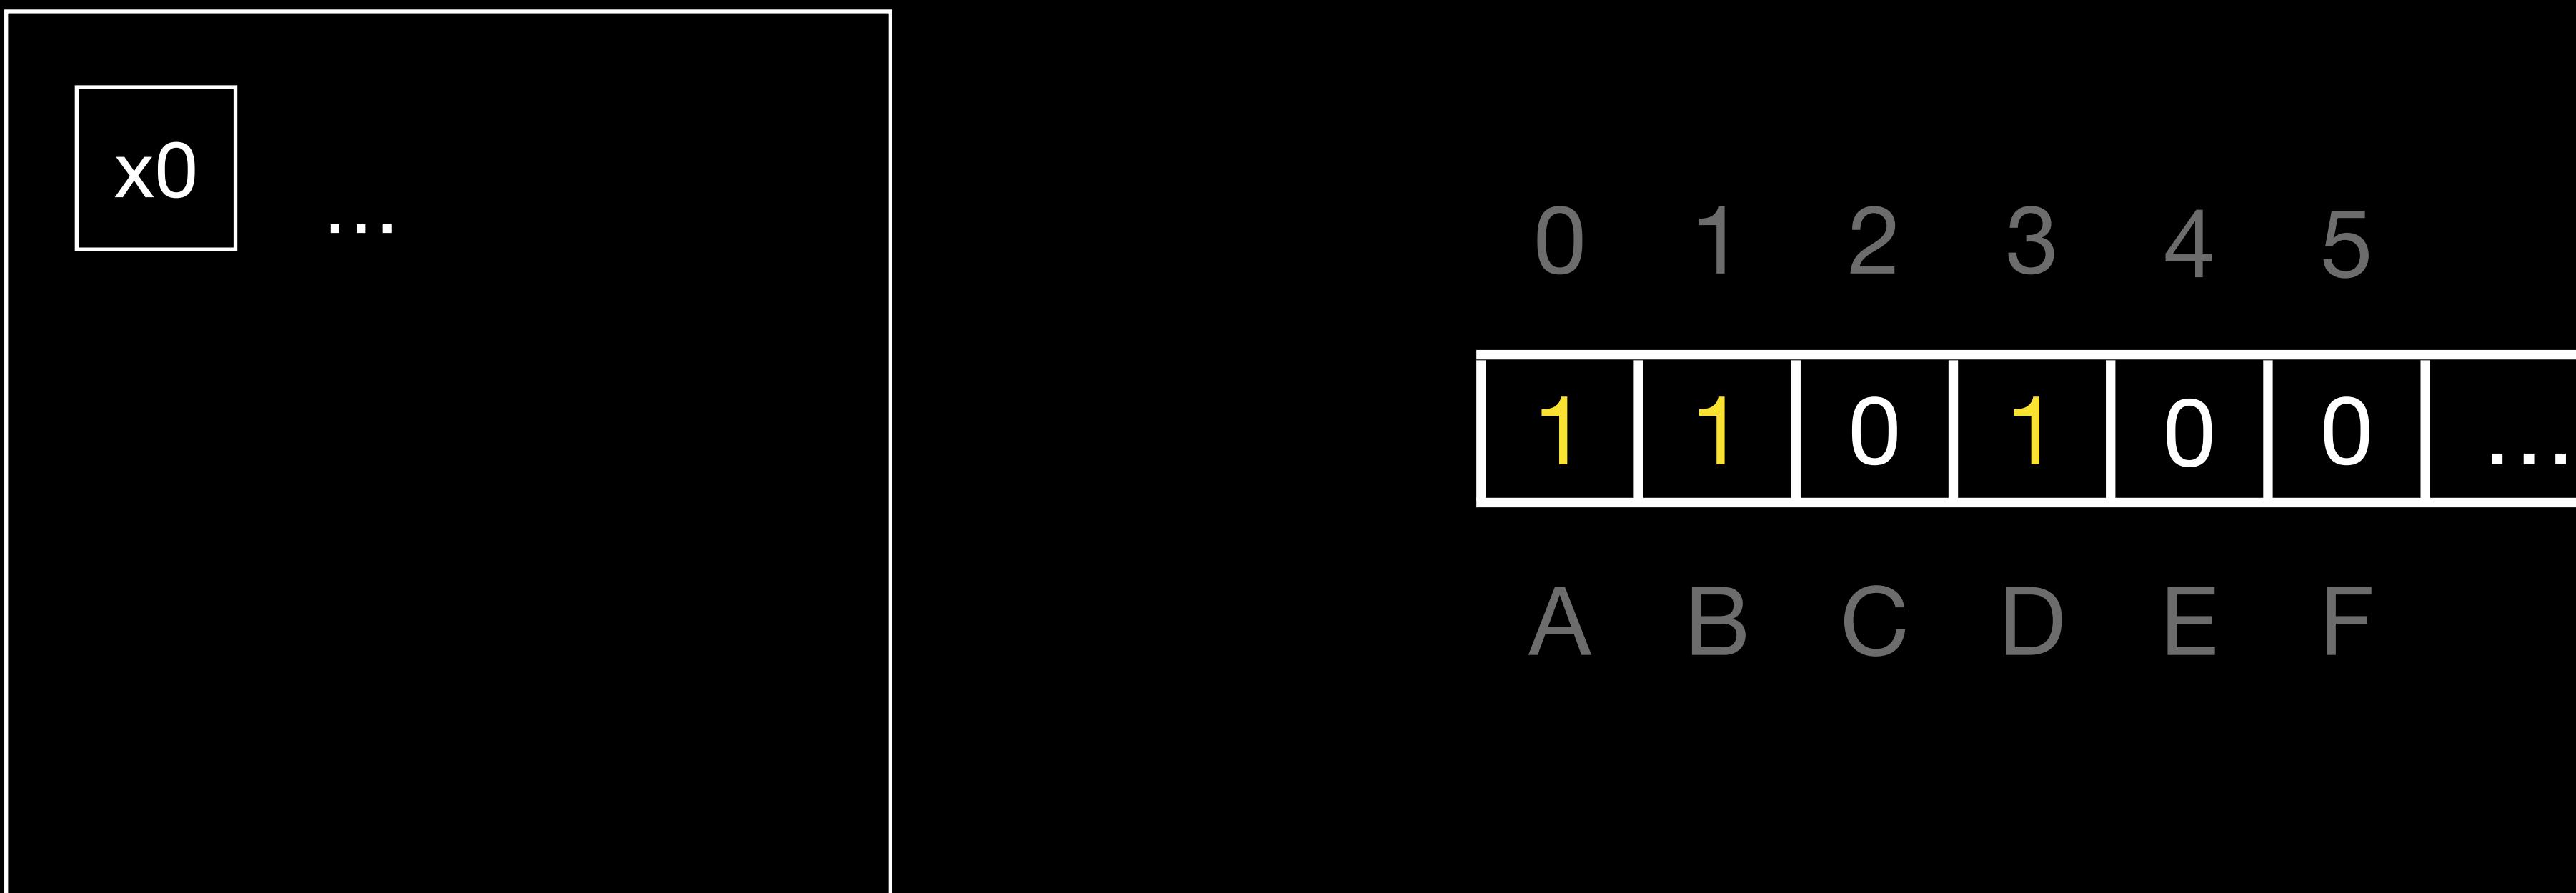


# How to store the “visited” set?



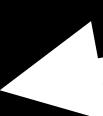
1 cycle

# How to store the “visited” set?



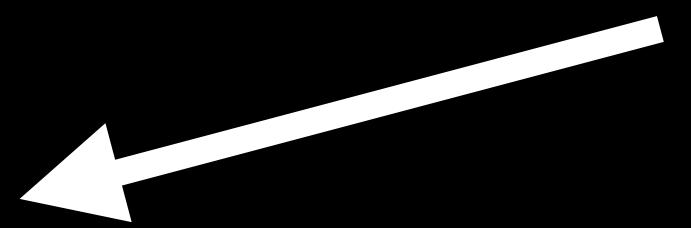
clear at the start

mov x0, #0



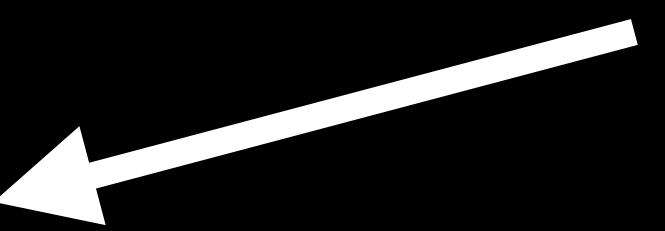
mark D as visited

eor x0, x0, #8

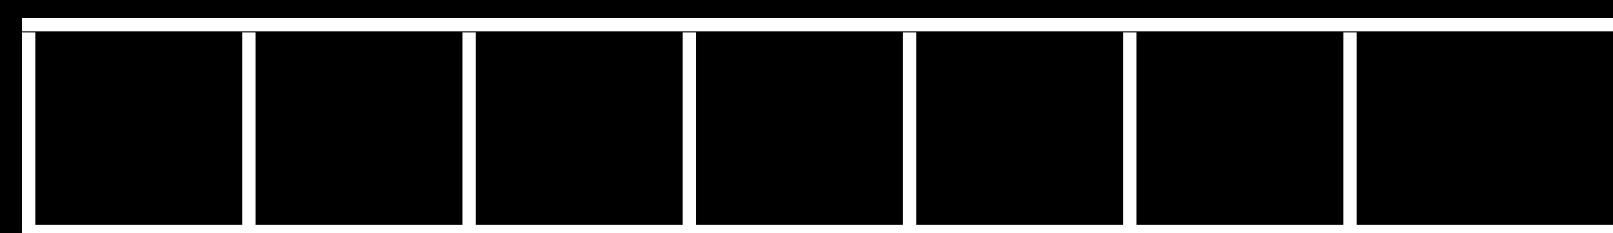


mark D as visited

eor x0, x0, #8



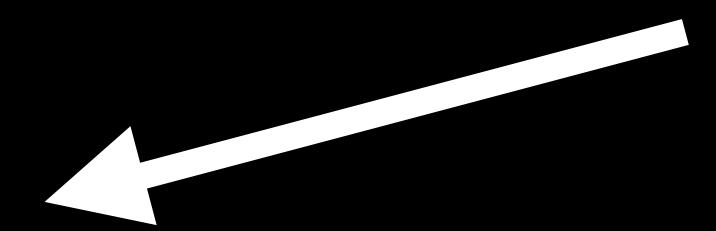
0 1 2 3 4 5



A B C D E F

mark D as unvisited

eor x0, x0, #8



(flips the bit back again)

# How to visit a node?

# How to visit a node?

Branch to a subroutine

```
_main:  
    bl visit_a  
  
visit_a:  
    // Visiting node A
```

```
_main:  
bl visit_a
```

The “visit\_a” subroutine

```
visit_a: ←  
// Visiting node A
```

```
_main:
```

```
bl visit_a
```

“branch link”

```
visit_a:
```

```
// Visiting node A
```

# What is “branch link” ?

Moves the program counter

Sets the x30 register

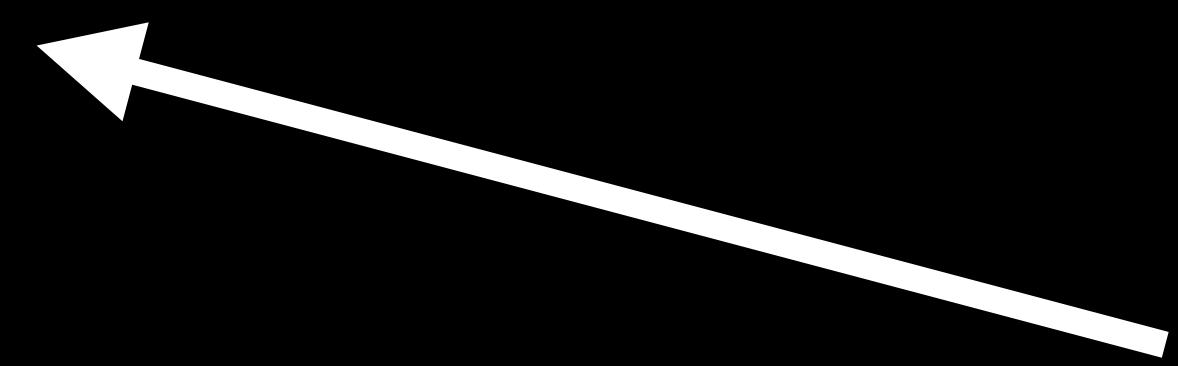
```
_main:
```

```
    bl visit_a
```

```
visit_a:
```

```
    // Visiting node A
```

```
ret
```



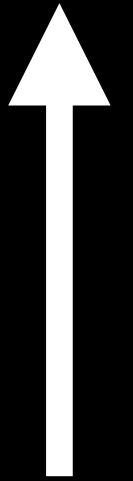
Return to the address in x30

# How to branch conditionally?

Multiple different ways

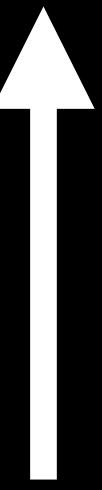
tbz x0, #3, visit\_d

tbz x0, #3, visit\_d



test if a bit is zero

tbz x0, #3, visit\_d



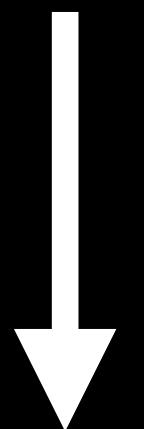
index 3

Problem: doesn't set x30

tbnz x0, #3, after  
bl visit\_d

after:

test if non-zero

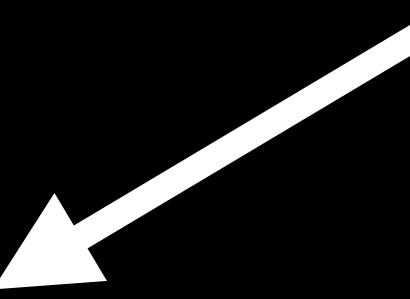


tbnz x0, #3, after  
bl visit\_d

after:

```
tbnz x0, #3, after  
bl visit_d  
after:
```

skip the next line



tbnz x0, #3, after  
bl visit\_d ←  
after: sets x30

# How to check if we found a path?

The register is ‘111111’

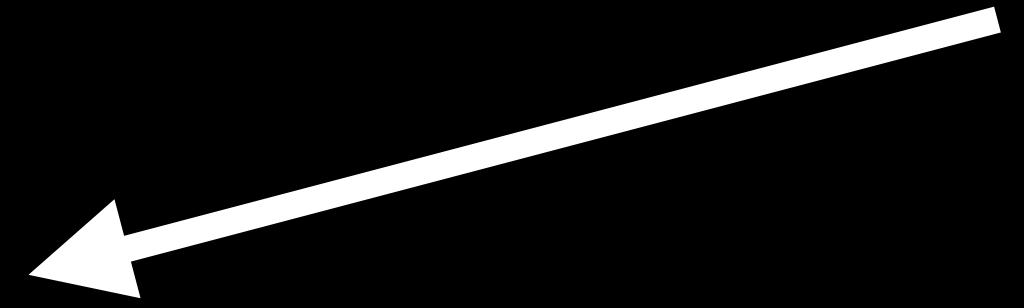
```
cmp x0, #63
b.eq found_a_path
```

Compare with '111111'

```
cmp x0, #63  
b.eq found_a_path
```

```
cmp x0, #63
```

```
b.eq found_a_path
```



Branch if equal

Putting it together

\_main:

    mov x0, #0

    bl visit\_a

    bl visit\_b

    bl visit\_c

    bl visit\_d

    bl visit\_e

    bl visit\_f

no\_path\_found:

    mov x16, #1

    mov x0, #1

    svc 0

found\_a\_path:

    mov x16, #1

    mov x0, #0

    svc 0

\_main:

```
mov x0, #0
```



We haven't visited any nodes yet

```
bl visit_a  
bl visit_b  
bl visit_c  
bl visit_d  
bl visit_e  
bl visit_f
```

no\_path\_found:

```
mov x16, #1  
mov x0, #1  
svc 0
```

found\_a\_path:

```
mov x16, #1  
mov x0, #0  
svc 0
```

```
_main:
```

```
    mov x0, #0
```

```
    bl visit_a
```

```
    bl visit_b
```

```
    bl visit_c
```

```
    bl visit_d
```

```
    bl visit_e
```

```
    bl visit_f
```



Try from every possible node

(do a depth-first search)

```
no_path_found:
```

```
    mov x16, #1
```

```
    mov x0, #1
```

```
    svc 0
```

```
found_a_path:
```

```
    mov x16, #1
```

```
    mov x0, #0
```

```
    svc 0
```

\_main:

```
    mov x0, #0
```

```
    bl visit_a
```

```
    bl visit_b
```

```
    bl visit_c
```

```
    bl visit_d
```

```
    bl visit_e
```

```
    bl visit_f
```

no\_path\_found:



Exit 1 if no path was found

```
    mov x16, #1
```

```
    mov x0, #1
```

```
    svc 0
```

found\_a\_path:

```
    mov x16, #1
```

```
    mov x0, #0
```

```
    svc 0
```

```
_main:
```

```
    mov x0, #0
```

```
    bl visit_a
```

```
    bl visit_b
```

```
    bl visit_c
```

```
    bl visit_d
```

```
    bl visit_e
```

```
    bl visit_f
```

```
no_path_found:
```

```
    mov x16, #1
```

```
    mov x0, #1
```

```
    svc 0
```

```
found_a_path:
```



Branch here if we find a path

```
    mov x16, #1
```

```
    mov x0, #0
```

```
    svc 0
```

visit\_a

visit\_a:

    eor x0, x0, #1

    tbnz x0, #1, after\_ab

    bl visit\_b

after\_ab:

    cmp x0, #63

    b.eq found\_a\_path

    eor x0, x0, #1

ret

```
visit_a:  
    eor x0, x0, #1 ← Mark A as visited: 20
```

```
tbnz x0, #1, after_ab  
bl visit_b  
after_ab:
```

```
cmp x0, #63  
b.eq found_a_path
```

```
eor x0, x0, #1  
ret
```

```
visit_a:
```

```
    eor x0, x0, #1
```

```
    tbnz x0, #1, after_ab
```



Skip b if already visited  
(index 1)

```
    bl visit_b
```

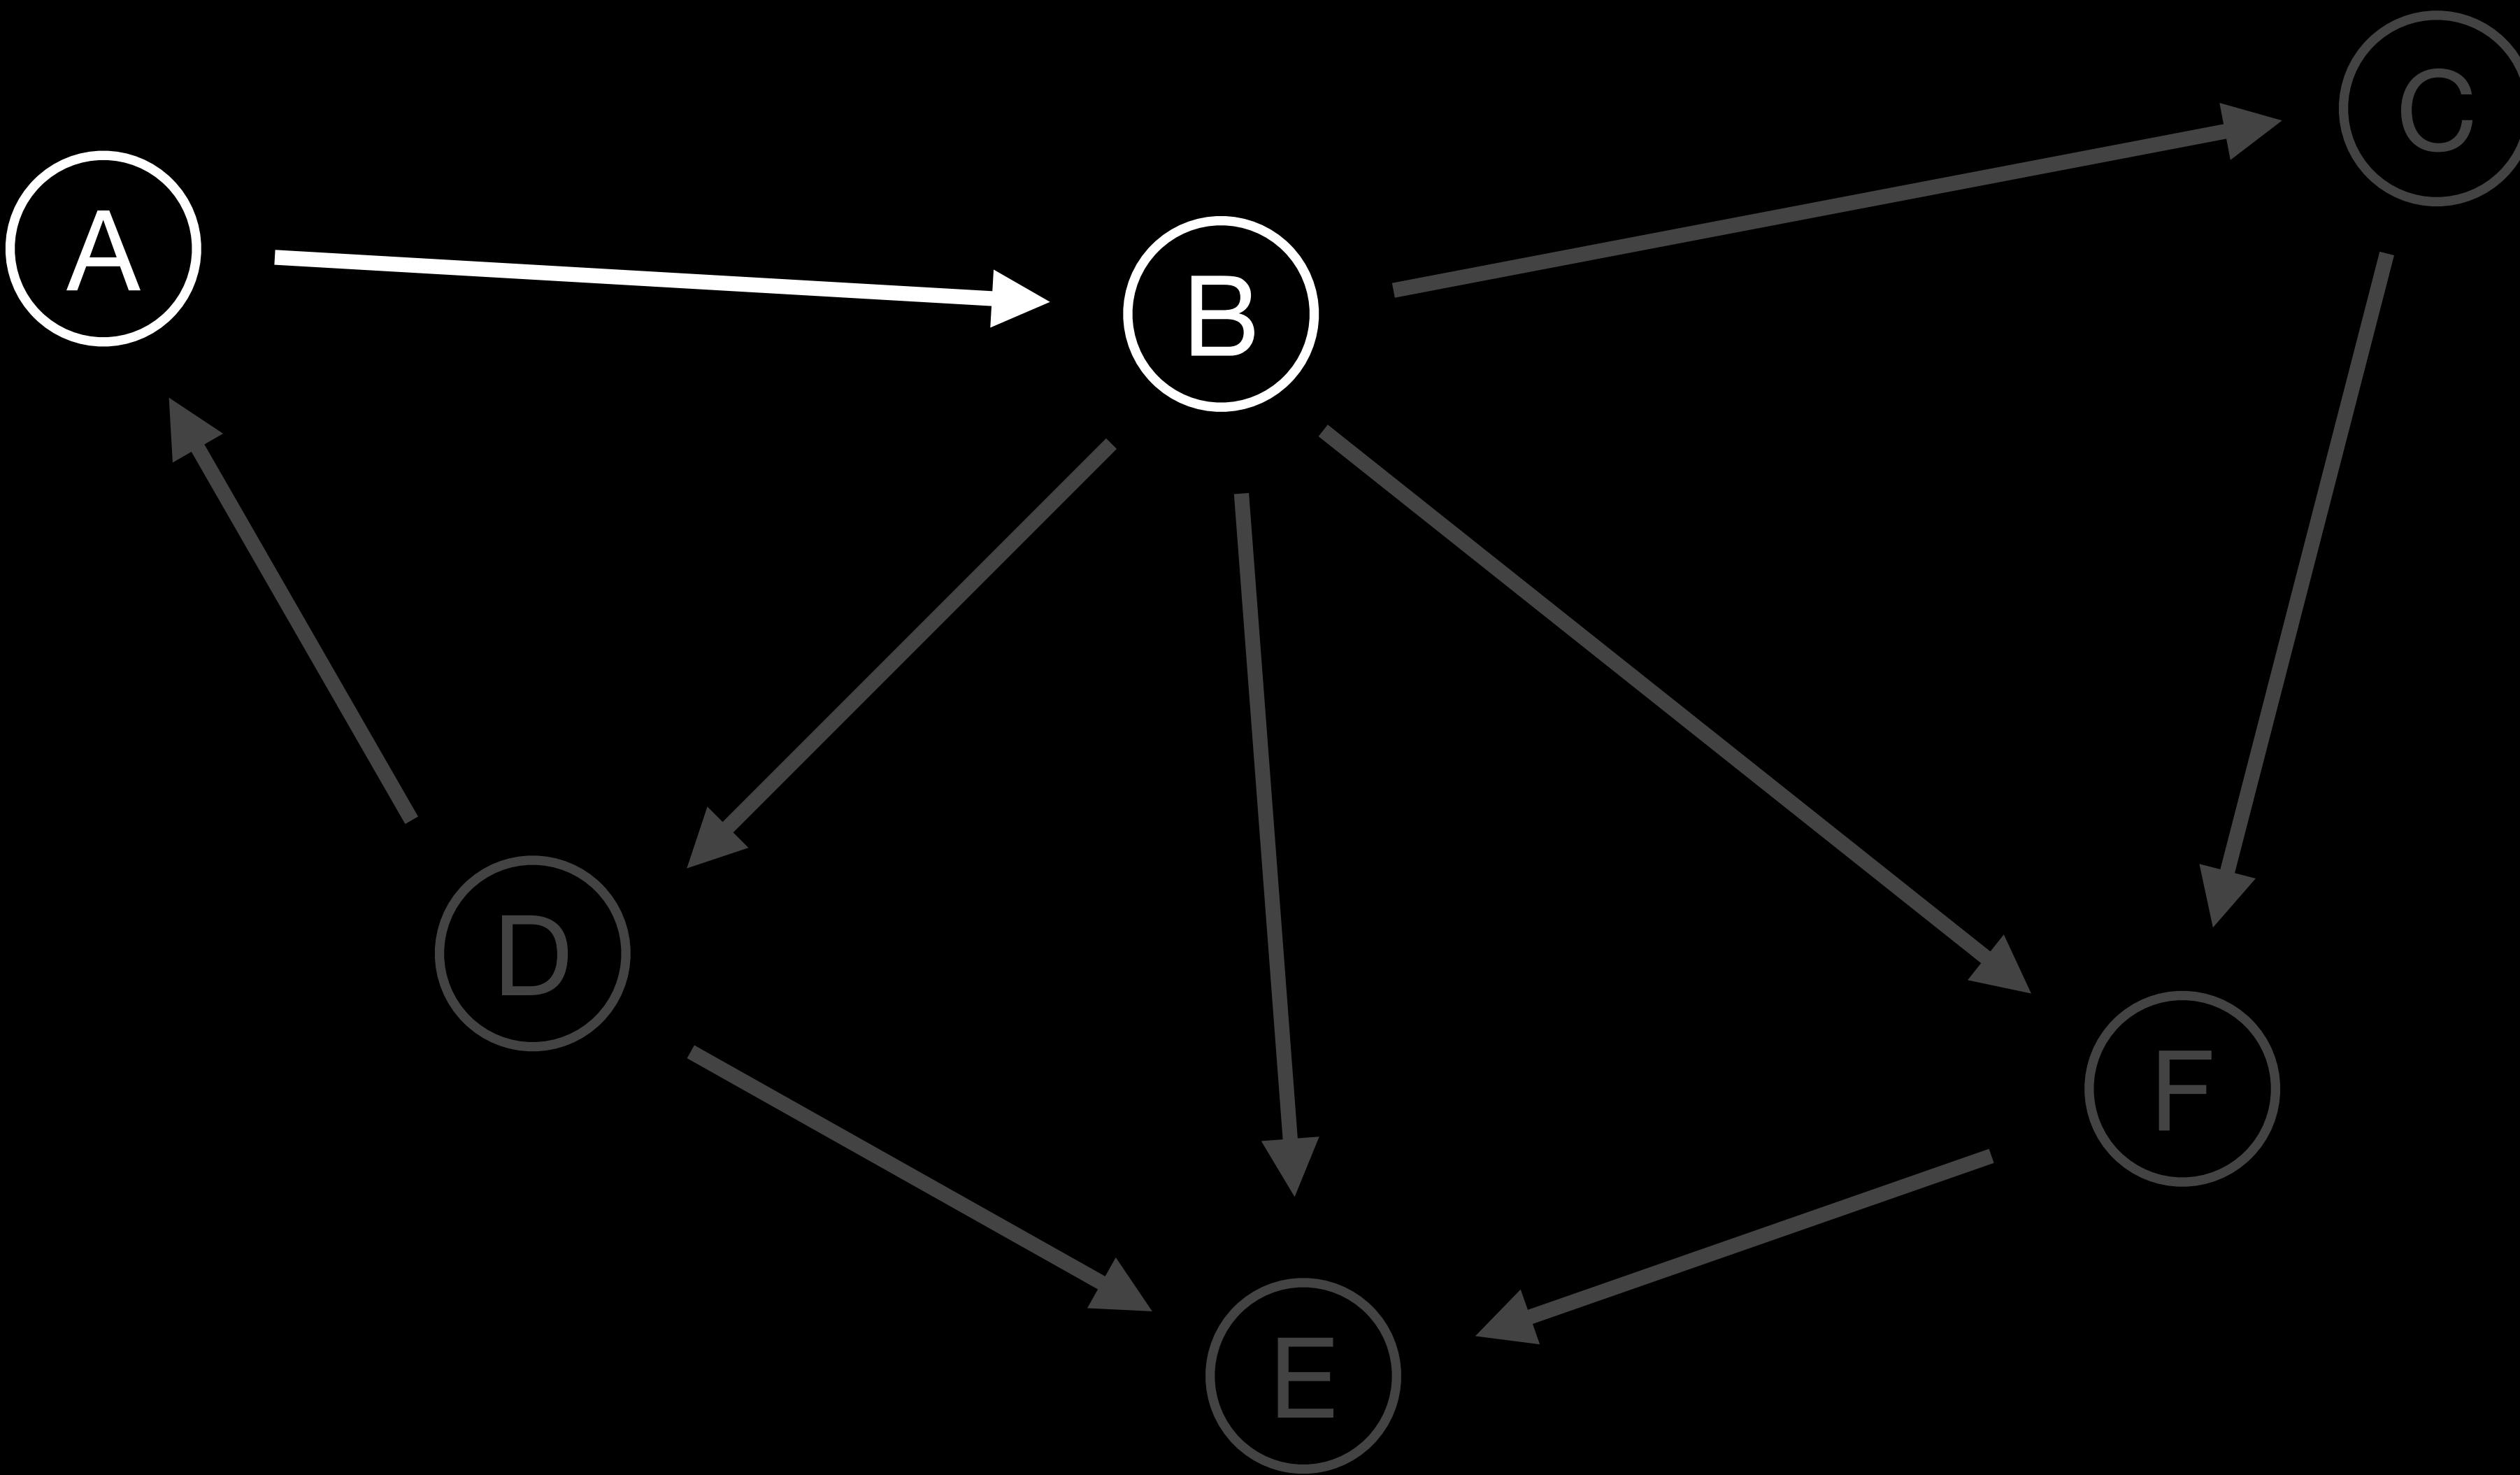
```
after_ab:
```

```
    cmp x0, #63
```

```
    b.eq found_a_path
```

```
    eor x0, x0, #1
```

```
    ret
```



```
visit_a:
```

```
    eor x0, x0, #1
```

```
    tbnz x0, #1, after_ab
```

```
    bl visit_b
```

```
after_ab:
```

```
    cmp x0, #63
```

```
    b.eq found_a_path
```

```
    eor x0, x0, #1
```

```
    ret
```



Check if all nodes visited

```
visit_a:
```

```
    eor x0, x0, #1
```

```
    tbnz x0, #1, after_ab
```

```
    bl visit_b
```

```
after_ab:
```

```
    cmp x0, #63
```

```
    b.eq found_a_path
```

```
eor x0, x0, #1
```



Mark A as **un**visited:  $2^0$

```
ret
```

Repeat for other nodes

visit\_b:

    eor x0, x0, #2

    tbnz x0, #2, after\_bc

    bl visit\_c

after\_bc:

    tbnz x0, #3, after\_bd

    bl visit\_d

after\_bd:

// ...

```
visit_b:
```

```
    eor x0, x0, #2 ← Mark B as visited: 21
```

```
    tbnz x0, #2, after_bc
```

```
    bl visit_c
```

```
after_bc:
```

```
    tbnz x0, #3, after_bd
```

```
    bl visit_d
```

```
after_bd:
```

```
// ...
```

```
visit_b:
```

```
    eor x0, x0, #2
```

```
    tbnz x0, #2, after_bc
```

```
    bl visit_c
```

```
after_bc:
```

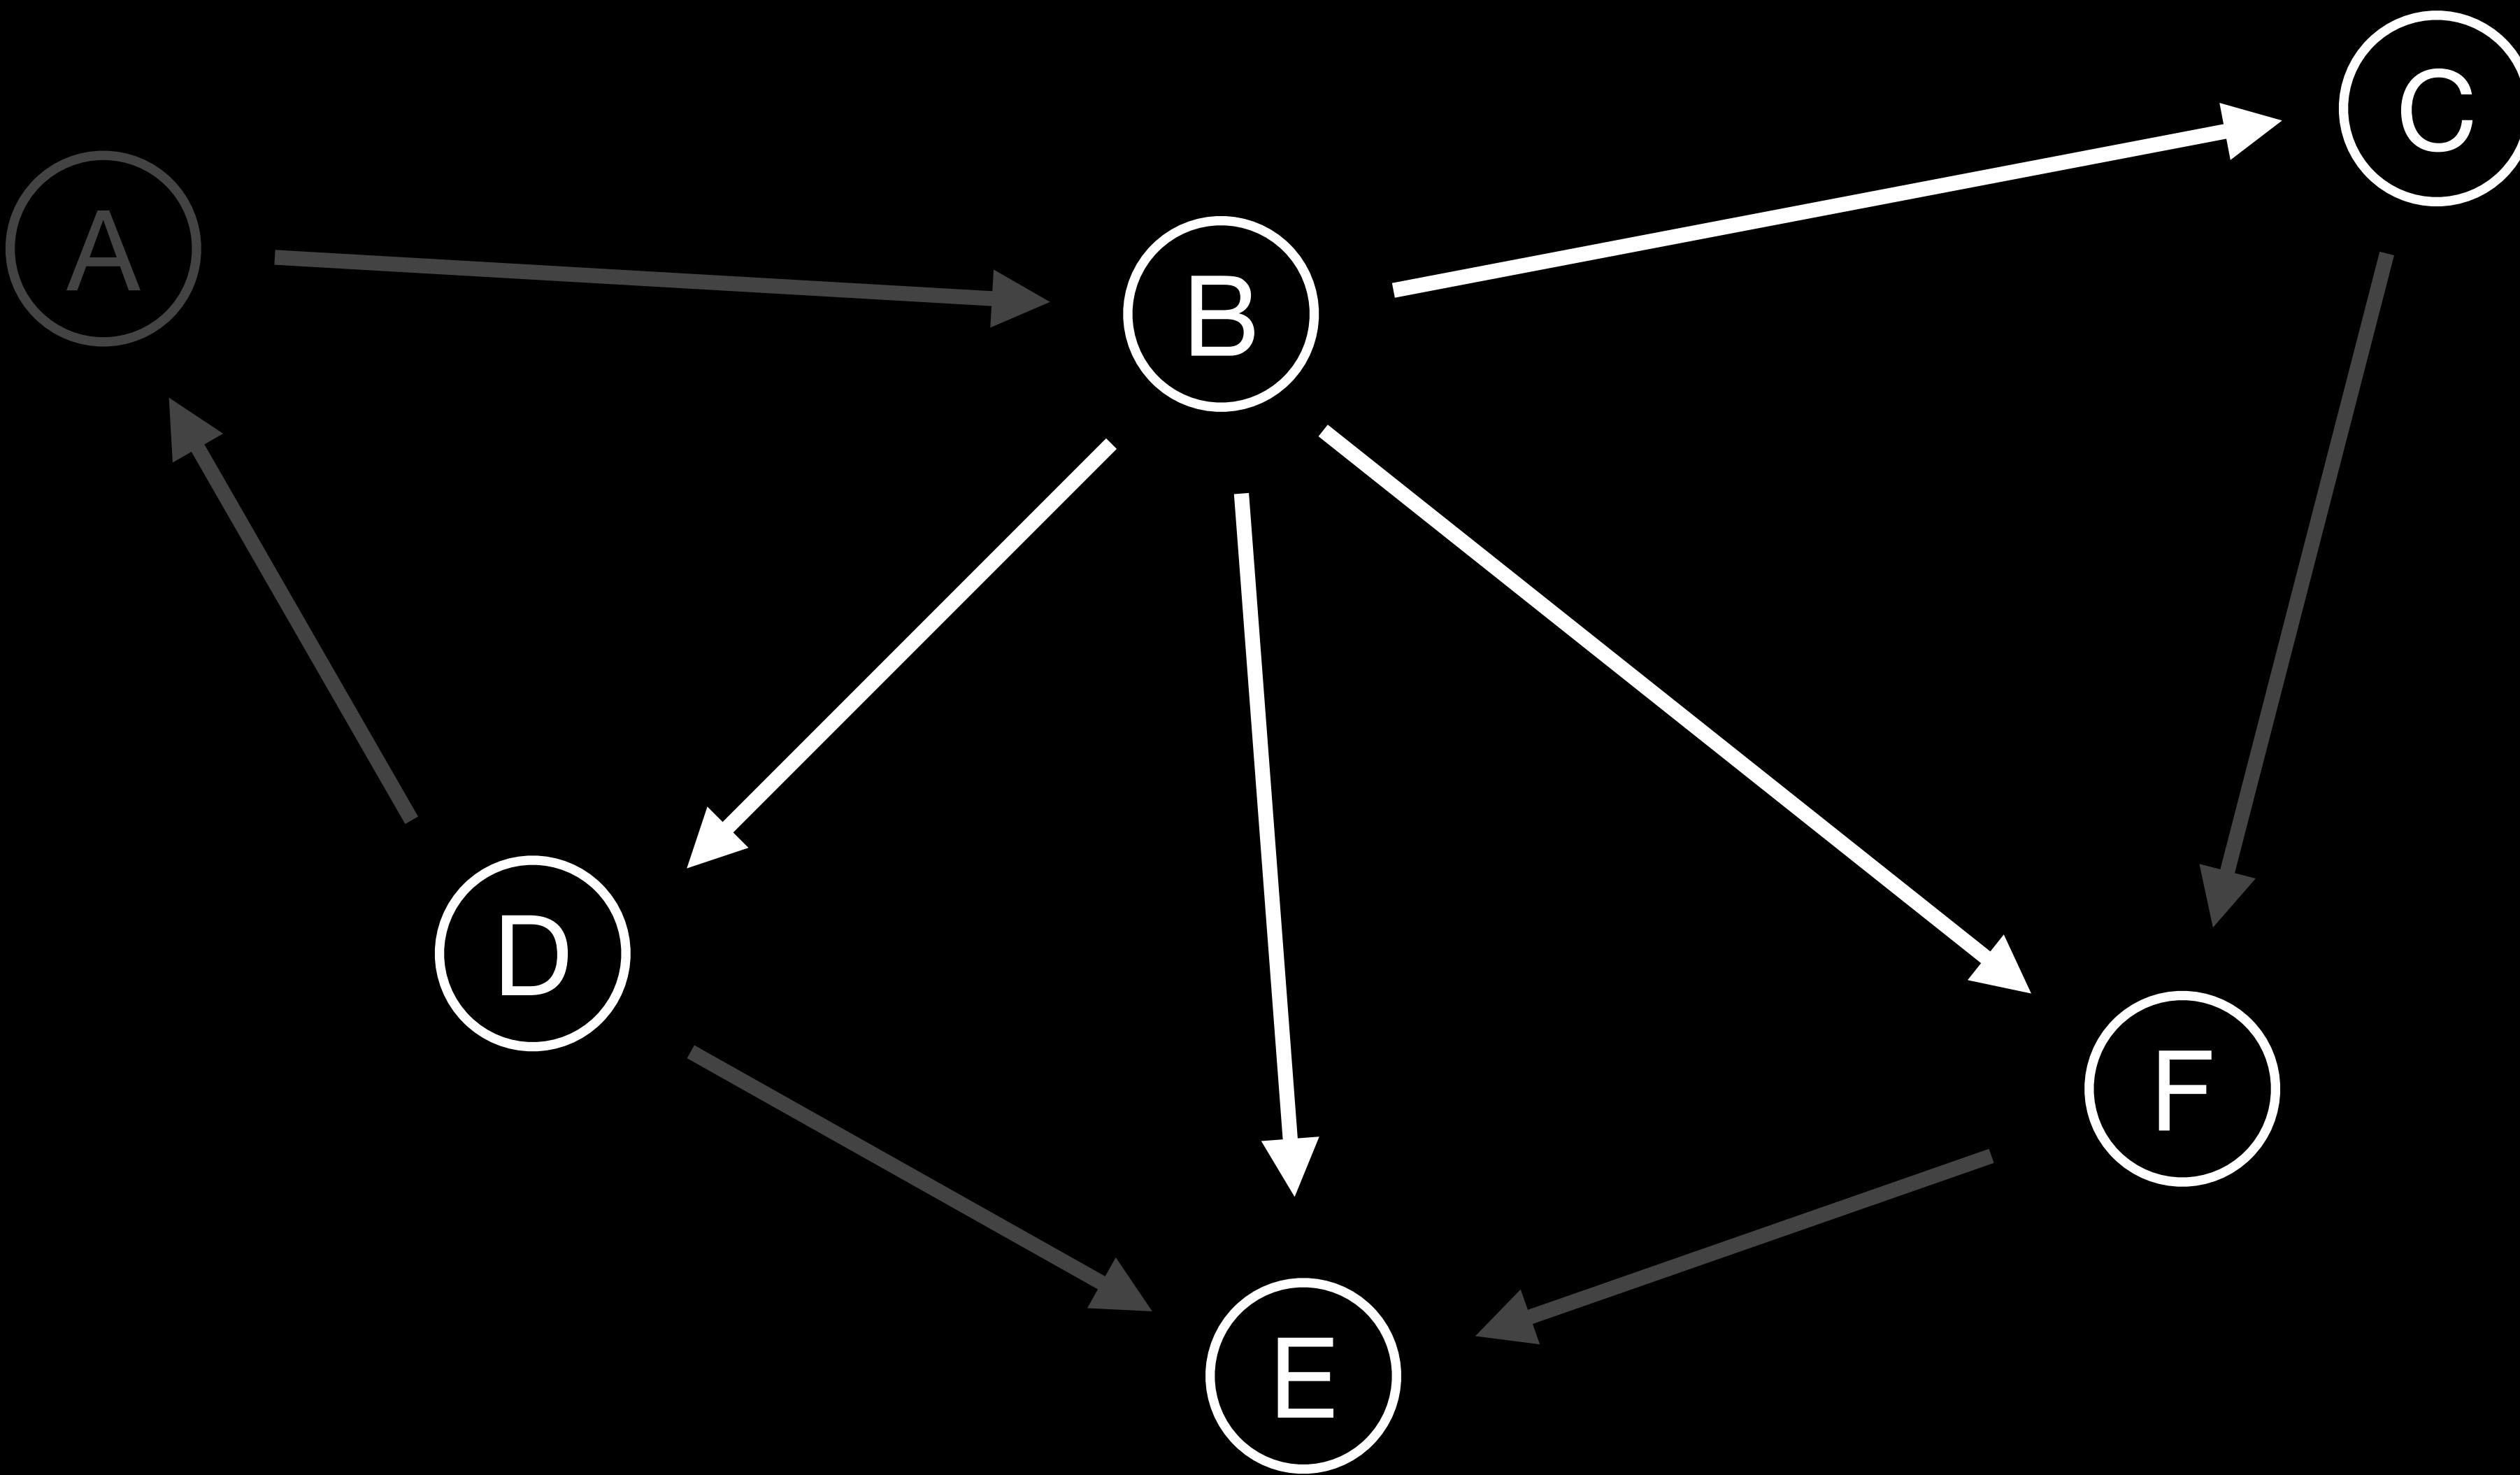
```
    tbnz x0, #3, after_bd
```

```
    bl visit_d
```

```
after_bd:
```

```
// ...
```

```
Visit all the  
nodes from B
```



And that's it!

And that's it!

Well, not quite!

And that's it!

Well, not quite!

There's a problem

visit\_a

visit\_b

visit\_d

visit\_a

visit\_b

visit\_d

call stack



visit\_a        sets x30  
visit\_b        sets x30  
visit\_d

visit\_a

visit\_b

visit\_d



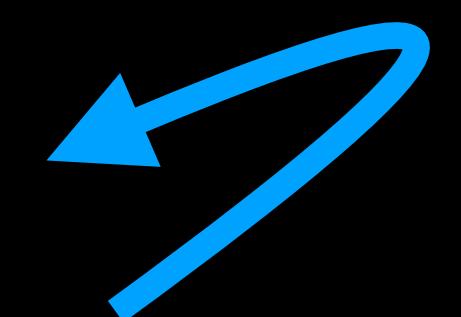
return

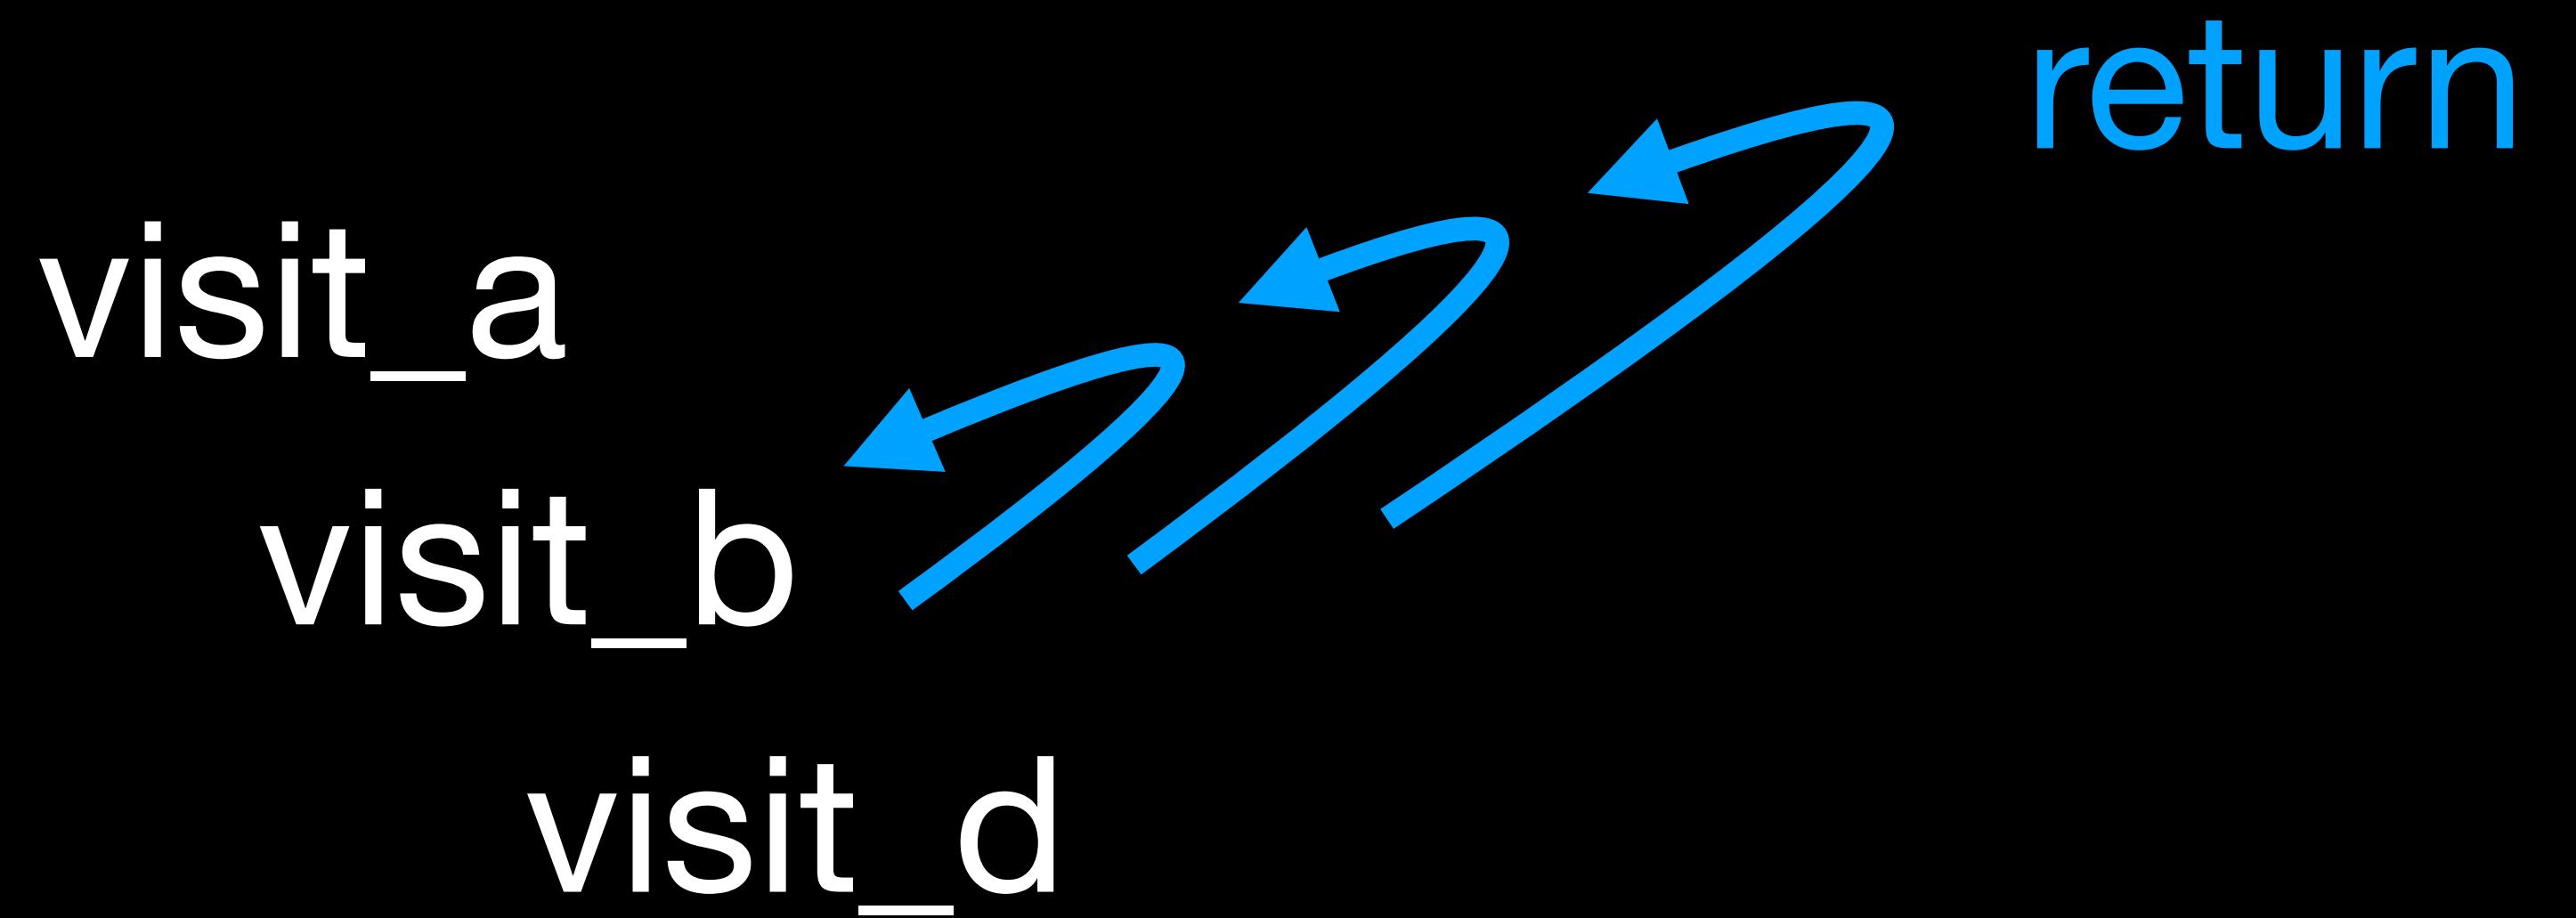
visit\_a

return

visit\_b

visit\_d





We need a stack!

visit\_a

visit\_b

visit\_d

call stack

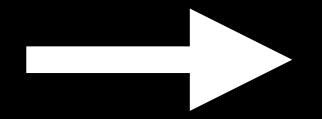


# How the stack works

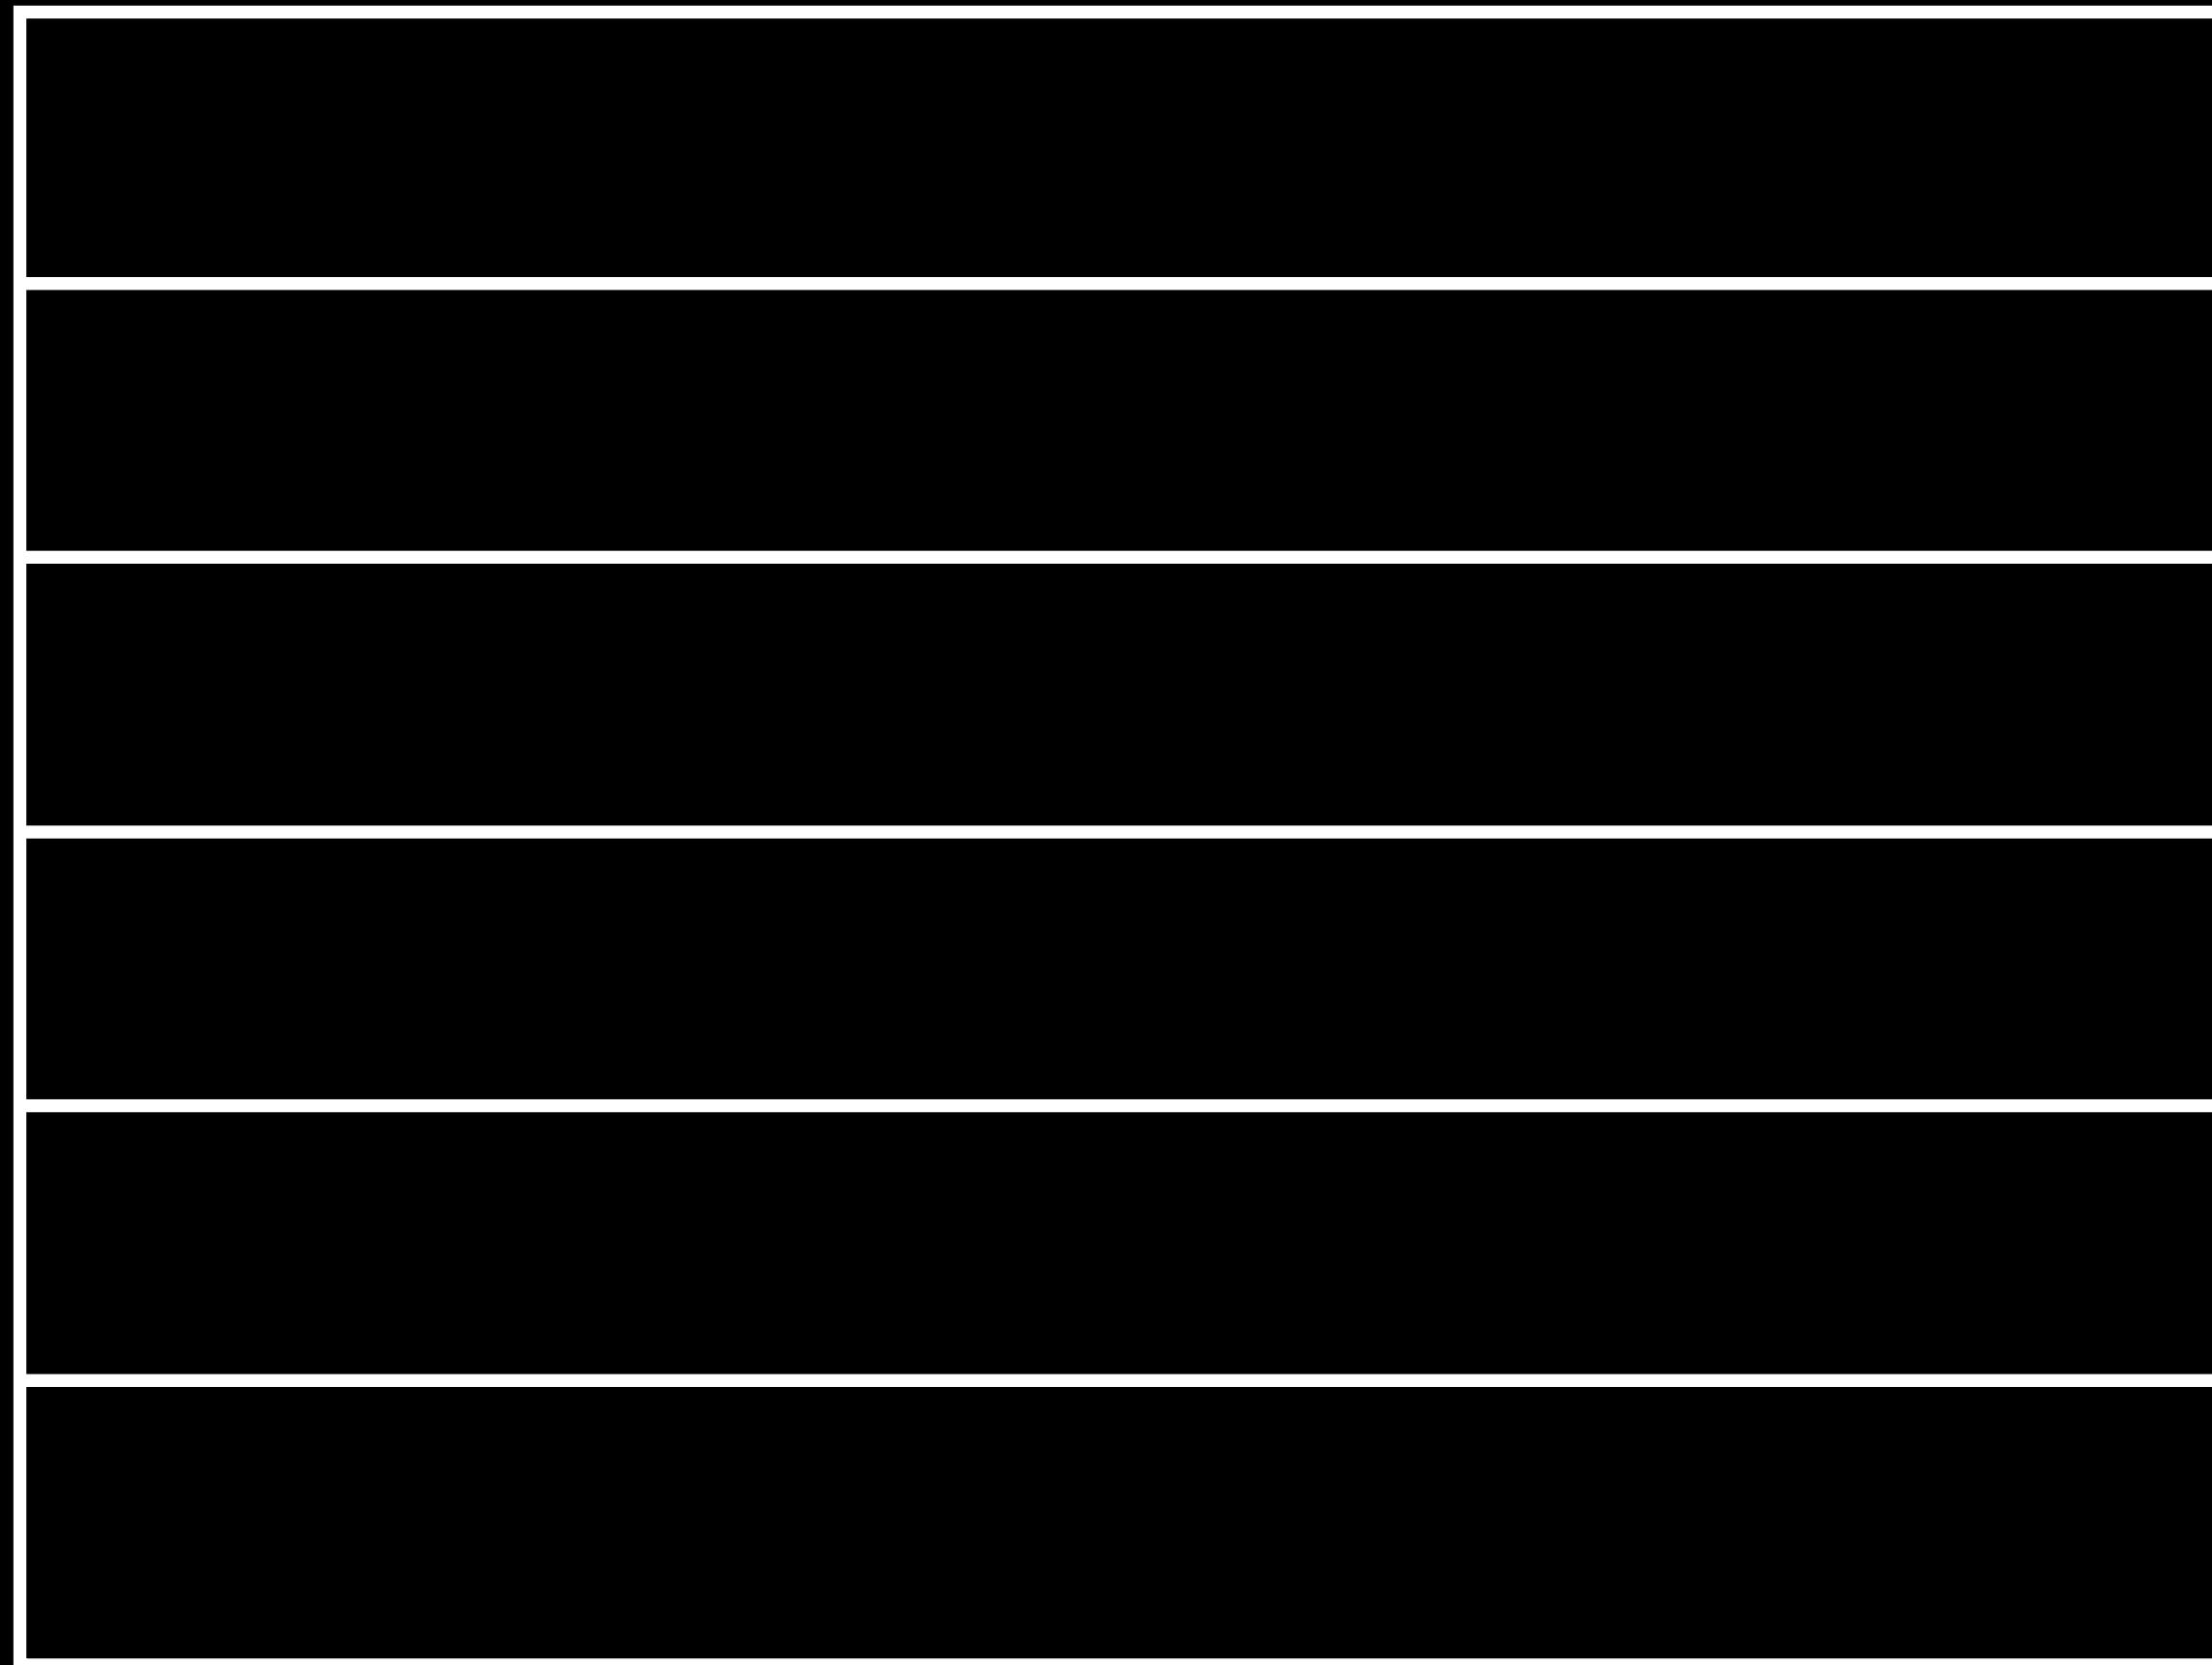
Stored in memory

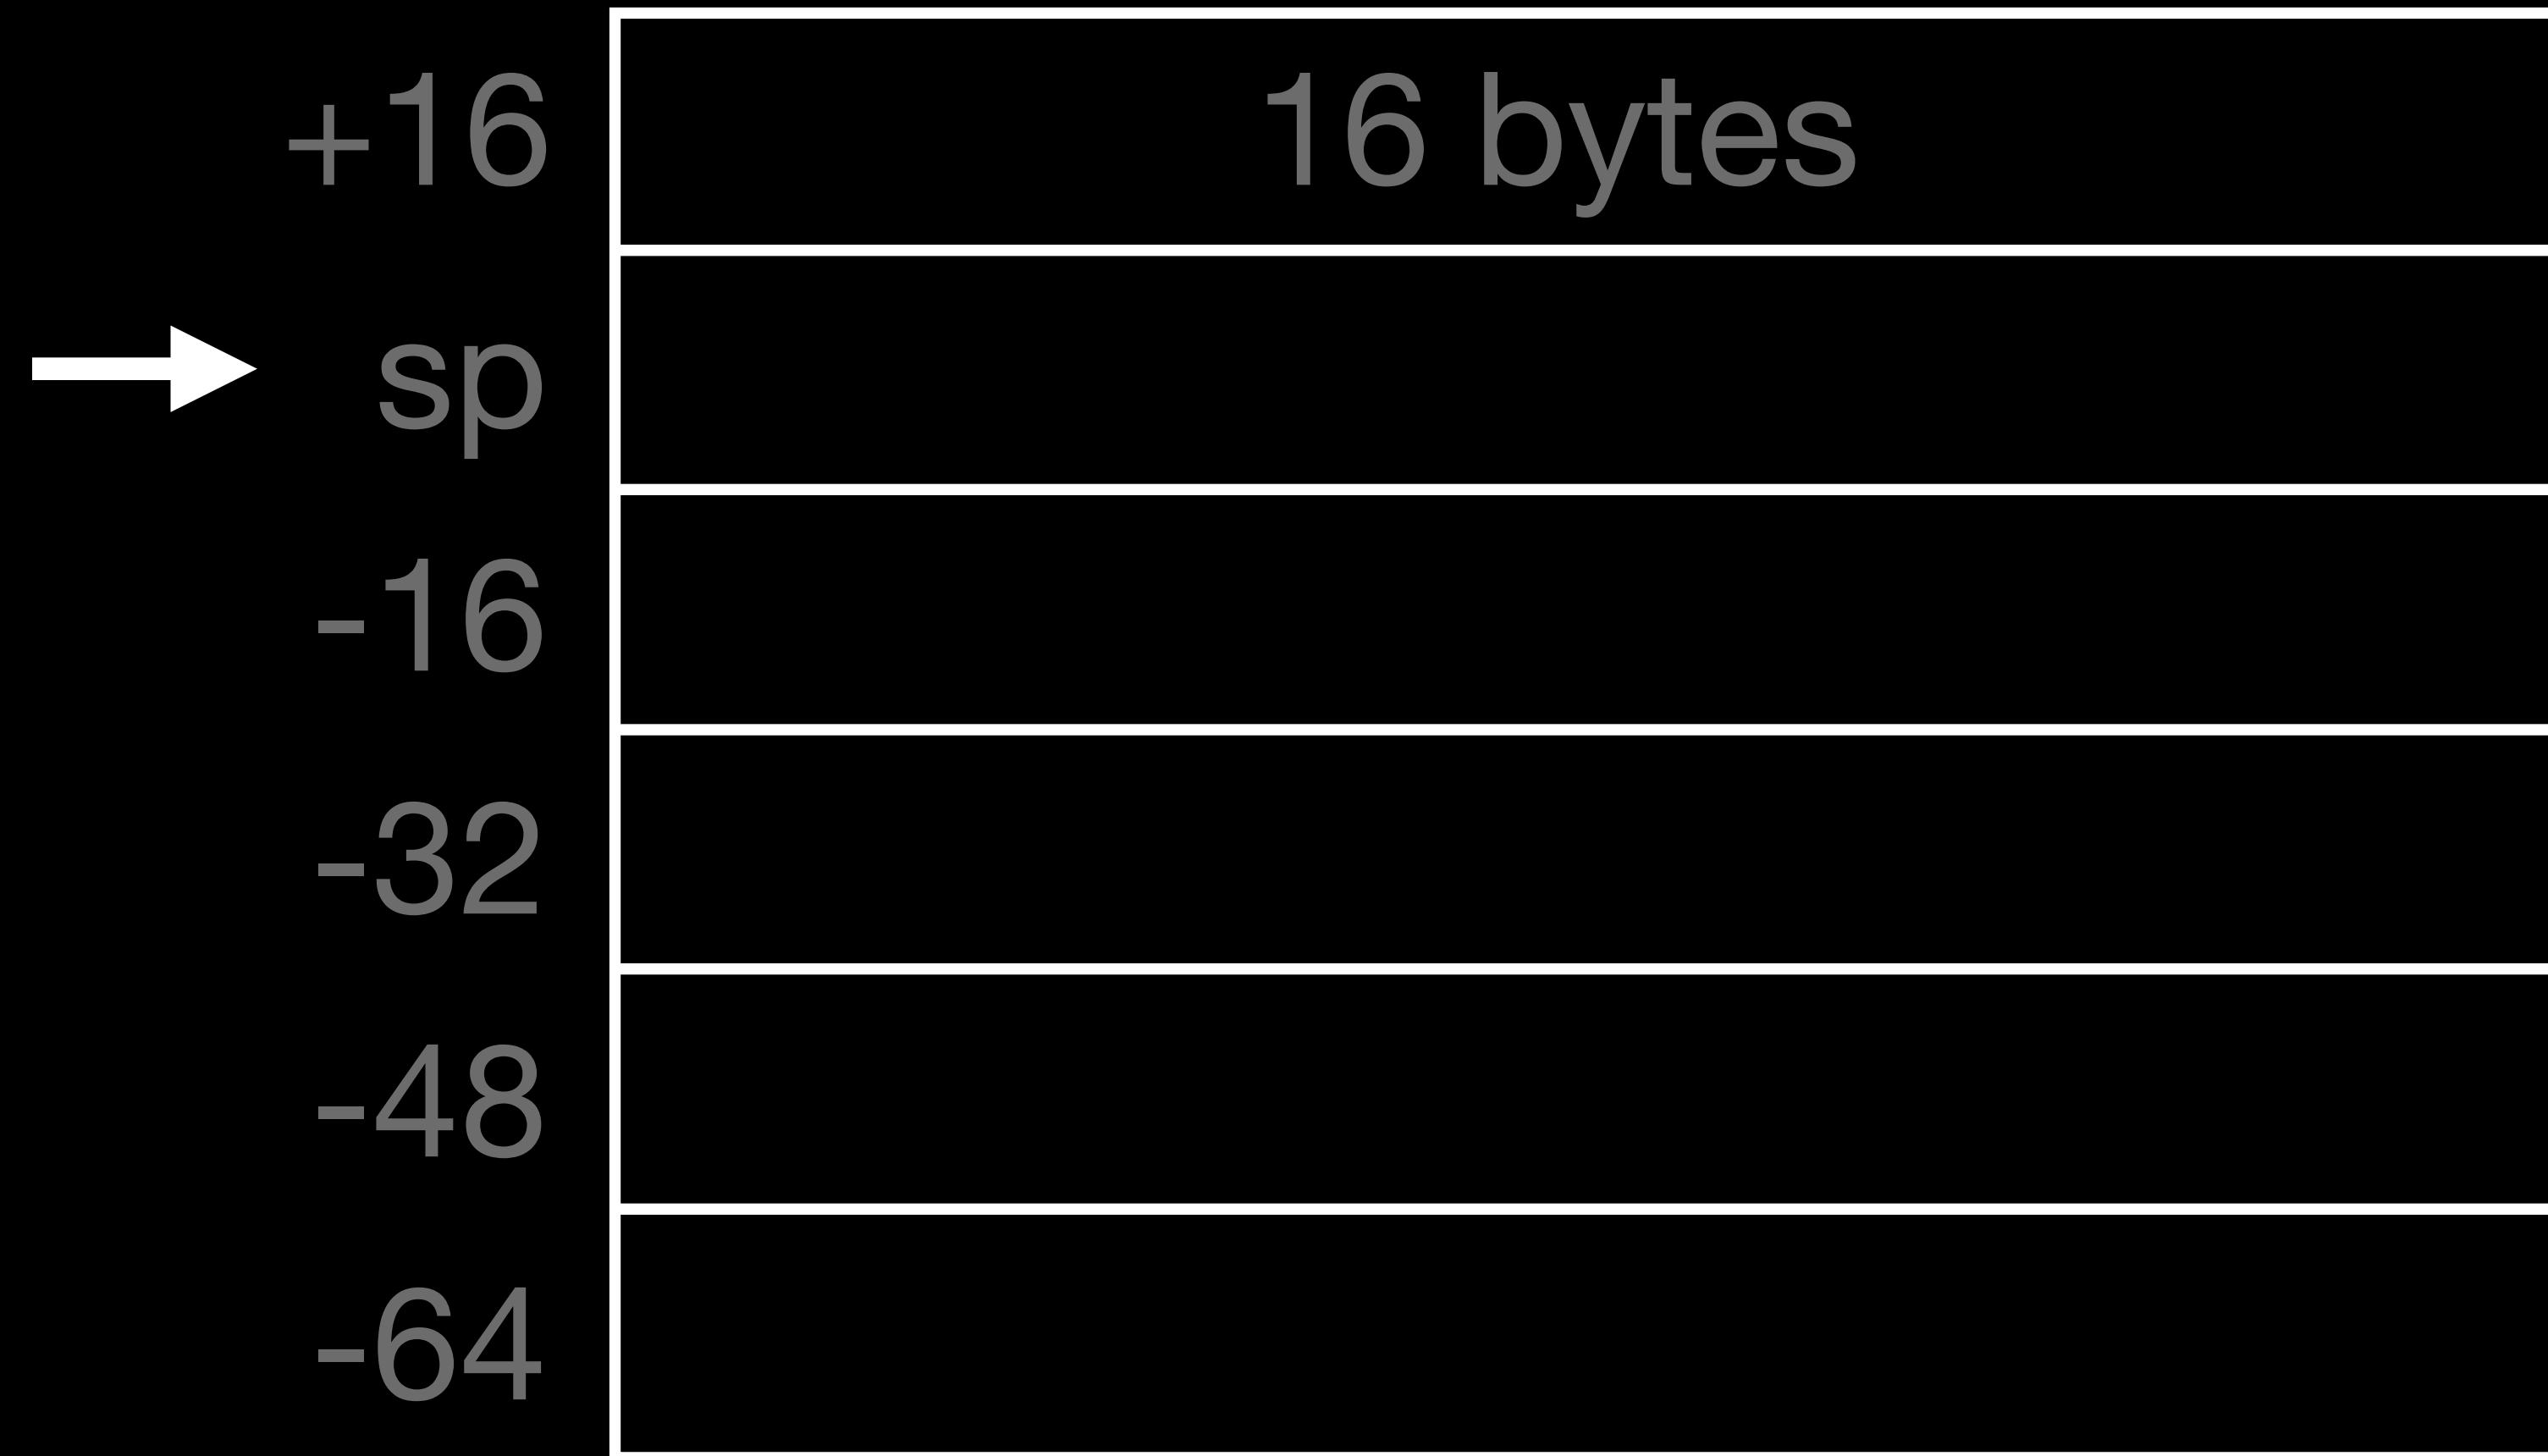
Stack pointer (sp) register

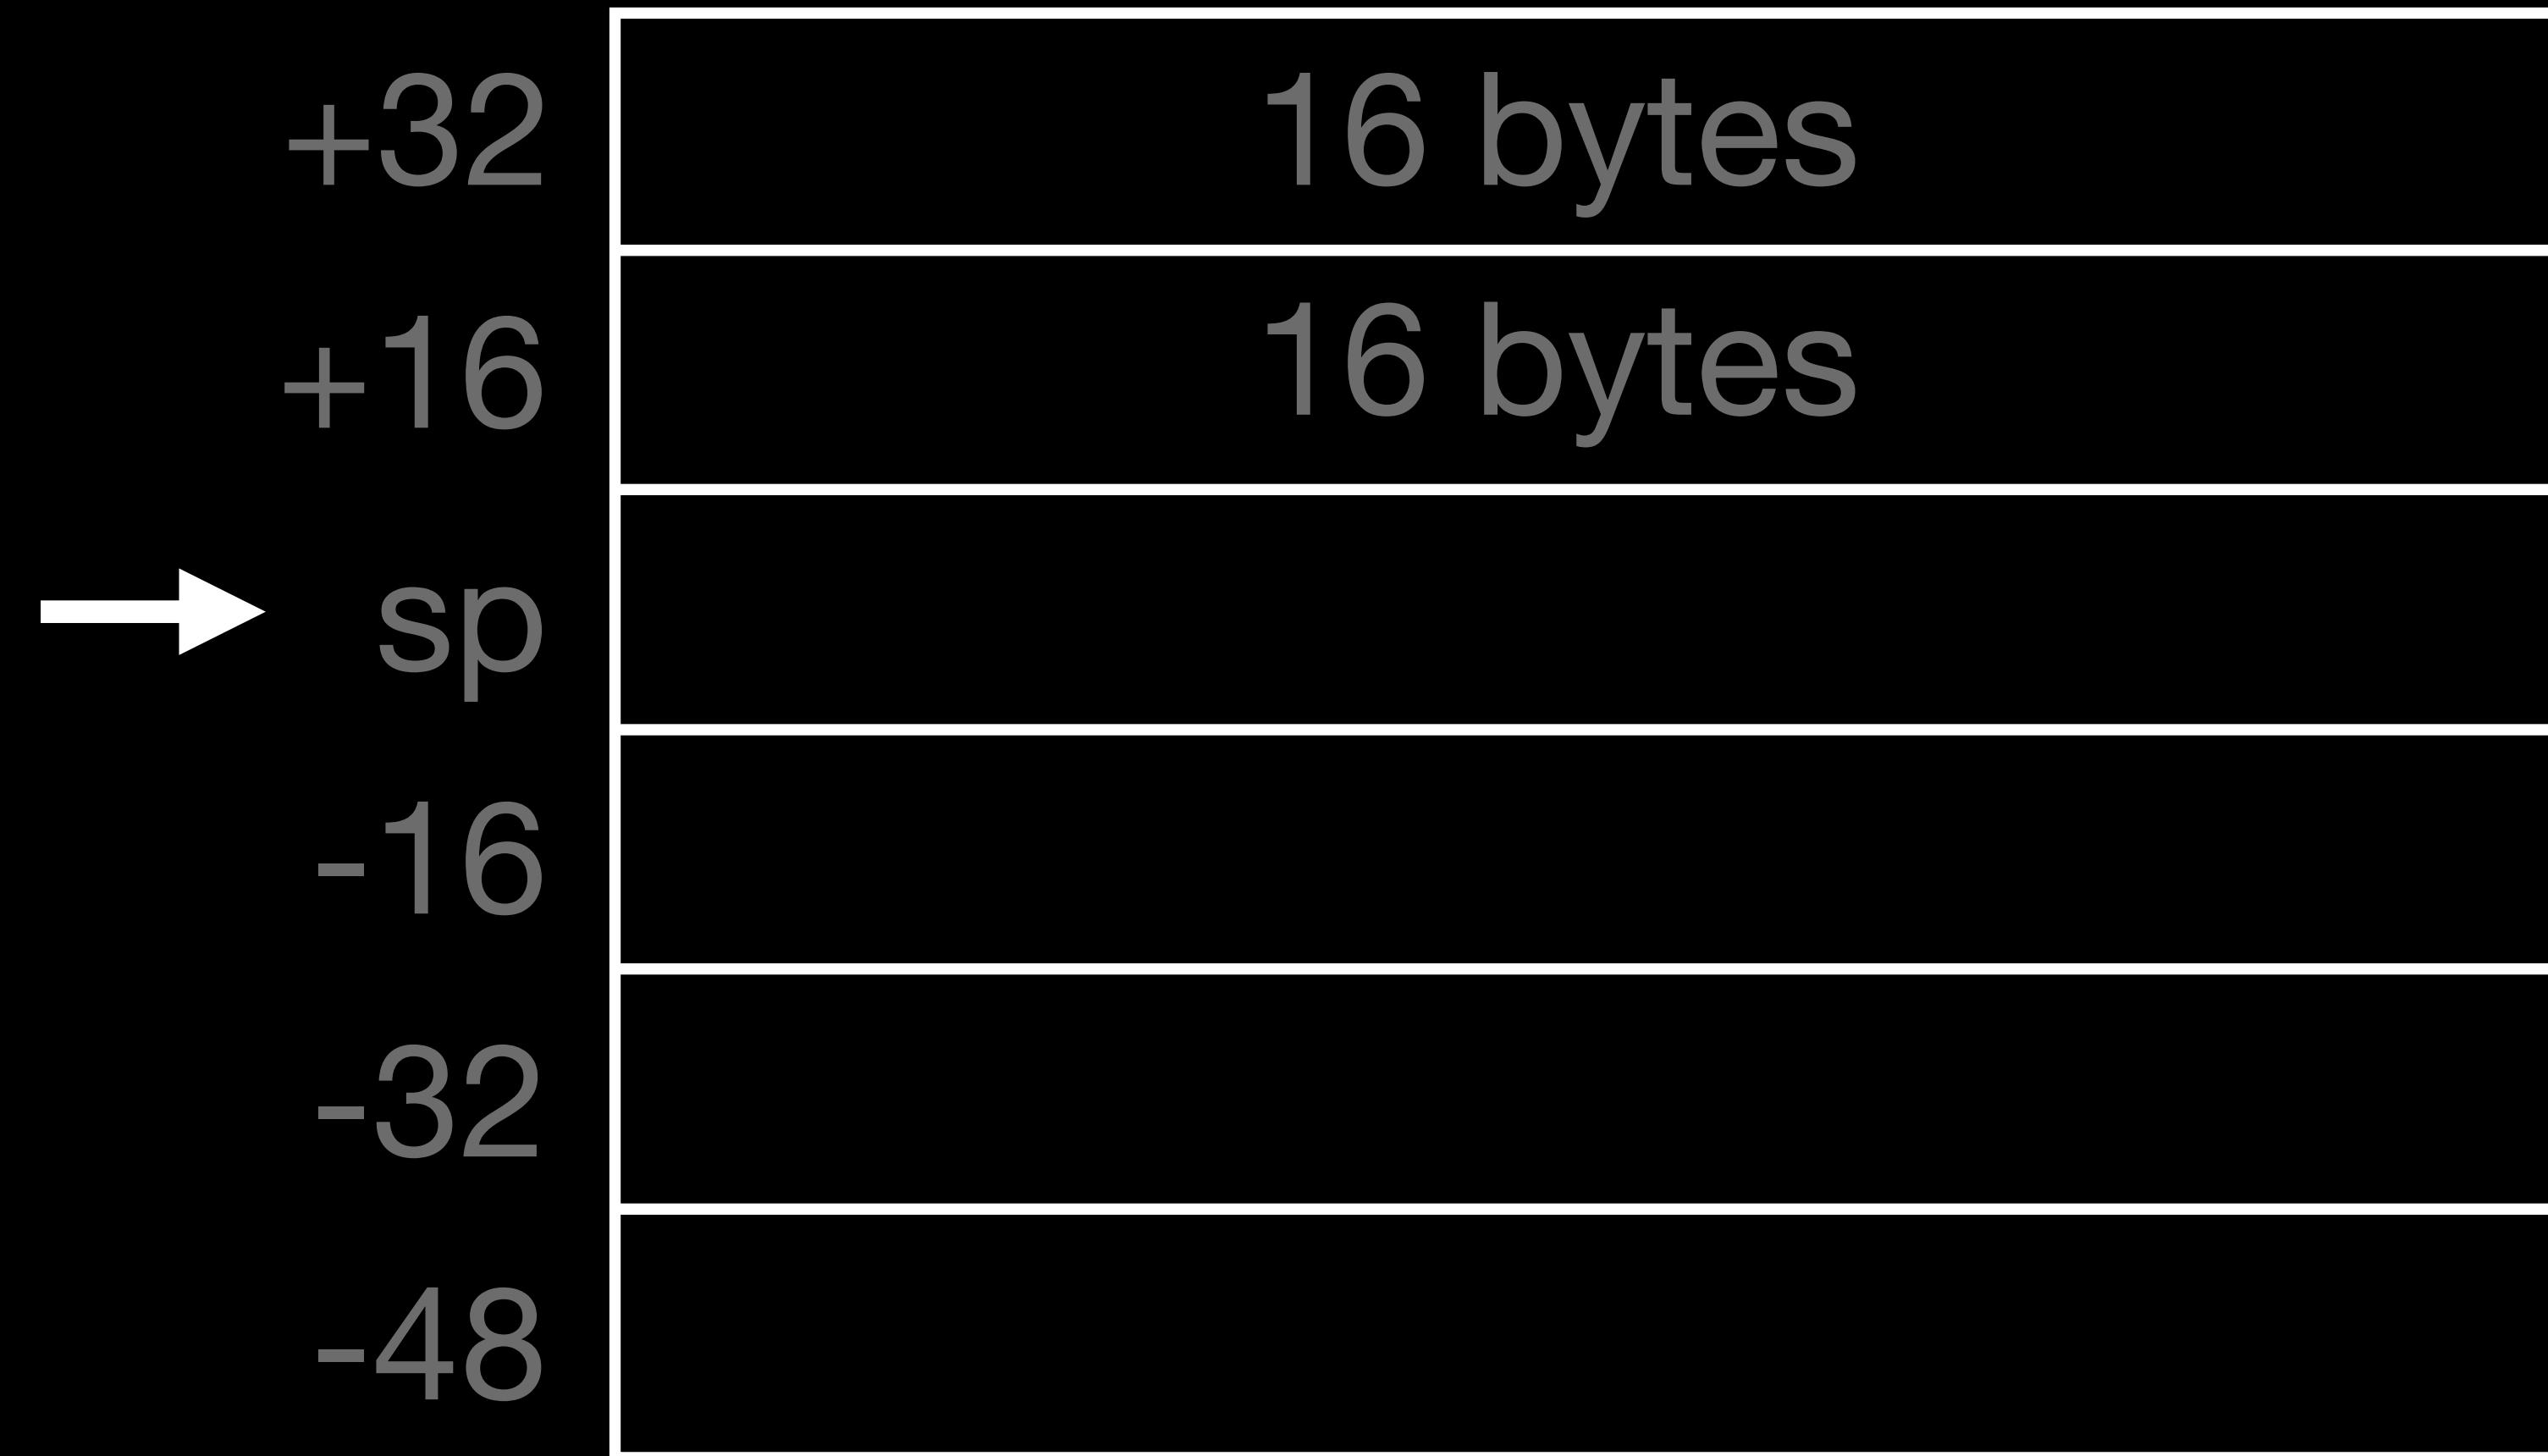
Address moves down



sp  
-16  
-32  
-48  
-64  
-80



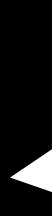




```
str x30, [sp, -16]!  
ldr x30, [sp], 16
```

```
str x30, [sp, -16]! ← Push x30  
ldr x30, [sp], 16
```

```
str x30, [sp, -16]!  
ldr x30, [sp], 16
```



Pop x30

visit\_a:

str x30, [sp, -16]!

// ...

Code from before

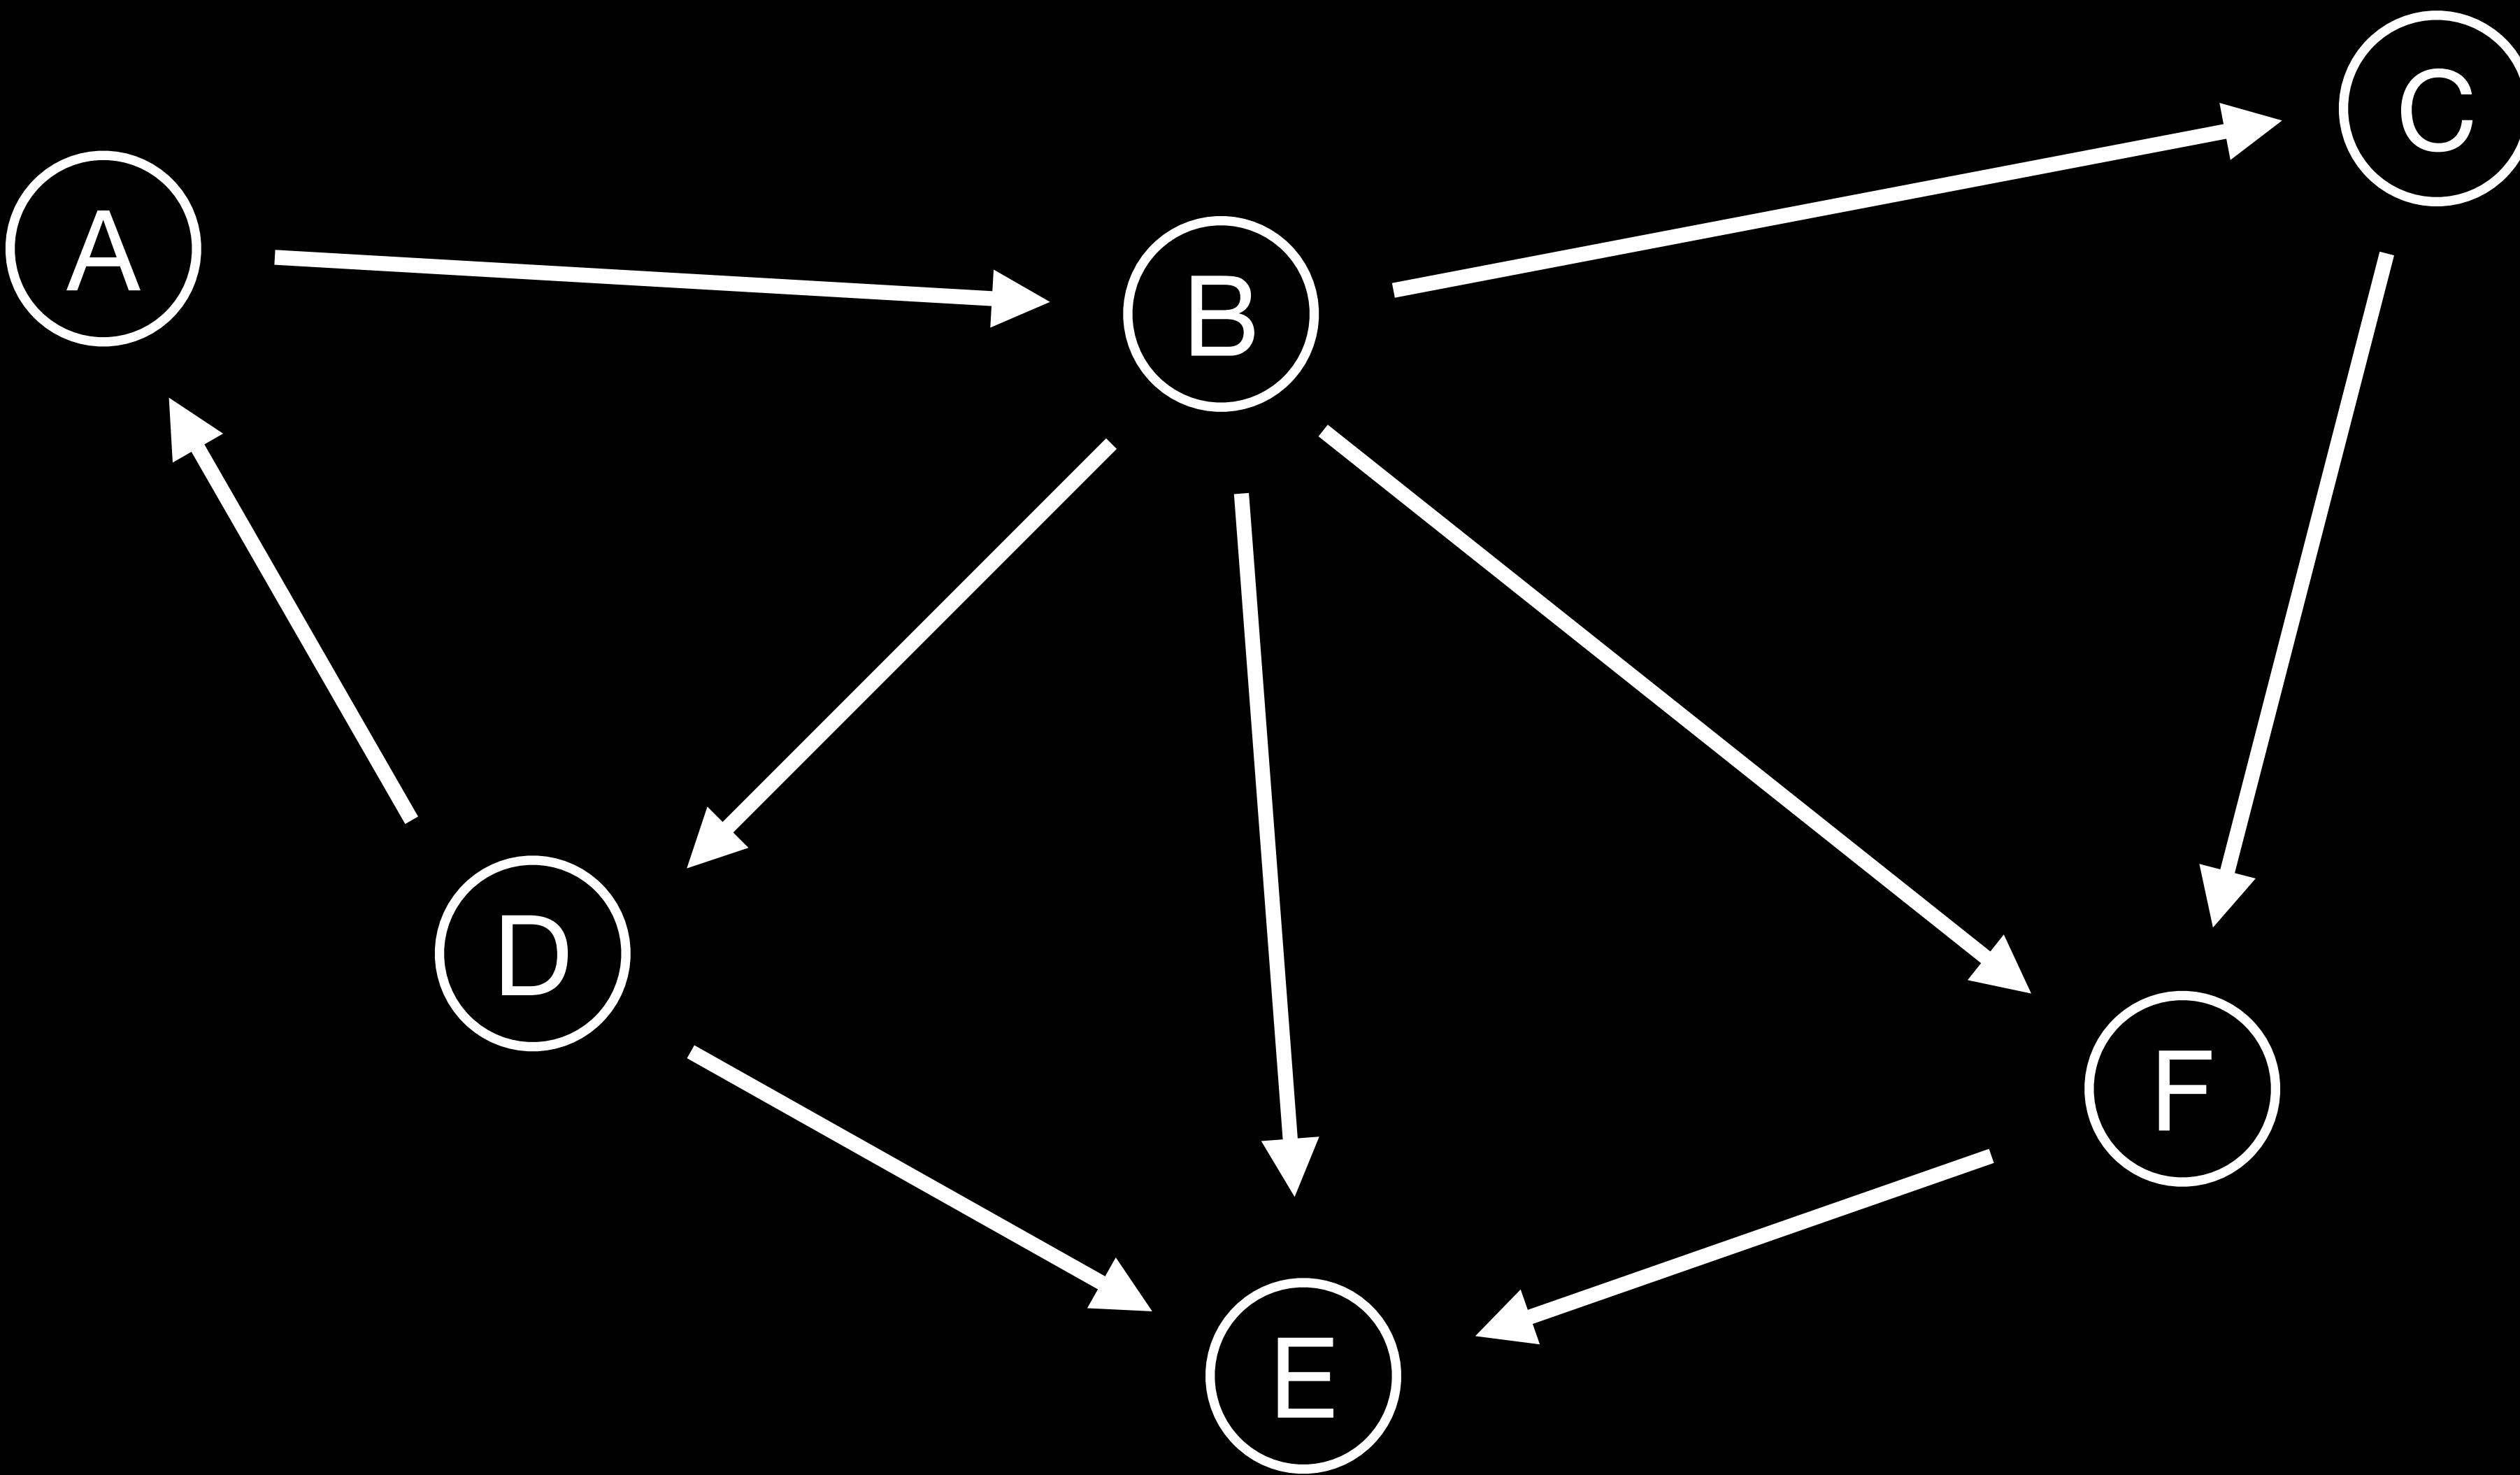
ldr x30, [sp], 16  
ret

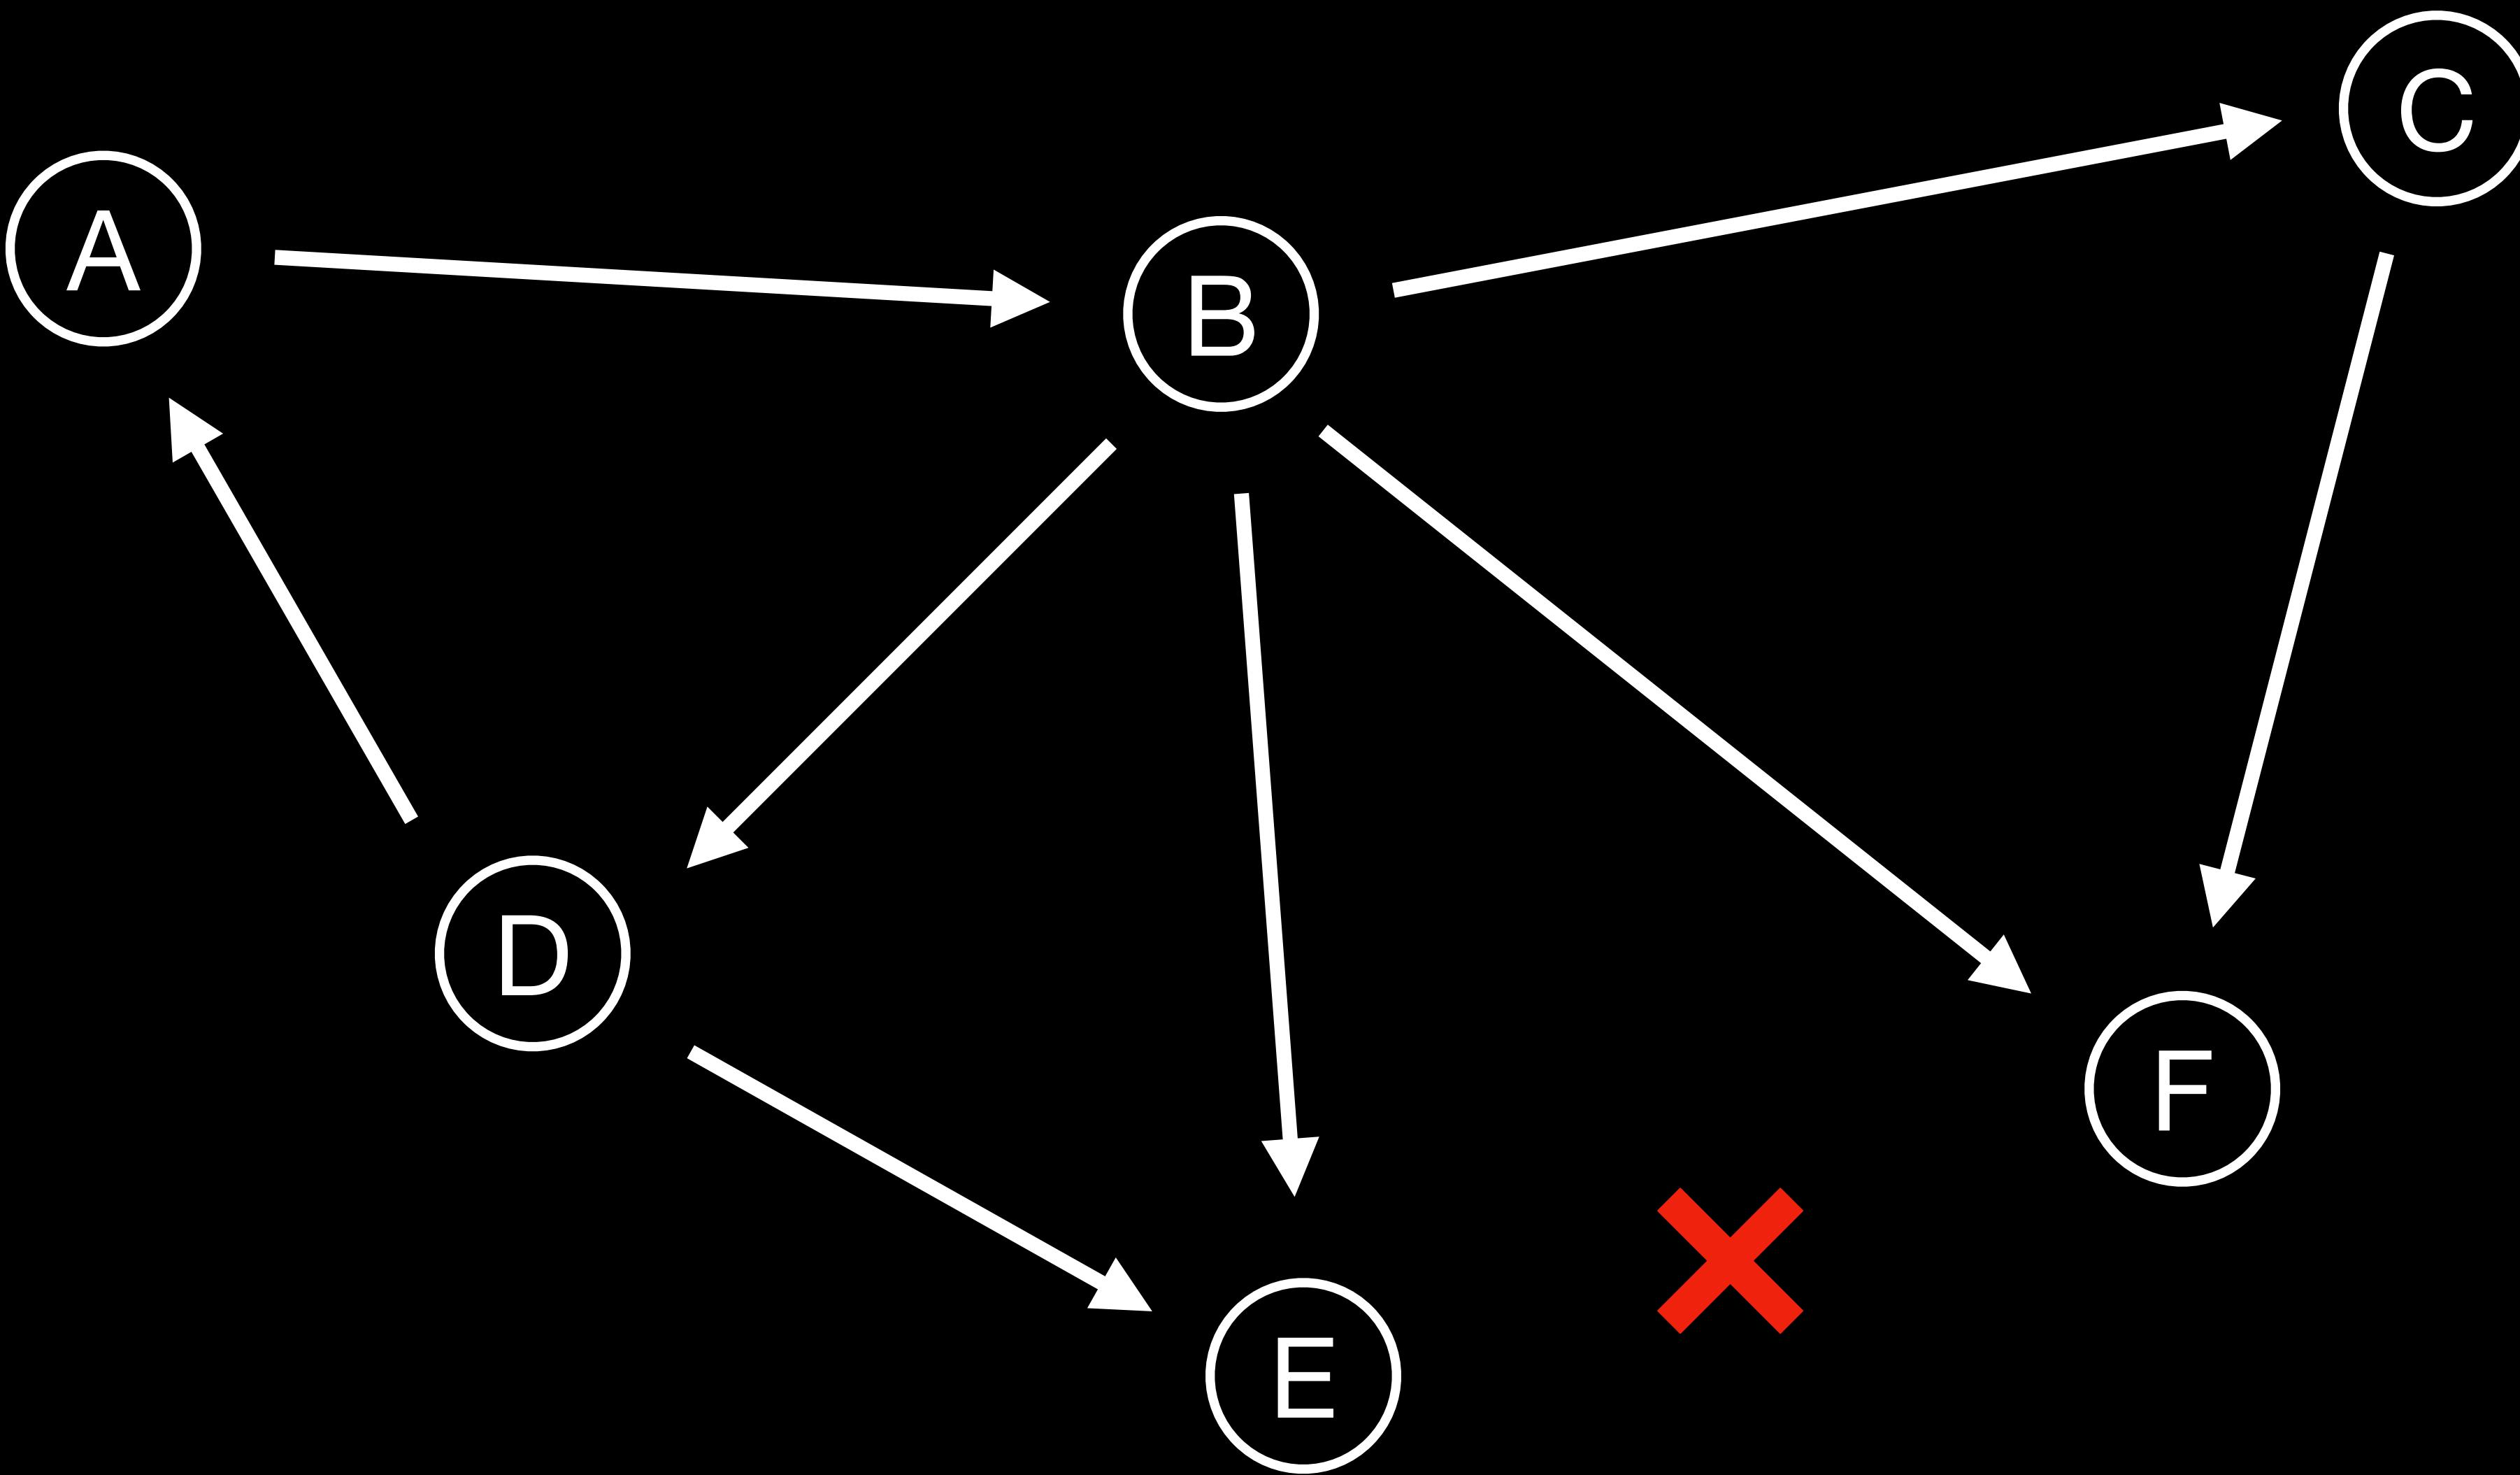
And that's it!

(really)

# Demo

<https://github.com/tuzz/assembly/blob/main/src/hamiltonian.s>





By the end

Access to a powerful tool

Know more about your computer

Where to go next

# Useful resources

<https://modexp.wordpress.com/2018/10/30/arm64-assembly>

<https://www.youtube.com/watch?v=GBRdzaAxHB8>

<https://www.amazon.co.uk/dp/1484258800>

<https://developer.arm.com/architectures/cpu-architecture/a-profile/exploration-tools>

Thanks!

@chrispatuzzo



